

Contents

1	Introduction	3
2	Literature	5
3	Preliminaries	8
3.1	$M/M/1$ queue	8
3.2	Non-preemptive priority and the PASTA property	8
3.3	$M/M/1+M$ queue (Erlang-A)	9
3.4	Pollaczek–Khinchine formula	10
3.5	Confluent hypergeometric function	10
3.6	First-order linear ODEs and PDEs	10
3.6.1	Integrating factor	11
3.6.2	Variation of parameters	11
3.6.3	Method of Characteristics	11
4	Model description: Model-X	11
4.1	Balance equations	13
4.1.1	Balance equations for state space S	13
4.1.2	Balance equations for state space $\tilde{S}$	15
4.2	Derivation of equations for (Partial) Probability Generating functions	16
4.2.1	Equation for PGF in state space S	16
4.2.2	Equation for PPGF in state space $\tilde{S}$	26
5	Analysis of Model-A ($\gamma_1 = \gamma_2 = 0, \theta_1 = \theta_2 = 0$)	32
5.1	Analysis of Model- A : state space S	32
5.1.1	Analytical approach for Model- A in state space S	33
5.1.2	Probabilistic approach for Model- A in state space S	34
5.1.3	Partial-PGF recurrence approach for Model- A in state space S	35
5.2	Analysis of Model- A on $\tilde{S}$: scope and limitations of the partial-PGF recursion	37
5.3	Mean queue lengths and mean waiting times	39
5.3.1	Class-1 marginal PGF and mean queue length	39
5.3.2	Total queue length from the diagonal	39
5.3.3	Class-2 mean queue length by subtraction	40
5.3.4	Mean waiting times via Little’s Law	40
6	Analysis of Model-B ($\gamma_1, \gamma_2 > 0, \theta_1 = \theta_2 = 0$)	40
7	Analysis of Model-B_2 ($\gamma_1 > 0, \gamma_2 = 0, \theta_1 = \theta_2 = 0$)	46
7.1	Reduced fundamental equation	46
7.2	Mean queue lengths and mean waiting times	51
8	Analysis of Model-C_2 ($\theta_1 > 0, \gamma_1 = \gamma_2 = \theta_2 = 0$)	53
8.1	Setup: reduction to a first-order linear ODE in x	54
8.2	Probabilistic derivation of $P_y(y)$	55
8.3	Analytical re-derivation of $P_y(y)$ via boundedness at $x = 1$	56
8.4	Recovery of $P(x, y)$ from $P_y(y)$	58
8.5	Mean queue lengths and mean waiting times	60
9	Analysis of Model-B_2^H ($\gamma_1 > 0$ head-of-line, $\gamma_2 = \theta_1 = \theta_2 = 0$)	62
9.1	Balance equations	62
9.2	Fundamental PGF equation	62
9.3	Closed-form solution	63
9.3.1	Analytical approach for Model- B_2^H	64
9.3.2	Probabilistic approach for Model- B_2^H	65
9.4	Limits and sanity checks	65
9.5	Mean queue lengths and mean waiting times	65
9.5.1	Class-1 marginal PGF and mean queue length	65
9.5.2	Total queue length from the diagonal	65
9.5.3	Class-2 mean queue length by subtraction	66
9.5.4	Mean waiting times via Little’s Law	66

10 Analysis of Model-C_2^H ($\theta_1 > 0$ head-of-line, $\gamma_1 = \gamma_2 = \theta_2 = 0$)	66
10.1 Balance equations	67
10.2 Fundamental PGF equation	67
10.3 Closed-form solution	68
10.3.1 Analytical approach for Model- C_2^H	68
10.3.2 Probabilistic approach for Model- C_2^H	69
10.4 Limits and sanity checks	69
10.5 Mean queue lengths and mean waiting times	70
10.5.1 Class-1 marginal PGF and mean queue length	70
10.5.2 Class-2 mean queue length	70
10.5.3 Mean waiting times via Little's Law	70
10.6 Comparison of the head-of-line models	71
11 Comparison of the Solved Models	71
11.1 Traffic-intensity sweep	74
11.2 The jockeying mechanism (Model B_2)	76
11.3 Class asymmetry	78
11.4 The abandonment mechanism and the throughput-quality tradeoff (Model C_2)	80
11.4.1 The price in lost customers	81
11.5 Head-of-line versus full-rate mechanisms (Models B_2^H and C_2^H)	84
12 Results	85
12.1 Validation	85
12.2 Mean queue lengths and mean waiting times	88
12.3 Head-of-line models B_2^H and C_2^H	89
12.4 Cross-validation: simulation versus the CTMC	90
12.5 Convergence to Model A	91
12.5.1 Model C_2 as $\theta_1 \rightarrow 0^+$	92
12.5.2 Model B_2 as $\gamma_1 \rightarrow 0^+$	93
12.5.3 Instant-jockeying limit $\gamma_1 \rightarrow \infty$	93
12.5.4 Head-of-line models B_2^H and C_2^H	94
12.6 Comparative dashboard	95
13 Conclusion and Future Work	95
13.1 The obstruction taxonomy	95
13.2 The mechanism dichotomy	96
13.3 Operational reading and the deterioration-direction caveat	96
13.4 Limitations	96
13.5 Future work	97

Exact Generating-Function Analysis of a Two-Class Non-Preemptive Priority $M/M/1$ Queue with Jockeying and Abandonment

Victor Dominguez Sainz

June 19, 2026

Abstract

We study the stationary queue-length distribution of a two-class, non-preemptive priority $M/M/1$ queue augmented with two waiting-room mechanisms: *jockeying*, in which a waiting customer switches to the other queue, and *abandonment*, in which a waiting customer leaves the system. Jockeying conserves the total number of customers, whereas abandonment does not, and combining both yields a model (Model- X) that is analytically intractable. We therefore present a collection of simpler models: Model- A (baseline), Model- B (bidirectional jockeying), Model- B_2 (jockeying from the priority class to the non-priority class) and Model- C_2 (abandonments in the priority class). We also study two head-of-the-line models: Model- B_2^H (only the first customer in the priority queue can jockey) and Model- C_2^H (only the first customer in the priority queue can abandon).

The central object is the bivariate Probability Generating Function (PGF) $P(x, y)$ of the two queue lengths, obtained by analytical and, where the structure allows, probabilistic arguments. The purpose of this work is to present a unified framework for deriving the functional equation of each model and to expose the mathematical challenges intrinsic to each. Model- A is classical and collapses to an algebraic expression by finding a kernel root inside the unit circle; Model- B_2 is solved by imposing analyticity at an interior point and Model- C_2 by imposing boundedness inside the unit circle. The head-of-the-line models, like the baseline Model- A , also yield an algebraic expression for the PGF. Two models are left open: Model- B reduces to an integral expression for the PGF involving a Volterra-type integral equation, and Model- X compounds this with nonlinear characteristic curves.

Every closed form is validated against a truncated continuous-time Markov Chain (CTMC) that recovers the baseline model as the mechanism rates vanish. The comparison reveals a key difference between jockeying and abandonment: the former merely redistributes delay between classes, whereas the latter substantially shortens the queue lengths.

1 Introduction

Motivation

Priority queues arise naturally wherever a common server must satisfy heterogeneous demand. The primary motivation for this thesis is healthcare operations, specifically the management of waiting lists for elderly-care facilities, where medical severity dictates a user's —hereafter *customer*— priority class. Here prolonged waiting may itself alter the system's dynamics: a lack of appropriate care can deteriorate a waiting customer's health, leading them to change priority status (jockeying) or leave the system (abandonment), whether to seek alternative care or through death. Despite this motivation, the work is fundamentally modelling-driven: we characterise the long-run behaviour of queueing systems in which jockeying and/or abandonment play a role.

This thesis analyses a 2-class $M/M/1$ queue with non-preemptive priority: each class arrives according to a Poisson process (rates λ_1 and λ_2), service times are exponential at a common rate μ , and the single server completes the current service even when a higher-priority customer arrives. To this we add the mechanisms of jockeying and abandonment. The model and its mechanisms are described in more detail in §4.

How these two mechanisms interact within a priority discipline, and how that interaction affects the difficulty of obtaining a closed-form Probability Generating Function (PGF), is the main focus of this thesis.

The analytical challenge

In full generality, the bivariate PGF for Model- X satisfies a functional equation with terms of three kinds: algebraic coefficients, partial derivatives in x —from jockeying and abandonment in class-1— and partial derivatives in y —from the same mechanisms in class-2. When both kinds of derivative term are present, the equation belongs to a family of integro-differential equations whose solution would require more advanced tools. The problem is moreover asymmetric: derivative terms from the non-priority class complicate it more than those from the priority class alone.

We therefore solve less complex models obtained by setting some of the jockeying and abandonment parameters to zero. Table 1 lists the models studied here and indicates whether a PGF expression is obtained for each:

Model	Active mechanism	Equation type	Status
A	none	algebraic PGF equation	solved
B	bidirectional jockeying	PDE, Volterra family	open
X	jockeying and abandonment	PDE, transcendental characteristics	open
B_2	one-way jockeying ($1 \rightarrow 2$)	ODE in x , Kummer	solved
B_2^H	one-way jockeying, head-of-line	algebraic PGF equation	solved
C_2	class-1 abandonment	ODE in x , Kummer	solved
C_2^H	class-1 abandonment, head-of-line	algebraic PGF equation	solved

Table 1: Models studied in this thesis, with the active mechanism, the resulting functional-equation type, and the solution status. The head-of-line variants B_2^H and C_2^H restrict the mechanism to the customer at the head of Queue 1, which collapses the derivative term to an algebraic one and renders the model solvable in closed form.

Some models —notably Model- X and Model- B — are left incomplete because they involve derivative terms in y ; others, such as a potential Model- B_1 , Model- C or Model- C_1 , are not considered for the same reason, as they all reduce to a Volterra integral equation lying beyond the scope of this thesis. Model- B serves as our representative example of how such problems reduce to that integral equation; the remaining models are solvable. The last two entries, Models B_2^H and C_2^H (§§9 and 10), are *head-of-line* simplifications in which only the customer at the head of Queue 1 may jockey or abandon. The mechanism then acts at the constant rate $\gamma_1 \mathbf{1}_{\{n_1 \geq 1\}}$ or $\theta_1 \mathbf{1}_{\{n_1 \geq 1\}}$ rather than the length-proportional $\gamma_1 n_1$ or $\theta_1 n_1$, so the fundamental equation stays algebraic and is solved by the Kernel Method exactly as in Model- A .

Methods

Three complementary proof strategies are used throughout.

Kernel method / ODE with integrating factor. For Model- A , which has no derivative terms, the algebraic coefficient of $P(x, y)$ vanishes on a curve (the *kernel*); evaluating the equation there determines the remaining unknowns. Model- B_2 and Model- C_2 , which do carry derivative terms in x , reduce to a first-order linear ODE in x with an explicit integrating factor, and their unknowns follow from a regularity argument at a singularity of that factor. We present both strategies together, as they share the same underlying idea: the latter extends the former to the case where the unknowns are independent of the derivative terms and can be taken outside the integral.

Probabilistic argument via Pollaczek–Khinchine and the effective service time. For Model- A and Model- C_2 , class-2 arrivals are Poisson, so the PASTA property holds and we can apply the Pollaczek–Khinchine formula together with the *effective service time* seen by a class-2 customer. This is an $M/G/1$ queue whose service times equal the busy period of the class-1 subsystem —an $M/M/1$ queue for Model- A and an $M/M/1 + M$ system for Model- C_2 . In both cases the argument independently confirms the analytic results. Since jockeying destroys the Poisson structure, it does not apply to the other models.

Partial PGF (PPGF) in the alternative state space $\tilde{S}$. For the baseline Model- A we introduce an alternative state space indexed by the number of class-2 customers —whose distribution is hard— and the total number of customers —which is easy, since the total is conserved and behaves like a regular $M/M/1$ queue. This yields a PPGF with a non-homogeneous term. Cohen’s recursion argument [Coh56] was developed for the regular state space; we explore, with limited success, the

challenges of carrying it over to this alternative space and assess whether Cohen’s trick can deliver a feasible rough approximation.

Structure of the thesis

Section 2 reviews the relevant literature and Section 3 collects the probabilistic and analytic background used throughout this work. Section 4 defines the full model, Model- X , and derives the general balance equations and the fundamental PGF equations, both for the regular and the alternative state space. The analysis then proceeds from baseline to synthesis: Section 5 solves the baseline Model- A in closed form; Section 6 shows that bidirectional jockeying (Model- B) reduces the problem to a Volterra-type integral equation and is left open. Sections 7 and 8 then recover tractability in the one-way-jockeying (Model- B_2) and class-1-abandonment (Model- C_2) specialisations, while Sections 9 and 10 show how the head-of-line variants, Models B_2^H and C_2^H , collapse the functional equation to an algebraic one solved in closed form. Finally, Section 11 compares the solved models, Section 12 presents the numerical validation and performance metrics, and Section 13 concludes by abstracting the obstruction taxonomy and setting out the future work.

2 Literature

The study of two-class priority queues with departure mechanisms sits at the intersection of three well-developed streams: the exact generating-function analysis of non-preemptive priority queues, the theory of queues with customer impatience (abandonment), and the theory of jockeying and dynamic priority changes. The applied motivation throughout is the management of access to care — elderly-care waiting lists, emergency departments, and transplant registries — in which the two mechanisms studied here carry direct clinical readings: a waiting patient whose condition deteriorates changes priority class (jockeying), while one who leaves the list by death or transfer reneges (abandonment). This thesis contributes an exact, single-server treatment that admits jockeying and abandonment within a unified bivariate-PGF framework, filling a gap that the multi-server and approximate literature has so far left open.

Non-preemptive priority queues: exact queue-length analysis

The direct methodological ancestor of this work is Cohen [Coh56], who studied a single server loaded by two independent Poisson streams and introduced the partial-generating-function recursion — the “Cohen trick” — that resolves the lower-priority queue length by transforming only one coordinate of the joint distribution while keeping the other discrete. That device is exactly the one reused, in modern notation, in the partial-PGF analysis of Model- A on the total-count state space (§5.2). The stationary queue-length law of the two-class non-preemptive $M/M/1$ priority queue — the content of Theorem 1 (Model- A) — is in this sense classical, recoverable by several routes; the contribution of the present work is not that result in isolation but the *unified framework* built around it, into which jockeying and abandonment are introduced one mechanism at a time and solved by a common functional-equation method.

The most recent exact closed-form treatment in this line is that of Zuk & Kirszenblat [ZK23], who obtain the joint and marginal queue-length distributions of a non-preemptive $M/M/c$ queue with two priority classes and a common service rate by generating-function methods and a quadratic recurrence; their single-server specialisation ($c = 1$) coincides with Model- A . A complementary strand studies dynamic rather than static priority: Stanford et al. [STZ14] analyse the *accumulating-priority* $M/G/1$ queue, in which each class earns priority linearly in its waiting time so that older customers eventually overtake newer ones, and derive exact waiting-time distributions through a maximum-priority process and Laplace–Stieltjes transforms. Accumulating priority is a continuous-time analogue of the discrete, rate- γ_i jockeying studied here; like the present work it is exact and single-server, but its object is the sojourn time rather than the bivariate queue-length PGF.

Priority queues with abandonment

Abandonment enters queueing theory through the Erlang- A ($M/M/c + M$) model, in which each waiting customer holds an independent exponential patience clock; the modern treatment and its heavy-traffic scaling are due to Garnett et al. [GMR02], and the single-server case $M/M/1 + M$ supplies the busy-period structure on which Model- C_2 rests (§3.3). The analytic signature of

impatience is well documented: per-customer patience clocks make the transition rates state-inhomogeneous, so the generating-function solutions, where they close at all, close in terms of hypergeometric series.

Kapodistria [Kap11] exhibits this pattern at its sharpest in a single-server queue with *synchronised* abandonments — all waiting customers abandon simultaneously, each with probability p , at Poisson opportunity epochs — solving the stationary law exactly through basic (q -)hypergeometric series and recovering, in the no-synchronisation limit $p \rightarrow 0^+$, the $M/M/1 + M$ queue-length PGF as a ratio of Kummer functions: precisely the special-function structure that carries the boundary function of Model- C_2 .

Exact stationary results for priority queues *with* reneging are scarce and almost always multi-server. Jouini & Roubos [JR14] treat an $M/M/s + M$ queue with two non-preemptive classes and per-class exponential patience, deriving Laplace–Stieltjes transforms for conditional and unconditional waiting times together with the probabilities of service and of abandonment — the multi-server analogue of Model- C_2 , but with the waiting time, not the queue-length PGF, as its object. Where exact transforms are unavailable the literature turns to fluid and diffusion approximations or to many-server asymptotics, in which the Erlang-A stationary law and its descendants are expressed through confluent hypergeometric and incomplete-gamma functions [ZM05]; the same special-function structure appears exactly, and at finite scale, in the closed form for Model- C_2 .

The exact multi-server priority results that do exist carry no class-asymmetric reneging — those of Wang et al. [WBSW15] and Zuk & Kirszenblat [ZK23] are abandonment-free — so Model- C_2 appears to be the first exact, single-server, queue-length treatment of *class-asymmetric* abandonment, with impatience acting on the priority class alone. The clinical reading that licenses this asymmetry is long-standing: Zenios [Zen99] models the organ-transplant waiting list as a queue whose reneging is patient mortality, precisely the regime in which losing high-priority demand is the outcome of interest.

Jockeying and priority changes

The closest predecessor on the jockeying side is de Waal [dW92], who studies an $M/G/1$ priority queue in which impatient customers *change priority class* while waiting. Its entry in Table 2 is marked approximate not by choice of method but by necessity: with general service times the governing equations do not close, and the analysis delivers Laplace–Stieltjes-based *approximations* of the waiting-time distributions rather than an exact stationary law. The healthcare incarnation of the same idea is the model of Hu, Chan & Dong [HCD22], who let customers transition between a moderate and an urgent class as their condition deteriorates or improves and characterise a near-optimal $c\mu/\theta$ -type scheduling index in heavy traffic; the system is multi-server and its results are asymptotic, so it too is approximate in the sense of Table 2. Against this backdrop Model- B_2 is the exact, single-server counterpart of the approximate multi-server “priority-change” literature: it retains the deterioration mechanism — a waiting class-1 customer relocating to the class-2 queue at rate γ_1 — but pays for exactness with the single-server, Markovian, one-directional restriction, returning a closed-form boundary function in place of an approximation.

Multi-server extensions and performance design

The last reference bounds the present work from the side of scale. Wang et al. [WBSW15] give the exact queue-length PGF of an $M/M/c$ queue with two *preemptive* priority classes and class-dependent service rates, a useful contrast to the non-preemptive, common-rate discipline used throughout this thesis. That the result carries no departure mechanism is symptomatic of the multi-server literature as a whole: once $c > 1$ is combined with jockeying or abandonment, the analysis moves to approximations and asymptotics, confirming that exact, mechanism-rich analysis is essentially a single-server phenomenon.

Methodological lineage

The functional-equation technique used in Sections 7 and 8 belongs to the kernel-method and boundary-value tradition for two-dimensional Markovian queues. The systematic theory — reducing a bivariate balance equation to the curve on which its kernel vanishes, and from there to a boundary-value problem for the unknown boundary functions — is developed by Fayolle, Iasnogorodski & Malyshev [FIM99] for random walks in the quarter-plane. That programme still delivers explicit two-class results: Devos, Walraevens, Fiems & Bruneel [DWFB20] obtain the closed-form joint queue-length distribution of a discrete-time two-class queue with randomly alternating service

by analysing the kernel of just such a functional equation and pinning its two unknown boundary functions through analytic continuation.

The pinning arguments of this thesis are the elementary, one-dimensional residue of that programme. In Model- B_2 the boundary function is fixed by demanding *analyticity* of the PGF at the interior point where the kernel root meets the diagonal (§7), and in Model- C_2 by demanding mere *boundedness* at the unit-disc boundary (§8); both select the analytic solution of the reduced equation, exactly the role the boundary conditions play in the general theory.

Two narrower tools recur. The confluent hypergeometric (Kummer) function that carries the Model- C_2 boundary function is the same special function through which the Erlang-A and reneging-queue stationary laws are expressed [ZM05, Kap11, DLM]; and the selection of the decaying solution of the three-term recurrence behind the partial-PGF analysis is governed by Pincherle’s theorem, for which we cite the recurrence-relation treatment of Gautschi [Gau67] alongside the DLMF [DLM].

Position of this thesis

Table 2: Literature overview: models and mechanisms. ✓ = present, – = absent, ≈ = approximate only.

Reference	Servers	Non-preem.	Jockeying	Aband.	Queue-len.	Exact
Cohen (1956)	1	✓	–	–	✓	✓
de Waal (1992)	1	✓	✓	–	–	≈
Jouini & Roubos (2014)	c	✓	–	✓	–	✓
Wang et al. (2015)	c	–	–	–	✓	✓
Stanford et al. (2014)	1	✓	✓	–	–	✓
Zuk & Kirszenblat (2023)	c	✓	–	–	✓	✓
Hu, Chan & Dong (2022)	c	✓	✓	✓	–	≈
This thesis (Model-A)	1	✓	–	–	✓	✓
This thesis (Model-B_2)	1	✓	✓	–	✓	✓
This thesis (Model-C_2)	1	✓	–	✓	✓	✓

Table 2 places this work against its closest neighbours, and reading it by row makes the gap precise. The Cohen (1956) row is the single-server, exact, non-preemptive queue-length baseline with neither extra mechanism — the ancestor of Model-A. de Waal (1992) adds jockeying (priority change) at a single server, but only approximately and for waiting times rather than queue lengths. Jouini & Roubos (2014) add abandonment exactly, but at c servers and again for waiting times. Wang et al. (2015) is exact and queue-length-focused, yet multi-server, preemptive, and mechanism-free. Stanford et al. (2014) is single-server and exact with a dynamic-priority (jockeying-like) mechanism, but its object is the sojourn time. Zuk & Kirszenblat (2023) is the exact multi-server queue-length result whose $c = 1$ case is Model-A, with neither jockeying nor abandonment. Hu, Chan & Dong (2022) carries both mechanisms — jockeying-as-deterioration and abandonment — but multi-server and only approximately.

No row above combines, at a single server, the exact bivariate queue-length PGF with either jockeying or abandonment — which is precisely the column pattern of the three “This thesis” rows. The gap can therefore be stated in one sentence: prior to this work there is no exact, single-server, queue-length analysis that places non-preemptive priority together with jockeying or abandonment inside a single PGF framework.

The non-claims must be stated with equal precision. The framework is unified but not complete: the full model (Model-X, Remark 1), in which jockeying and abandonment act together, and the bidirectional-jockeying model (Model-B, §6) are left open, each reducing to a Volterra-type integral equation that this thesis characterises but does not solve in closed form. The exact contributions are confined to the three “This thesis” rows of Table 2 — the baseline Model-A, the one-way jockeying Model- B_2 , and the class-1 abandonment Model- C_2 , together with their head-of-line simplifications. The claim is exactness and unification at a single server with one mechanism active at a time; it is explicitly not a solution of the combined or bidirectional problems, which remain, as in the multi-server literature, beyond exact reach.

3 Preliminaries

This section collects the background material used throughout the derivations, so that the later arguments can stay focused on their primary objectives.

The central object throughout is the (P)PGF of the model in question, together with the system’s idle probability for some corollaries.

3.1 $M/M/1$ queue

The canonical $M/M/1$ queue is the foundation of queueing theory and is treated in many textbooks; we follow [AR02]. Customers arrive according to a Poisson process —i.e. the inter-arrival times are exponentially distributed with parameter λ —, service times are exponentially distributed with parameter μ , and a single server operates in first-come-first-serve (FCFS) —also known as first-in-first-out (FIFO)— order. The system is stable if and only if the load $\rho = \lambda/\mu < 1$.

Let π_n be the limiting probability of finding n customers in the system. Under stability, the detailed balance equations $\lambda\pi_n = \mu\pi_{n+1}$ ($n \geq 0$) together with the normalisation $\sum_{n=0}^{\infty} \pi_n = 1$ give us the geometric steady-state distribution of the number of customers in the system:

$$\pi_n = (1 - \rho) \rho^n, \quad n \geq 0.$$

By the PASTA property (§3.2), this also corresponds to the distribution seen by arriving customers. Let L be the number of customers in the system, including any in service. Its expectation is

$$\mathbb{E}[L] = \frac{\rho}{1 - \rho}.$$

From this, the expected sojourn time W —the time between a customer’s arrival and their departure after service—follows by Little’s Law ($\mathbb{E}[L] = \lambda\mathbb{E}[W]$):

$$\mathbb{E}[W] = \frac{\mathbb{E}[L]}{\lambda} = \frac{1/\mu}{1 - \rho}.$$

A less common but highly relevant concept is the Busy Period (BP): the time between a customer arriving to an idle system and the departure that next leaves it empty. Its Laplace–Stieltjes Transform (LST) and expectation are known:

$$\widetilde{BP}(s) = \frac{(\mu + \lambda + s) - \sqrt{(\mu + \lambda + s)^2 - 4\mu\lambda}}{2\lambda}, \quad (3.1)$$

$$\mathbb{E}[BP] = \frac{1}{\mu(1 - \rho)}. \quad (3.2)$$

3.2 Non-preemptive priority and the PASTA property

Throughout this thesis the two customer classes are served under a *non-preemptive* (head-of-line) priority discipline: whenever the server becomes free it admits a waiting class-1 customer if any are present, and a class-2 customer only otherwise, but a customer already in service is *never* interrupted—a class-1 arrival that finds a class-2 customer in service waits until that service completes. Priority therefore governs only the order in which the queues are drained, not the service in progress. Because the service rate μ is common to both classes, the *memorylessness* of the exponential service time makes the residual service of the customer in progress $\text{Exp}(\mu)$ regardless of its class; hence the class of the in-service customer does not influence the future evolution of the queue lengths, which is why a single in-service customer of either class suffices to mark the server busy.

We also invoke repeatedly the *PASTA* property (*Poisson Arrivals See Time Averages*) [AR02]: if a system is fed by a Poisson arrival process that satisfies the *lack-of-anticipation assumption* (at each instant the future increments of the arrival process are independent of the current system state), then the long-run fraction of arrivals that find the system in a given state equals the stationary probability of that state, so the distribution seen by an arriving customer coincides with the time-stationary one. Both hypotheses hold here: the exogenous class- i arrivals are independent Poisson processes at rates λ_i , and they are independent of the system’s internal exponential service, jockeying and abandonment clocks, so they cannot anticipate the future. By Poisson *superposition* the merged arrival stream is itself Poisson at rate $\lambda = \lambda_1 + \lambda_2$, so PASTA applies to the aggregate state; and since each class stream is an independent Poisson process—equivalently, an i.i.d. λ_i/λ *thinning* of the merged stream—PASTA applies to each class separately.

3.3 $M/M/1+M$ queue (Erlang-A)

Another model relevant to this thesis is the $M/M/1+M$ queue, also known as the *Erlang-A* model. It extends the $M/M/1$ queue —Poisson arrivals at rate λ , exponential service at rate μ — with customer impatience: each *waiting* customer (the one in service is not impatient) independently abandons after an exponentially distributed patience time with rate θ . Crucially, unlike the $M/M/1$ it does not require $\rho = \lambda/\mu < 1$ for stability: it is positive recurrent for all $\lambda, \mu, \theta > 0$.

The steady-state distribution follows from its balance equations, where n denotes the number of customers in the *system*:

$$\lambda \pi_n = (\mu + n\theta) \pi_{n+1}, \quad n \geq 0.$$

Iterating yields

$$\pi_n = \pi_0 \frac{(\lambda/\theta)^n}{(\mu/\theta)^{(n)}} \quad n \geq 0; \quad \text{where} \quad a^{(n)} = a(a+1) \cdots (a+n-1)$$

Normalisation gives the non-trivial idle probability $\pi_0 = [{}_1F_1(1; \mu/\theta; \lambda/\theta)]^{-1}$ (see §3.5 for the Kummer function ${}_1F_1$), and summing the geometric-like terms shows that the queue-length PGF is correspondingly a *ratio* of Kummer functions, $\Pi(z) = {}_1F_1(1; \mu/\theta; (\lambda/\theta)z) / {}_1F_1(1; \mu/\theta; \lambda/\theta)$. The same closed form is derived by Kapodistria [Kap11, Thm. 3] as the no-synchronisation limit of a queue with synchronised abandonments, there in the sojourn-time-impatience convention (the customer in service also abandons), which amounts to replacing μ by $\mu + \theta$.

The *Busy Period* of this $M/M/1+M$ is far more complex than its $M/M/1$ counterpart. We record it for the class-1 subsystem used in Model- C_2 , with arrival rate λ_1 , service rate μ and abandonment rate θ_1 , so that the relevant idle probability is $\pi_0 = [{}_1F_1(1; \mu/\theta_1; \lambda_1/\theta_1)]^{-1}$. Writing B_C for this busy period, its LST and expectation are

$$\tilde{B}_C(s) = \frac{\mu}{\lambda_1 + \mu + s - \frac{\lambda_1(\mu + \theta_1)}{\lambda_1 + \mu + \theta_1 + s - \frac{\lambda_1(\mu + 2\theta_1)}{\lambda_1 + \mu + 2\theta_1 + s - \cdots}}}, \quad (3.3)$$

$$\mathbb{E}[B_C] = \frac{1}{\lambda_1} \left({}_1F_1(1; \mu/\theta_1; \lambda_1/\theta_1) - 1 \right) \quad (3.4)$$

The continued fraction (3.3) is generated by a first-passage recursion. Let $\phi_n(s)$ denote the LST of the first-passage (sub-busy-period) time from n customers in the system down to $n-1$. Conditioning on the first transition out of state n —an arrival (rate λ_1) to state $n+1$, or a departure by service completion or abandonment (combined rate $\mu + (n-1)\theta_1$: one customer in service plus $n-1$ waiting, each impatient at rate θ_1) to state $n-1$ —the strong Markov property gives

$$\phi_n(s) = \frac{\mu + (n-1)\theta_1}{\lambda_1 + \mu + (n-1)\theta_1 + s - \lambda_1 \phi_{n+1}(s)}, \quad n \geq 1, \quad \tilde{B}_C(s) = \phi_1(s). \quad (3.5)$$

Unrolling (3.5) from $n=1$ reproduces the continued fraction (3.3) term by term. The continued fraction also sums in closed form, and the argument is short enough to give in full. Fix $s > 0$ and scale the rates by θ_1 , setting

$$a = \frac{\mu}{\theta_1}, \quad b = \frac{s}{\theta_1}, \quad c = \frac{\lambda_1}{\theta_1},$$

so that (3.5) reads $\phi_n(s) = (a+n-1)/((a+n-1) + b + c - c\phi_{n+1}(s))$. Consider the integrals

$$I(p) = \int_0^1 t^p (1-t)^{b-1} e^{-ct} dt, \quad p > -1,$$

which by the Euler representation (3.10) are Beta-function multiples of Kummer values, $I(p) = B(p+1, b) {}_1F_1(p+1; p+1+b; -c)$. Integrating the exact derivative $\frac{d}{dt} [t^p (1-t)^b e^{-ct}]$ over $[0, 1]$ — the boundary terms vanish for $p, b > 0$ — and splitting $(1-t)^b = (1-t)^{b-1} - t(1-t)^{b-1}$ in the two terms where it appears yields the three-term recurrence

$$(p+b+c)I(p) = pI(p-1) + cI(p+1), \quad (3.6)$$

a relation of contiguous type for ${}_1F_1$ (obtainable equivalently by combining the contiguous relations of [DLM, §13.3]). Dividing (3.6) at $p = a+n-1$ by $I(a+n-2)$ shows that the ratios $R_n :=$

$I(a+n-1)/I(a+n-2)$ satisfy exactly the recursion of ϕ_n . Two further facts pin the identification. First, by Laplace's method at the endpoint $t = 1$, $I(p) \sim e^{-c} \Gamma(b) p^{-b}$ as $p \rightarrow \infty$: the solution $\{I(p)\}$ decays only polynomially, whereas the dominant solution of (3.6) grows superexponentially (its consecutive ratio grows like p/c), so $\{I(p)\}$ is the *minimal (recessive)* solution. Second, by *Pincherle's theorem* the convergent continued fraction equals precisely the ratio built from the minimal solution. Hence

$$\tilde{B}_C(s) = \phi_1(s) = \frac{I(a)}{I(a-1)} = \frac{B(a+1, b) {}_1F_1(a+1; a+1+b; -c)}{B(a, b) {}_1F_1(a; a+b; -c)}, \quad (3.7)$$

a contiguous ratio of the form (3.11). This identity is used, with $s = \lambda_2(1-y)$, in the boundedness route for Model- C_2 (§8, Equation (8.15)).

3.4 Pollaczek–Khinchine formula

The $M/G/1$ queue [AR02] is a single-server queue with Poisson arrivals at rate λ and generally distributed i.i.d. service times B with LST $\tilde{B}(s) = \mathbb{E}[e^{-sB}]$. For stability purposes, the system load $\rho_B = \lambda \mathbb{E}[B] < 1$ is required (written ρ_B to keep the symbol ρ for the two-class load of this thesis). Under these two conditions —Poisson arrivals and stability— the Pollaczek–Khinchine (PK) formula gives the PGF of the stationary number of customers in the system:

$$L(z) = \frac{(1 - \rho_B)(1 - z) \tilde{B}(\lambda(1 - z))}{\tilde{B}(\lambda(1 - z)) - z}. \quad (3.8)$$

3.5 Confluent hypergeometric function

The *Kummer confluent hypergeometric function of the first kind* appears in several derivations in this thesis; it reads as follows:

$${}_1F_1(a; b; z) = \sum_{n=0}^{\infty} \frac{a^{(n)}}{b^{(n)} n!} z^n, \quad a^{(n)} = \prod_{k=0}^{n-1} (a + k), \quad a^{(0)} = 1, \quad (3.9)$$

where b is not 0 or a negative integer —i.e. $b \notin \{0, -1, -2, \dots\}$ — and the series converges for all $z \in \mathbb{C}$. The function ${}_1F_1(a; b; z)$ is the unique solution of *Kummer's equation* [DLM]:

$$z u'' + (b - z) u' - a u = 0,$$

normalised by ${}_1F_1(a; b; 0) = 1$.

For this thesis, two identities related to it are used. First, the *Euler integral representation*, for $a, b > 0$ we have:

$${}_1F_1(a; a+b; z) = \frac{1}{B(a, b)} \int_0^1 t^{a-1} (1-t)^{b-1} e^{zt} dt, \quad (3.10)$$

where B is the *Beta function* $B(a, b) = \Gamma(a)\Gamma(b)/\Gamma(a+b)$, and Γ is the *Gamma function*. Secondly, *contiguous ratios*: ratios of two ${}_1F_1$ values whose parameters differ by integer shifts. Two patterns arise in this thesis. The *simultaneous shift*

$$\frac{{}_1F_1(a+1; a+1+b; z)}{{}_1F_1(a; a+b; z)}, \quad (3.11)$$

in which both parameters move by one so that their difference b is preserved, is by the Euler representation (3.10) a ratio of two weighted Beta integrals differing only in the exponent of t ; it appears as the closed-form value of the Erlang-A busy-period continued fraction (§3.3, Equation (3.7)) and hence in the boundary function of Model- C_2 . The *first-parameter shift* ${}_1F_1(a+1; b; z)/{}_1F_1(a; b; z)$, with the second parameter fixed, appears instead in the boundary function of Model- B_2 (§7).

3.6 First-order linear ODEs and PDEs

Since jockeying and abandonment add state-dependent rates to the balance equations, derivative terms appear in the derivations, requiring us to solve first-order linear ODEs and PDEs for the PGF. The three methods below are invoked at specific steps of the analysis.

3.6.1 Integrating factor

This method is standard to solve —for each fixed y — a *first-order linear ODE in x* . Such an equation has standard form

$$\frac{dP}{dx} + p(x; y) P = q(x; y). \quad (3.12)$$

The *integrating factor* $\mu_I(x; y) = \exp(\int_0^x p(t; y) dt)$ converts (3.12) to $\frac{d}{dx}[P \cdot \mu_I] = q \cdot \mu_I$. Integrating from 0 to x gives the general solution

$$P(x, y) = \frac{1}{\mu_I(x; y)} \left[C_0(y) + \int_0^x q(t; y) \mu_I(t; y) dt \right], \quad (3.13)$$

where $C_0(y)$ is an arbitrary function of y , determined by a boundary condition.

3.6.2 Variation of parameters

Variation of parameters provides an alternative derivation of (3.13) that makes the role of the free function explicit. The homogeneous equation $\frac{dP}{dx} + p(x; y) P = 0$ has general solution $P_h(x; y) = C/\mu_I(x; y)$ for an arbitrary constant C . To solve the full equation, one *varies* that constant, setting $P(x, y) = C(x; y)/\mu_I(x; y)$ and substituting into (3.12). Since the homogeneous part cancels, one obtains the following.

$$\frac{dC}{dx}(x; y) = q(x; y) \mu_I(x; y),$$

and integrating recovers (3.13) with $C_0(y)$.

3.6.3 Method of Characteristics

This method is standard for solving a *first-order linear PDE in both x and y* simultaneously. It reduces the PDE to a family of ODEs along the so-called *characteristic curves* in the (x, y) -plane; along each curve, the PDE collapses to a single-variable ODE.

A PDE of the form

$$a \frac{\partial P}{\partial x} + b \frac{\partial P}{\partial y} = cP + d \quad (3.14)$$

is solved by tracing the characteristic curve $(x(t), y(t))$ satisfying

$$\frac{dx}{dt} = a, \quad \frac{dy}{dt} = b. \quad (3.15)$$

By the chain rule, $\frac{d}{dt}P(x(t), y(t)) = a P_x + b P_y$, so (3.14) reduces to

$$\frac{dP}{dt} = cP + d, \quad (3.16)$$

a first-order linear ODE in t , solved by the integrating factor (§3.6.1). When a and b are constants, the characteristics are straight lines and the quantity

$$\xi := bx - ay$$

is preserved along each curve—it is the *first integral* that labels each characteristic. The integration constant of (3.16) therefore varies between characteristics: it is an arbitrary function $F(\xi)$. Re-expressing the solution in the original variables (x, y) via (3.15) gives the general solution of (3.14). In the PGF setting, F is identified with a boundary generating function, and analyticity of $P(x, y)$ for $|x|, |y| \leq 1$ uniquely determines it.

4 Model description: Model- X

This section defines the general two-class system with both jockeying (γ_i) and abandonment (θ_i); we call this full model *Model- X* . The simpler variants are recovered as special cases by zeroing the appropriate parameters, as summarised in Table 3:

Model	Jockeying	Abandonment	Parameter restriction
X	both ($1 \leftrightarrow 2$)	both ($1, 2$)	$\gamma_i \geq 0, \theta_i \geq 0$ (general)
A	none	none	$\gamma_1 = \gamma_2 = 0, \theta_1 = \theta_2 = 0$
B	both ($1 \leftrightarrow 2$)	none	$\gamma_i \geq 0, \theta_1 = \theta_2 = 0$
B_2	$1 \rightarrow 2$ only	none	$\gamma_1 \geq 0, \gamma_2 = 0, \theta_1 = \theta_2 = 0$
C_2	none	class 1 only	$\theta_1 \geq 0, \theta_2 = 0, \gamma_1 = \gamma_2 = 0$
B_2^H	$1 \rightarrow 2$, head-of-line	none	$\gamma_1 \mathbf{1}_{\{n_1 \geq 1\}} > 0, \gamma_2 = \theta_1 = \theta_2 = 0$
C_2^H	none	class 1, head-of-line	$\theta_1 \mathbf{1}_{\{n_1 \geq 1\}} > 0, \gamma_1 = \gamma_2 = \theta_2 = 0$

Table 3: Nested family of models. Model- X is the parent; Models A , B , B_2 and C_2 are obtained by restricting the jockeying directions and the abandoning classes. The two head-of-line variants B_2^H and C_2^H are listed separately, as their zeroing mechanism involves $\mathbf{1}_{\{n_1 \geq 1\}}$ rather than being proportional to the queue length.[†]

Figure 1 gives a visual overview of Model- X and each variant we consider. It omits the two head-of-line variants, which are less conventional and for which such a diagram would add little:




Figure 1: Physical queue layout for Model- X (left) and its four specialisations (right, 2×2 grid). Queue 1 (top buffer) has non-preemptive priority over Queue 2. Solid arrows are for Poisson arrivals (λ_i) and exponential service (μ); dashed arrows are for jockeying (γ_i); and dotted arrows are for abandonments (θ_i)

The model in this section is an $M/M/1$ queue with two priority classes. Customers are served First-Come, First-Served (FCFS) within their own class, and class-1 (the priority class) has non-preemptive priority over class-2: a service in progress is never interrupted, so a class-2 customer in service is not displaced by an arriving class-1 customer. Beyond these basic dynamics, Model- X adds two queue-leaving mechanisms. First, *jockeying*: a waiting customer leaves its own queue to *join the other queue*, at per-customer rate γ_1 for direction $1 \rightarrow 2$ and γ_2 for $2 \rightarrow 1$, so the total number of customers is preserved. Second, *abandonment* (reneging): a waiting customer leaves the *system* at per-customer rate θ_i , decreasing the total by one. Concretely, each waiting class- i customer (not the one in service) holds an independent $\text{Exp}(\theta_i)$ patience clock and abandons when it expires, so the aggregate class- i abandonment rate in a state with n_i waiting customers is the linear term $\theta_i n_i$ in the balance equations. This is precisely the linear-reneging mechanism of the $M/M/1 + M$ queue from §3.3.

In Kendall's notation this is an $M/M/1$ -type model: Poisson arrivals at rate λ_i for class- i , exponentially distributed service times at rate μ , and a single server. To define the stochastic processes that drive this system, we introduce the following random variables:

- $\mathcal{B}$: binary indicator for the server being idle or busy, $\mathcal{B} \in \{0, 1\}$;
- N_1, N_2 : number of customers of class 1 and 2, respectively, in their respective queues;
- N : total number of customers in the queues.

Note that N_1 , N_2 and N count customers in the *queues*, not the *system*. Since the service rate is common to both classes, we write the idle state as (0) and all other states as (n_1, n_2) , so $(0, 0)$ means someone is in service and the queues are empty. Alternatively, an auxiliary state space $\tilde{S}$ tracks only n_2 and the total $n = n_1 + n_2$, subject to $n \geq n_2 \geq 0$. Formally:

$$\begin{aligned} S &:= \{(0)\} \cup \{(n_1, n_2) \mid n_1, n_2 \geq 0\}, \\ \tilde{S} &:= \{(0)\} \cup \{(n_2, n) \mid n \geq n_2 \geq 0\}. \end{aligned} \tag{4.1}$$

We define the long-run probabilities of finding the system in each state:

- π_0 for the idle state (0) , it belongs to both spaces S and $\tilde{S}$;
- $\pi(n_1, n_2)$ for states $(n_1, n_2) \in S$, where $n_1, n_2 \geq 0$;
- $\tilde{\pi}(n_2, n)$ for states $(n_2, n) \in \tilde{S}$, where $n \geq n_2 \geq 0$.

Arrivals are Poisson for each class and service times exponential (common across classes), with the following notation:

- λ_1, λ_2 : arrival rates of classes 1 and 2, respectively;
- μ : service rate of a customer, regardless of class;
- $\rho_i = \frac{\lambda_i}{\mu}$: system load of class $i \in \{1, 2\}$;
- γ_1, γ_2 : per-customer jockeying rates from Class 1 $\rightarrow$ 2 and 2 $\rightarrow$ 1, respectively; jockeying transfers a customer between queues, preserving the total number of customers;
- θ_1, θ_2 : per-customer abandonment (reneging) rates from classes 1 and 2, respectively; abandonment removes a customer from the system, reducing the total n by one.

For the aggregate system parameters, we define:

- $\lambda = \lambda_1 + \lambda_2$: total arrival rate;
- $\rho = \frac{\lambda}{\mu}$: total system load.

Stability. Jockeying does not affect stability (customers are merely relocated). Model- X has abandonments from both classes, so the departure rate $\mu + \theta_1 n_1 + \theta_2 n_2$ grows unbounded, ensuring positive recurrence for all loads. In contrast, models without abandonments require $\rho < 1$ (they behave like M/M/1). Mixed cases, such as Model- C_2 , are discussed separately (§8).

Computation of π_0 and $\pi(0, 0)$. Balance between state (0) and $(0, 0)$ gives $\lambda\pi_0 = \mu, \pi(0, 0)$. Hence,

$$\pi(0, 0) = \rho, \pi_0. \tag{4.2}$$

This holds regardless of jockeying and abandonments (neither mechanism acts when there are no waiting customers). However, π_0 itself depends on the abandonment model. For models *without* abandonment, $\pi_0 = 1 - \rho$ (M/M/1 reduction). For models *with* abandonment, the departure rate depends on the split (n_1, n_2) , so π_0 must be determined from the bivariate PGF (§4.2).

4.1 Balance equations

This subsection presents the state-transition diagrams of Model- X on both state spaces S and $\tilde{S}$, and derives the corresponding sets of balance equations.

4.1.1 Balance equations for state space S

The state-transition diagram of Model- X on S is given in Figure 2.




Figure 2: State transition diagram on state space S for the model with jockeying ($\gamma_i n_i$, dashed) and abandonments ($\theta_i n_i$, dotted).

$$\begin{aligned}
& [\lambda_1 + \lambda_2 + \mu + (\gamma_1 + \theta_1)n_1 + (\gamma_2 + \theta_2)n_2] \pi(n_1, n_2) \\
&= \mu \pi(n_1 + 1, n_2) \\
&\quad + \lambda_1 \pi(n_1 - 1, n_2) \\
&\quad + \lambda_2 \pi(n_1, n_2 - 1) \\
&\quad + \gamma_1(n_1 + 1) \pi(n_1 + 1, n_2 - 1) \\
&\quad + \gamma_2(n_2 + 1) \pi(n_1 - 1, n_2 + 1) \\
&\quad + \theta_1(n_1 + 1) \pi(n_1 + 1, n_2) \\
&\quad + \theta_2(n_2 + 1) \pi(n_1, n_2 + 1)
\end{aligned}
\tag{4.3}$$

$$\begin{aligned}
& [\lambda_1 + \lambda_2 + \mu + (\gamma_2 + \theta_2)n_2] \pi(0, n_2) \\
&= \lambda_2 \pi(0, n_2 - 1) \\
&\quad + \mu \pi(1, n_2) \\
&\quad + \mu \pi(0, n_2 + 1) \\
&\quad + \gamma_1 \pi(1, n_2 - 1) \\
&\quad + \theta_1 \pi(1, n_2) \\
&\quad + \theta_2(n_2 + 1) \pi(0, n_2 + 1)
\end{aligned}
\tag{4.4}$$

$$\begin{aligned}
& [\lambda_1 + \lambda_2 + \mu + (\gamma_1 + \theta_1)n_1] \pi(n_1, 0) \\
&= \lambda_1 \pi(n_1 - 1, 0) \\
&\quad + \mu \pi(n_1 + 1, 0) \\
&\quad + \gamma_2 \pi(n_1 - 1, 1) \\
&\quad + \theta_1(n_1 + 1) \pi(n_1 + 1, 0) \\
&\quad + \theta_2 \pi(n_1, 1)
\end{aligned}
\tag{4.5}$$

$$[\lambda_1 + \lambda_2] \pi(0, 0) = (\mu + \theta_1) \pi(1, 0) + (\mu + \theta_2) \pi(0, 1)
\tag{4.6}$$

$$(\lambda_1 + \lambda_2) \pi_0 = \mu \pi(0, 0)
\tag{4.7}$$

4.1.2 Balance equations for state space $\tilde{S}$

Recall the state space $\tilde{S}$ defined in (4.1).

For readability we denote its limiting probabilities by $\tilde{\pi}(n_2, n)$, for $(n_2, n) \in \tilde{S}$. Figure 3 shows the corresponding state transition diagram.




Figure 3: State transition diagram on state space $\tilde{S}$.

The balance equations, after combining the equations where $\tilde{\pi}_0$ shows up, are:

$$\begin{aligned}
 & [\lambda_1 + \lambda_2 + \mu + (\gamma_1 + \theta_1)(n - n_2) + (\gamma_2 + \theta_2)n_2] \tilde{\pi}(n_2, n) \\
 &= \lambda_1 \tilde{\pi}(n_2, n - 1) \\
 &+ \lambda_2 \tilde{\pi}(n_2 - 1, n - 1) \\
 &+ \mu \tilde{\pi}(n_2, n + 1) \\
 &+ \gamma_1(n - n_2 + 1) \tilde{\pi}(n_2 - 1, n) \\
 &+ \gamma_2(n_2 + 1) \tilde{\pi}(n_2 + 1, n) \\
 &+ \theta_1(n + 1 - n_2) \tilde{\pi}(n_2, n + 1) \\
 &+ \theta_2(n_2 + 1) \tilde{\pi}(n_2 + 1, n + 1)
 \end{aligned}
 \quad \text{for } 1 \leq n_2 \leq n - 1, n \geq 2 \quad (4.8)$$

$$\begin{aligned}
 & [\lambda_1 + \lambda_2 + \mu + (\gamma_2 + \theta_2)n] \tilde{\pi}(n, n) \\
 &= \lambda_2 \tilde{\pi}(n - 1, n - 1) \\
 &+ \mu \tilde{\pi}(n + 1, n + 1) \\
 &+ \mu \tilde{\pi}(n, n + 1) \\
 &+ \gamma_1 \tilde{\pi}(n - 1, n) \\
 &+ \theta_1 \tilde{\pi}(n, n + 1) \\
 &+ \theta_2(n + 1) \tilde{\pi}(n + 1, n + 1)
 \end{aligned}
 \quad \text{for } n_2 = n, n \geq 1 \quad (4.9)$$

$$\begin{aligned}
 & [\lambda_1 + \lambda_2 + \mu + (\gamma_1 + \theta_1)n] \tilde{\pi}(0, n) \\
 &= \lambda_1 \tilde{\pi}(0, n - 1) \\
 &+ \mu \tilde{\pi}(0, n + 1) \\
 &+ \gamma_2 \tilde{\pi}(1, n) \\
 &+ \theta_1(n + 1) \tilde{\pi}(0, n + 1) \\
 &+ \theta_2 \tilde{\pi}(1, n + 1)
 \end{aligned}
 \quad \text{for } n_2 = 0, n \geq 1 \quad (4.10)$$

$$\begin{aligned} & [\lambda_1 + \lambda_2] \tilde{\pi}(0, 0) \\ &= (\mu + \theta_1) \tilde{\pi}(0, 1) + (\mu + \theta_2) \tilde{\pi}(1, 1) \end{aligned} \quad (4.11)$$

$$(\lambda_1 + \lambda_2) \tilde{\pi}_0 = \mu \tilde{\pi}(0, 0) \quad (4.12)$$

4.2 Derivation of equations for (Partial) Probability Generating functions

To study the limiting behaviour of this two-class system, we define the bivariate Probability Generating Function (PGF) of the joint queue lengths N_1 and N_2 :

$$P(x, y) = \mathbb{E}[x^{N_1} y^{N_2}] = \sum_{n_1=0}^{\infty} \sum_{n_2=0}^{\infty} \pi(n_1, n_2) x^{n_1} y^{n_2}.$$

The derivation also uses the boundary terms, where at least one queue is empty:

$$\begin{aligned} P_x(x) &= P(x, 0) = \sum_{n_1=0}^{\infty} \pi(n_1, 0) x^{n_1} \\ P_y(y) &= P(0, y) = \sum_{n_2=0}^{\infty} \pi(0, n_2) y^{n_2}. \end{aligned}$$

Note that both terms include the busy-but-empty state $\pi(0, 0)$, whereas the fully idle state is treated separately.

A further term arises in the balance equations, describing the boundary dynamics when the class-1 queue holds exactly one waiting customer:

$$P_1(y) = \sum_{n_2=1}^{\infty} \pi(1, n_2) y^{n_2}.$$

This is an intermediate term; as the derivation proceeds, it will be expressed in terms of the primary boundary functions $P_x(x)$ and $P_y(y)$.

In the same fashion, we introduce the **Partial Probability Generating Function (PPGF)**. This technique, notably used by Cohen [Coh56], transforms only a subset of the random variables while keeping the discrete indices of the rest. For a priority queue it is particularly advantageous: it lets us exploit the well-known properties of the priority class (often a standard $M/M/1$ system) and concentrate the analytical effort on the non-trivial PGF of the non-priority class.

Cohen developed this framework for multi-server systems; we apply it here to the single-server case. To avoid confusion from mid-20th-century notation, we give a modern formulation consistent with the variables defined above.

We define the PPGF with respect to N_2 , conditioned on the first queue being in state n_1 :

$$\tilde{P}(n_1, y) = \mathbb{E}[y^{N_2} \mathbb{1}_{\{N_1=n_1\}}] = \sum_{n_2=0}^{\infty} \pi(n_1, n_2) y^{n_2}. \quad (4.13)$$

Similarly, the transform with respect to N_1 , keeping the second queue in discrete state n_2 , is:

$$\tilde{P}(x, n_2) = \mathbb{E}[x^{N_1} \mathbb{1}_{\{N_2=n_2\}}] = \sum_{n_1=0}^{\infty} \pi(n_1, n_2) x^{n_1}.$$

These partial transforms relate to the joint PGF by:

$$P(x, y) = \sum_{n_1=0}^{\infty} \tilde{P}(n_1, y) x^{n_1} = \sum_{n_2=0}^{\infty} \tilde{P}(x, n_2) y^{n_2}.$$

4.2.1 Equation for PGF in state space S

Lemma 1 (Fundamental equation for the PGF on S). *Consider the two-class non-preemptive priority queue on the state space $S = \{(n_1, n_2) : n_1, n_2 \geq 0\}$, with Poisson arrival rates λ_1, λ_2 , exponential service rate μ , jockeying rates γ_1, γ_2 and abandonment rates θ_1, θ_2 , where n_1 and n_2*

count the class-1 and class-2 customers waiting in the queues. Assume the associated CTMC is positive recurrent, so that the stationary distribution $\{\pi(n_1, n_2)\}$ exists and the joint PGF

$$P(x, y) = \sum_{n_1=0}^{\infty} \sum_{n_2=0}^{\infty} \pi(n_1, n_2) x^{n_1} y^{n_2}$$

is analytic on the closed unit polydisc $|x| \leq 1$, $|y| \leq 1$. Then $P(x, y)$, together with the boundary function $P_y(y) = P(0, y) = \sum_{n_2=0}^{\infty} \pi(0, n_2) y^{n_2}$, satisfies the first-order linear partial differential equation

$$\begin{aligned} & [(\lambda_1 + \lambda_2 + \mu)xy - \mu y - \lambda_1 x^2 y - \lambda_2 x y^2] P(x, y) \\ & + xy [\gamma_1(x - y) + \theta_1(x - 1)] \frac{\partial}{\partial x} P(x, y) \\ & + xy [\gamma_2(y - x) + \theta_2(y - 1)] \frac{\partial}{\partial y} P(x, y) \\ & = \mu(x - y)P_y(y) - \mu x(1 - y)\pi(0, 0). \end{aligned} \tag{4.14}$$

Reading the fundamental equation. Equation (4.14) repays a moment's mechanical reading, since its structure dictates the entire analysis that follows. The bracketed algebraic coefficient multiplying $P(x, y)$ is exactly the kernel of the baseline Model-A: the flux balance between the arrival contributions $\lambda_1 x^2 y$ and $\lambda_2 x y^2$, the service contribution μy , and the total outflow $(\lambda_1 + \lambda_2 + \mu)xy$.

The two convection coefficients carry the added mechanisms. In the x -convection coefficient $xy [\gamma_1(x - y) + \theta_1(x - 1)]$, the jockeying factor $(x - y)$ vanishes on the diagonal $x = y$ — the analytic signature that jockeying merely relocates a customer and so conserves the total count — while the abandonment factor $(x - 1)$ vanishes at $x = 1$, the signature that abandonment respects normalisation ($P(1, 1) = 1$) yet removes a customer and thereby destroys conservation. The y -convection coefficient $xy [\gamma_2(y - x) + \theta_2(y - 1)]$ states the symmetric facts for the class-2 mechanisms, with $(y - x)$ vanishing on the diagonal and $(y - 1)$ at $y = 1$.

The right-hand side carries the two boundary unknowns that any solution must pin down: the marginal $P_y(y) = P(0, y)$, the generating function of the states with an empty priority queue, and the scalar $\pi(0, 0)$, the probability that the server is busy with both queues empty.

The position of these vanishing factors fixes, model by model, where the operative singularity sits and hence which technique applies. With both convection terms absent (Model-A, §5) the equation is algebraic and $P_y(y)$ is recovered from the kernel root $x^*(y)$ inside the unit disc. With one-way jockeying alone (Model- B_2 , §7) the factor $(x - y)$ places the singularity at the interior point $x = y$ with a negative exponent, and $P_y(y)$ is pinned by analyticity there. With class-1 abandonment alone (Model- C_2 , §8) the factor $(x - 1)$ places it on the boundary $x = 1$ with a positive exponent, and $P_y(y)$ is pinned by mere boundedness. Restricting a mechanism to the head of the queue removes the derivative outright, collapsing the equation to the algebraic Model-A form (Models B_2^H and C_2^H , §9). These four regimes are gathered side by side in Table 5 of §11.

Proof. We multiply each balance equation by $x^{n_1} y^{n_2}$ and sum over its domain.

From the interior equation (4.3):

$$\begin{aligned}
& (\lambda_1 + \lambda_2 + \mu) \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} \pi(n_1, n_2) x^{n_1} y^{n_2} \\
& + \gamma_1 \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} n_1 \pi(n_1, n_2) x^{n_1} y^{n_2} \\
& + \gamma_2 \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} n_2 \pi(n_1, n_2) x^{n_1} y^{n_2} \\
& + \theta_1 \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} n_1 \pi(n_1, n_2) x^{n_1} y^{n_2} \\
& + \theta_2 \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} n_2 \pi(n_1, n_2) x^{n_1} y^{n_2} \\
& = \mu \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} \pi(n_1 + 1, n_2) x^{n_1} y^{n_2} \\
& + \lambda_1 \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} \pi(n_1 - 1, n_2) x^{n_1} y^{n_2} \\
& + \lambda_2 \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} \pi(n_1, n_2 - 1) x^{n_1} y^{n_2} \\
& + \gamma_1 \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} (n_1 + 1) \pi(n_1 + 1, n_2 - 1) x^{n_1} y^{n_2} \\
& + \gamma_2 \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} (n_2 + 1) \pi(n_1 - 1, n_2 + 1) x^{n_1} y^{n_2} \\
& + \theta_1 \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} (n_1 + 1) \pi(n_1 + 1, n_2) x^{n_1} y^{n_2} \\
& + \theta_2 \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} (n_2 + 1) \pi(n_1, n_2 + 1) x^{n_1} y^{n_2} \\
& LHS_1 + LHS_2 + LHS_3 + LHS_4 + LHS_5 \\
& = RHS_1 + RHS_2 + RHS_3 + RHS_4 + RHS_5 + RHS_6 + RHS_7
\end{aligned}$$

Evaluating each summand:

$$\begin{aligned}
& LHS_1 \\
& = (\lambda_1 + \lambda_2 + \mu) \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} \pi(n_1, n_2) x^{n_1} y^{n_2} \\
& = (\lambda_1 + \lambda_2 + \mu) \sum_{n_1=1}^{\infty} \left[\sum_{n_2=0}^{\infty} \pi(n_1, n_2) x^{n_1} y^{n_2} - \pi(n_1, 0) x^{n_1} y^0 \right] \\
& = (\lambda_1 + \lambda_2 + \mu) \left[\sum_{n_1=1}^{\infty} \sum_{n_2=0}^{\infty} \pi(n_1, n_2) x^{n_1} y^{n_2} - \sum_{n_1=1}^{\infty} \pi(n_1, 0) x^{n_1} \right] \\
& = (\lambda_1 + \lambda_2 + \mu) \left[\sum_{n_1=0}^{\infty} \sum_{n_2=0}^{\infty} \pi(n_1, n_2) x^{n_1} y^{n_2} - \sum_{n_2=0}^{\infty} \pi(0, n_2) y^{n_2} \right] \\
& \quad - (\lambda_1 + \lambda_2 + \mu) \left[\sum_{n_1=0}^{\infty} \pi(n_1, 0) x^{n_1} - \pi(0, 0) \right] \\
& = (\lambda_1 + \lambda_2 + \mu) [P(x, y) - P_y(y) - P_x(x) + \pi(0, 0)]
\end{aligned}$$

LHS_2

$$\begin{aligned}
&= \gamma_1 \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} n_1 \pi(n_1, n_2) x^{n_1} y^{n_2} \\
&= \gamma_1 x \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} n_1 \pi(n_1, n_2) x^{n_1-1} y^{n_2} \\
&= \gamma_1 x \left[\sum_{n_1=1}^{\infty} \sum_{n_2=0}^{\infty} n_1 \pi(n_1, n_2) x^{n_1-1} y^{n_2} - \sum_{n_1=1}^{\infty} n_1 \pi(n_1, 0) x^{n_1-1} \right] \\
&= \gamma_1 x \left[\sum_{n_1=0}^{\infty} \sum_{n_2=0}^{\infty} n_1 \pi(n_1, n_2) x^{n_1-1} y^{n_2} - \sum_{n_1=0}^{\infty} n_1 \pi(n_1, 0) x^{n_1-1} \right] \\
&= \gamma_1 x \left[\frac{\partial}{\partial x} P(x, y) - \frac{d}{dx} P_x(x) \right]
\end{aligned}$$

LHS_3

$$\begin{aligned}
&= \gamma_2 \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} n_2 \pi(n_1, n_2) x^{n_1} y^{n_2} \\
&= \gamma_2 \sum_{n_1=1}^{\infty} \left[\sum_{n_2=0}^{\infty} n_2 \pi(n_1, n_2) x^{n_1} y^{n_2} - 0 \cdot \pi(n_1, 0) x^{n_1} \right] \\
&= \gamma_2 y \left[\sum_{n_1=0}^{\infty} \sum_{n_2=0}^{\infty} n_2 \pi(n_1, n_2) x^{n_1} y^{n_2-1} - \sum_{n_2=0}^{\infty} n_2 \pi(0, n_2) y^{n_2-1} \right] \\
&= \gamma_2 y \left[\frac{\partial}{\partial y} P(x, y) - \frac{d}{dy} P_y(y) \right]
\end{aligned}$$

LHS_4

$$\begin{aligned}
&= \theta_1 \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} n_1 \pi(n_1, n_2) x^{n_1} y^{n_2} \\
&= \theta_1 x \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} n_1 \pi(n_1, n_2) x^{n_1-1} y^{n_2} \\
&= \theta_1 x \left[\sum_{n_1=1}^{\infty} \sum_{n_2=0}^{\infty} n_1 \pi(n_1, n_2) x^{n_1-1} y^{n_2} - \sum_{n_1=1}^{\infty} n_1 \pi(n_1, 0) x^{n_1-1} \right] \\
&= \theta_1 x \left[\sum_{n_1=0}^{\infty} \sum_{n_2=0}^{\infty} n_1 \pi(n_1, n_2) x^{n_1-1} y^{n_2} - \sum_{n_1=0}^{\infty} n_1 \pi(n_1, 0) x^{n_1-1} \right] \\
&= \theta_1 x \left[\frac{\partial}{\partial x} P(x, y) - \frac{d}{dx} P_x(x) \right]
\end{aligned}$$

LHS_5

$$\begin{aligned}
&= \theta_2 \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} n_2 \pi(n_1, n_2) x^{n_1} y^{n_2} \\
&= \theta_2 y \left[\sum_{n_1=0}^{\infty} \sum_{n_2=0}^{\infty} n_2 \pi(n_1, n_2) x^{n_1} y^{n_2-1} - \sum_{n_2=0}^{\infty} n_2 \pi(0, n_2) y^{n_2-1} \right] \\
&= \theta_2 y \left[\frac{\partial}{\partial y} P(x, y) - \frac{d}{dy} P_y(y) \right]
\end{aligned}$$

RHS_1

$$\begin{aligned}
&= \mu \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} \pi(n_1+1, n_2) x^{n_1} y^{n_2} \\
&= \mu x^{-1} \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} \pi(n_1+1, n_2) x^{n_1+1} y^{n_2} \\
&= \mu x^{-1} \sum_{m=2}^{\infty} \sum_{n_2=1}^{\infty} \pi(m, n_2) x^m y^{n_2} \\
&= \mu x^{-1} \sum_{m=2}^{\infty} \left[\sum_{n_2=0}^{\infty} \pi(m, n_2) x^m y^{n_2} - \pi(m, 0) x^m \right] \\
&= \mu x^{-1} \left[\sum_{m=2}^{\infty} \sum_{n_2=0}^{\infty} \pi(m, n_2) x^m y^{n_2} - \sum_{m=2}^{\infty} \pi(m, 0) x^m \right] \\
&= \mu x^{-1} \left[\sum_{m=0}^{\infty} \sum_{n_2=0}^{\infty} \pi(m, n_2) x^m y^{n_2} - x \sum_{n_2=0}^{\infty} \pi(1, n_2) y^{n_2} - \sum_{n_2=0}^{\infty} \pi(0, n_2) y^{n_2} \right] \\
&\quad - \mu x^{-1} \left[\sum_{m=0}^{\infty} \pi(m, 0) x^m - x \pi(1, 0) - \pi(0, 0) \right] \\
&= \mu x^{-1} [P(x, y) - x P_y(y|n_1=1) - P_y(y) - P_x(x) + x \pi(1, 0) + \pi(0, 0)]
\end{aligned}$$

RHS_2

$$\begin{aligned}
&= \lambda_1 \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} \pi(n_1-1, n_2) x^{n_1} y^{n_2} \\
&= \lambda_1 x \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} \pi(n_1-1, n_2) x^{n_1-1} y^{n_2} \\
&= \lambda_1 x \sum_{m=0}^{\infty} \sum_{n_2=1}^{\infty} \pi(m, n_2) x^m y^{n_2} \\
&= \lambda_1 x \sum_{m=0}^{\infty} \left[\sum_{n_2=0}^{\infty} \pi(m, n_2) x^m y^{n_2} - \pi(m, 0) x^m \right] \\
&= \lambda_1 x \left[\sum_{m=0}^{\infty} \sum_{n_2=0}^{\infty} \pi(m, n_2) x^m y^{n_2} - \sum_{m=0}^{\infty} \pi(m, 0) x^m \right] \\
&= \lambda_1 x [P(x, y) - P_x(x)]
\end{aligned}$$

RHS_3

$$\begin{aligned}
&= \lambda_2 \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} \pi(n_1, n_2-1) x^{n_1} y^{n_2} \\
&= \lambda_2 y \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} \pi(n_1, n_2-1) x^{n_1} y^{n_2-1} \\
&= \lambda_2 y \sum_{n_1=1}^{\infty} \sum_{m=0}^{\infty} \pi(n_1, m) x^{n_1} y^m \\
&= \lambda_2 y \left[\sum_{n_1=0}^{\infty} \sum_{m=0}^{\infty} \pi(n_1, m) x^{n_1} y^m - \sum_{m=0}^{\infty} \pi(0, m) y^m \right] \\
&= \lambda_2 y [P(x, y) - P_y(y)]
\end{aligned}$$

RHS_4

$$\begin{aligned}
&= \gamma_1 \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} (n_1+1)\pi(n_1+1, n_2-1)x^{n_1}y^{n_2} \\
&= \gamma_1 y \sum_{n_1=1}^{\infty} \sum_{m_2=0}^{\infty} (n_1+1)\pi(n_1+1, m_2)x^{n_1}y^{m_2} \\
&= \gamma_1 y \sum_{m_1=2}^{\infty} \sum_{m_2=0}^{\infty} m_1\pi(m_1, m_2)x^{m_1-1}y^{m_2} \\
&= \gamma_1 y \left[\sum_{m_1=0}^{\infty} \sum_{m_2=0}^{\infty} m_1\pi(m_1, m_2)x^{m_1-1}y^{m_2} - \sum_{m_2=0}^{\infty} \pi(1, m_2)y^{m_2} \right] \\
&= \gamma_1 y \left[\frac{\partial}{\partial x} P(x, y) - P_y(y|n_1=1) \right]
\end{aligned}$$

RHS_5

$$\begin{aligned}
&= \gamma_2 \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} (n_2+1)\pi(n_1-1, n_2+1)x^{n_1}y^{n_2} \\
&= \gamma_2 \sum_{n_1=1}^{\infty} \sum_{m_2=2}^{\infty} m_2\pi(n_1-1, m_2)x^{n_1}y^{m_2-1} \\
&= \gamma_2 x \left[\sum_{n_1=1}^{\infty} \sum_{m_2=0}^{\infty} m_2\pi(n_1-1, m_2)x^{n_1-1}y^{m_2-1} - \sum_{n_1=1}^{\infty} \pi(n_1-1, 1)x^{n_1-1} \right] \\
&= \gamma_2 x \left[\sum_{m_1=0}^{\infty} \sum_{m_2=0}^{\infty} m_2\pi(m_1, m_2)x^{m_1}y^{m_2-1} - \sum_{m_1=0}^{\infty} \pi(m_1, 1)x^{m_1} \right] \\
&= \gamma_2 x \left[\frac{\partial}{\partial y} P(x, y) - P_x(x|n_2=1) \right]
\end{aligned}$$

RHS_6

$$\begin{aligned}
&= \theta_1 \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} (n_1+1)\pi(n_1+1, n_2)x^{n_1}y^{n_2} \\
&= \theta_1 \sum_{m_1=2}^{\infty} \sum_{n_2=1}^{\infty} m_1\pi(m_1, n_2)x^{m_1-1}y^{n_2} \\
&= \theta_1 \left[\sum_{m_1=0}^{\infty} \sum_{n_2=1}^{\infty} m_1\pi(m_1, n_2)x^{m_1-1}y^{n_2} - \sum_{n_2=1}^{\infty} \pi(1, n_2)y^{n_2} \right] \\
&= \theta_1 \left[\sum_{m_1=0}^{\infty} \sum_{n_2=0}^{\infty} m_1\pi(m_1, n_2)x^{m_1-1}y^{n_2} - \sum_{m_1=0}^{\infty} m_1\pi(m_1, 0)x^{m_1-1} \right] \\
&\quad - \theta_1 \left[\sum_{n_2=0}^{\infty} \pi(1, n_2)y^{n_2} - \pi(1, 0) \right] \\
&= \theta_1 \left[\frac{\partial}{\partial x} P(x, y) - \frac{d}{dx} P_x(x) - P_y(y|n_1=1) + \pi(1, 0) \right]
\end{aligned}$$

RHS_7

$$\begin{aligned}
&= \theta_2 \sum_{n_1=1}^{\infty} \sum_{n_2=1}^{\infty} (n_2+1) \pi(n_1, n_2+1) x^{n_1} y^{n_2} \\
&= \theta_2 \sum_{n_1=1}^{\infty} \sum_{m_2=2}^{\infty} m_2 \pi(n_1, m_2) x^{n_1} y^{m_2-1} \\
&= \theta_2 \left[\sum_{n_1=1}^{\infty} \sum_{m_2=0}^{\infty} m_2 \pi(n_1, m_2) x^{n_1} y^{m_2-1} - \sum_{n_1=1}^{\infty} \pi(n_1, 1) x^{n_1} \right] \\
&= \theta_2 \left[\sum_{n_1=0}^{\infty} \sum_{m_2=0}^{\infty} m_2 \pi(n_1, m_2) x^{n_1} y^{m_2-1} - \sum_{m_2=0}^{\infty} m_2 \pi(0, m_2) y^{m_2-1} \right] \\
&\quad - \theta_2 \left[\sum_{n_1=0}^{\infty} \pi(n_1, 1) x^{n_1} - \pi(0, 1) \right] \\
&= \theta_2 \left[\frac{\partial}{\partial y} P(x, y) - \frac{d}{dy} P_y(y) - P_x(x|n_2=1) + \pi(0, 1) \right]
\end{aligned}$$

Collecting and multiplying through by xy to clear the x^{-1} factor:

$$\begin{aligned}
&(\lambda_1 + \lambda_2 + \mu) [P(x, y) - P_y(y) - P_x(x) + \pi(0, 0)] \\
&+ \gamma_1 x \left[\frac{\partial}{\partial x} P(x, y) - \frac{d}{dx} P_x(x) \right] \\
&+ \gamma_2 y \left[\frac{\partial}{\partial y} P(x, y) - \frac{d}{dy} P_y(y) \right] \\
&+ \theta_1 x \left[\frac{\partial}{\partial x} P(x, y) - \frac{d}{dx} P_x(x) \right] \\
&+ \theta_2 y \left[\frac{\partial}{\partial y} P(x, y) - \frac{d}{dy} P_y(y) \right] \\
&= \mu x^{-1} [P(x, y) - x P_y(y|n_1=1) - P_y(y) - P_x(x) + x \pi(1, 0) + \pi(0, 0)] \\
&\quad + \lambda_1 x [P(x, y) - P_x(x)] \\
&\quad + \lambda_2 y [P(x, y) - P_y(y)] \\
&\quad + \gamma_1 y \left[\frac{\partial}{\partial x} P(x, y) - P_y(y|n_1=1) \right] \\
&\quad + \gamma_2 x \left[\frac{\partial}{\partial y} P(x, y) - P_x(x|n_2=1) \right] \\
&\quad + \theta_1 \left[\frac{\partial}{\partial x} P(x, y) - \frac{d}{dx} P_x(x) - P_y(y|n_1=1) + \pi(1, 0) \right] \\
&\quad + \theta_2 \left[\frac{\partial}{\partial y} P(x, y) - \frac{d}{dy} P_y(y) - P_x(x|n_2=1) + \pi(0, 1) \right]
\end{aligned}$$

$$\begin{aligned}
& [(\lambda_1 + \lambda_2 + \mu) - \mu x^{-1} - \lambda_1 x - \lambda_2 y] P(x, y) \\
& + [\gamma_1(x - y) + \theta_1(x - 1)] \frac{\partial}{\partial x} P(x, y) \\
& + [\gamma_2(y - x) + \theta_2(y - 1)] \frac{\partial}{\partial y} P(x, y) \\
& = [(\lambda_1 + \lambda_2 + \mu) - \mu x^{-1} - \lambda_1 x] P_x(x) \\
& + [(\lambda_1 + \lambda_2 + \mu) - \mu x^{-1} - \lambda_2 y] P_y(y) \\
& + [-(\lambda_1 + \lambda_2 + \mu) + \mu x^{-1}] \pi(0, 0) \\
& + [\gamma_1 x + \theta_1(x - 1)] \frac{d}{dx} P_x(x) \\
& + [\gamma_2 y + \theta_2(y - 1)] \frac{d}{dy} P_y(y) \\
& - [\mu + \gamma_1 y + \theta_1] P_y(y|n_1 = 1) \\
& - [\gamma_2 x + \theta_2] P_x(x|n_2 = 1) \\
& + \mu \pi(1, 0) \\
& + \theta_1 \pi(1, 0) \\
& + \theta_2 \pi(0, 1) \\
& [(\lambda_1 + \lambda_2 + \mu)xy - \mu y - \lambda_1 x^2 y - \lambda_2 xy^2] P(x, y) \\
& + xy [\gamma_1(x - y) + \theta_1(x - 1)] \frac{\partial}{\partial x} P(x, y) \\
& + xy [\gamma_2(y - x) + \theta_2(y - 1)] \frac{\partial}{\partial y} P(x, y) \\
& = [(\lambda_1 + \lambda_2 + \mu)xy - \mu y - \lambda_1 x^2 y] P_x(x) \\
& + [(\lambda_1 + \lambda_2 + \mu)xy - \mu y - \lambda_2 xy^2] P_y(y) \\
& - [(\lambda_1 + \lambda_2 + \mu)xy - \mu y] \pi(0, 0) \\
& + xy [\gamma_1 x + \theta_1(x - 1)] \frac{d}{dx} P_x(x) \\
& + xy [\gamma_2 y + \theta_2(y - 1)] \frac{d}{dy} P_y(y) \\
& - xy [\mu + \gamma_1 y + \theta_1] P_y(y|n_1 = 1) \\
& - xy [\gamma_2 x + \theta_2] P_x(x|n_2 = 1) \\
& + (\mu + \theta_1)xy \pi(1, 0) \\
& + \theta_2 xy \pi(0, 1)
\end{aligned} \tag{4.15}$$

The x -boundary equation (4.5), treated analogously, gives:

$$\begin{aligned}
& \sum_{n_1=1}^{\infty} [\lambda_1 + \lambda_2 + \mu + (\gamma_1 + \theta_1)n_1] \pi(n_1, 0)x^{n_1} \\
& = \lambda_1 \sum_{n_1=1}^{\infty} \pi(n_1 - 1, 0)x^{n_1} + \mu \sum_{n_1=1}^{\infty} \pi(n_1 + 1, 0)x^{n_1} \\
& + \gamma_2 \sum_{n_1=1}^{\infty} \pi(n_1 - 1, 1)x^{n_1} + \theta_1 \sum_{n_1=1}^{\infty} (n_1 + 1)\pi(n_1 + 1, 0)x^{n_1} + \theta_2 \sum_{n_1=1}^{\infty} \pi(n_1, 1)x^{n_1} \\
& (\lambda_1 + \lambda_2 + \mu) \sum_{n_1=1}^{\infty} \pi(n_1, 0)x^{n_1} + (\gamma_1 + \theta_1)x \sum_{n_1=1}^{\infty} n_1 \pi(n_1, 0)x^{n_1-1} \\
& = \lambda_1 x \sum_{m=0}^{\infty} \pi(m, 0)x^m + \mu x^{-1} \sum_{m=2}^{\infty} \pi(m, 0)x^m \\
& + \gamma_2 x \sum_{m=0}^{\infty} \pi(m, 1)x^m + \theta_1 \sum_{m=2}^{\infty} m \pi(m, 0)x^{m-1} + \theta_2 \sum_{m=1}^{\infty} \pi(m, 1)x^m
\end{aligned}$$

$$\begin{aligned}
& (\lambda_1 + \lambda_2 + \mu) \left[\sum_{n_1=0}^{\infty} \pi(n_1, 0) x^{n_1} - \pi(0, 0) \right] + (\gamma_1 + \theta_1) x \sum_{n_1=0}^{\infty} n_1 \pi(n_1, 0) x^{n_1-1} \\
&= \lambda_1 x \sum_{m=0}^{\infty} \pi(m, 0) x^m \\
&\quad + \mu x^{-1} \left[\sum_{m=0}^{\infty} \pi(m, 0) x^m - x \pi(1, 0) - \pi(0, 0) \right] \\
&\quad + \gamma_2 x \sum_{m=0}^{\infty} \pi(m, 1) x^m \\
&\quad + \theta_1 \left[\sum_{m=0}^{\infty} m \pi(m, 0) x^{m-1} - \pi(1, 0) \right] \\
&\quad + \theta_2 \left[\sum_{m=0}^{\infty} \pi(m, 1) x^m - \pi(0, 1) \right] \\
& (\lambda_1 + \lambda_2 + \mu) [P_x(x) - \pi(0, 0)] + (\gamma_1 + \theta_1) x \frac{d}{dx} P_x(x) \\
&= \lambda_1 x P_x(x) + \mu x^{-1} [P_x(x) - x \pi(1, 0) - \pi(0, 0)] \\
&\quad + \gamma_2 x P_x(x | n_2 = 1) + \theta_1 \left[\frac{d}{dx} P_x(x) - \pi(1, 0) \right] + \theta_2 [P_x(x | n_2 = 1) - \pi(0, 1)] \\
& [(\lambda_1 + \lambda_2 + \mu) - \mu x^{-1} - \lambda_1 x] P_x(x) + [\gamma_1 x + \theta_1(x - 1)] \frac{d}{dx} P_x(x) \\
&= [(\lambda_1 + \lambda_2 + \mu) - \mu x^{-1}] \pi(0, 0) - (\mu + \theta_1) \pi(1, 0) - \theta_2 \pi(0, 1) \\
&\quad + [\gamma_2 x + \theta_2] P_x(x | n_2 = 1) \\
& [(\lambda_1 + \lambda_2 + \mu) x y - \mu y - \lambda_1 x^2 y] P_x(x) + x y [\gamma_1 x + \theta_1(x - 1)] \frac{d}{dx} P_x(x) \\
&= [(\lambda_1 + \lambda_2 + \mu) x y - \mu y] \pi(0, 0) \\
&\quad - (\mu + \theta_1) x y \pi(1, 0) \\
&\quad - \theta_2 x y \pi(0, 1) \\
&\quad + x y [\gamma_2 x + \theta_2] P_x(x | n_2 = 1)
\end{aligned} \tag{4.16}$$

The y -boundary equation (4.4) gives:

$$\begin{aligned}
& \sum_{n_2=1}^{\infty} [\lambda_1 + \lambda_2 + \mu + (\gamma_2 + \theta_2) n_2] \pi(0, n_2) y^{n_2} \\
&= \lambda_2 \sum_{n_2=1}^{\infty} \pi(0, n_2 - 1) y^{n_2} \\
&\quad + \mu \sum_{n_2=1}^{\infty} \pi(1, n_2) y^{n_2} \\
&\quad + \mu \sum_{n_2=1}^{\infty} \pi(0, n_2 + 1) y^{n_2} \\
&\quad + \gamma_1 \sum_{n_2=1}^{\infty} \pi(1, n_2 - 1) y^{n_2} \\
&\quad + \theta_1 \sum_{n_2=1}^{\infty} \pi(1, n_2) y^{n_2} \\
&\quad + \theta_2 \sum_{n_2=1}^{\infty} (n_2 + 1) \pi(0, n_2 + 1) y^{n_2}
\end{aligned}$$

$$\begin{aligned}
& (\lambda_1 + \lambda_2 + \mu) \sum_{n_2=1}^{\infty} \pi(0, n_2) y^{n_2} \\
& + (\gamma_2 + \theta_2) y \sum_{n_2=1}^{\infty} n_2 \pi(0, n_2) y^{n_2-1} \\
& = \lambda_2 y \sum_{m=0}^{\infty} \pi(0, m) y^m \\
& + \mu \sum_{n_2=1}^{\infty} \pi(1, n_2) y^{n_2} \\
& + \mu y^{-1} \sum_{m=2}^{\infty} \pi(0, m) y^m \\
& + \gamma_1 y \sum_{m=0}^{\infty} \pi(1, m) y^m \\
& + \theta_1 \sum_{n_2=1}^{\infty} \pi(1, n_2) y^{n_2} \\
& + \theta_2 \sum_{m=2}^{\infty} m \pi(0, m) y^{m-1} \\
& (\lambda_1 + \lambda_2 + \mu) \left[\sum_{n_2=0}^{\infty} \pi(0, n_2) y^{n_2} - \pi(0, 0) \right] \\
& + (\gamma_2 + \theta_2) y \sum_{n_2=0}^{\infty} n_2 \pi(0, n_2) y^{n_2-1} \\
& = \lambda_2 y \sum_{m=0}^{\infty} \pi(0, m) y^m \\
& + \mu \left[\sum_{n_2=0}^{\infty} \pi(1, n_2) y^{n_2} - \pi(1, 0) \right] \\
& + \mu y^{-1} \left[\sum_{m=0}^{\infty} \pi(0, m) y^m - y \pi(0, 1) - \pi(0, 0) \right] \\
& + \gamma_1 y \sum_{m=0}^{\infty} \pi(1, m) y^m \\
& + \theta_1 \left[\sum_{n_2=0}^{\infty} \pi(1, n_2) y^{n_2} - \pi(1, 0) \right] \\
& + \theta_2 \left[\sum_{m=0}^{\infty} m \pi(0, m) y^{m-1} - \pi(0, 1) \right] \\
& (\lambda_1 + \lambda_2 + \mu) [P_y(y) - \pi(0, 0)] + (\gamma_2 + \theta_2) y \frac{d}{dy} P_y(y) \\
& = \lambda_2 y P_y(y) + \mu [P_y(y|n_1 = 1) - \pi(1, 0)] \\
& + \mu y^{-1} [P_y(y) - y \pi(0, 1) - \pi(0, 0)] + \gamma_1 y P_y(y|n_1 = 1) \\
& + \theta_1 [P_y(y|n_1 = 1) - \pi(1, 0)] + \theta_2 \left[\frac{d}{dy} P_y(y) - \pi(0, 1) \right] \\
& [(\lambda_1 + \lambda_2 + \mu) - \mu y^{-1} - \lambda_2 y] P_y(y) + [\gamma_2 y + \theta_2(y - 1)] \frac{d}{dy} P_y(y) \\
& = [\mu + \gamma_1 y + \theta_1] P_y(y|n_1 = 1) \\
& - (\mu + \theta_1) \pi(1, 0) - (\mu + \theta_2) \pi(0, 1) \\
& + [(\lambda_1 + \lambda_2 + \mu) - \mu y^{-1}] \pi(0, 0)
\end{aligned}$$

$$\begin{aligned}
& [(\lambda_1 + \lambda_2 + \mu) - \mu y^{-1} - \lambda_2 y] P_y(y) + [\gamma_2 y + \theta_2(y-1)] \frac{d}{dy} P_y(y) \\
&= [\mu + \gamma_1 y + \theta_1] P_y(y|n_1=1) \\
&\quad + [(\lambda_1 + \lambda_2 + \mu) - \mu y^{-1} - (\lambda_1 + \lambda_2)] \pi(0,0) \\
& [(\lambda_1 + \lambda_2 + \mu) - \mu y^{-1} - \lambda_2 y] P_y(y) + [\gamma_2 y + \theta_2(y-1)] \frac{d}{dy} P_y(y) \\
&= [\mu + \gamma_1 y + \theta_1] P_y(y|n_1=1) + [\mu - \mu y^{-1}] \pi(0,0) \\
& [(\lambda_1 + \lambda_2 + \mu)xy - \mu x - \lambda_2 xy^2] P_y(y) + xy [\gamma_2 y + \theta_2(y-1)] \frac{d}{dy} P_y(y) \\
&= [\mu + \gamma_1 y + \theta_1] xy P_y(y|n_1=1) - \mu x(1-y) \pi(0,0)
\end{aligned} \tag{4.17}$$

Substituting (4.16) into (4.15) eliminates all P_x terms:

$$\begin{aligned}
& [(\lambda_1 + \lambda_2 + \mu)xy - \mu y - \lambda_1 x^2 y - \lambda_2 xy^2] P(x, y) \\
&+ xy [\gamma_1(x-y) + \theta_1(x-1)] \frac{\partial}{\partial x} P(x, y) \\
&+ xy [\gamma_2(y-x) + \theta_2(y-1)] \frac{\partial}{\partial y} P(x, y) \\
&= [(\lambda_1 + \lambda_2 + \mu)xy - \mu y - \lambda_2 xy^2] P_y(y) \\
&\quad + xy [\gamma_2 y + \theta_2(y-1)] \frac{d}{dy} P_y(y) \\
&\quad - xy [\mu + \gamma_1 y + \theta_1] P_y(y|n_1=1)
\end{aligned} \tag{4.18}$$

Isolating $[\mu + \gamma_1 y + \theta_1] xy P_y(y|n_1=1)$ from (4.17) and substituting into (4.18):

$$\begin{aligned}
& [(\lambda_1 + \lambda_2 + \mu)xy - \mu y - \lambda_1 x^2 y - \lambda_2 xy^2] P(x, y) \\
&+ xy [\gamma_1(x-y) + \theta_1(x-1)] \frac{\partial}{\partial x} P(x, y) \\
&+ xy [\gamma_2(y-x) + \theta_2(y-1)] \frac{\partial}{\partial y} P(x, y) \\
&= [(\lambda_1 + \lambda_2 + \mu)xy - \mu y - \lambda_2 xy^2] P_y(y) \\
&\quad + xy [\gamma_2 y + \theta_2(y-1)] \frac{d}{dy} P_y(y) \\
&\quad - \left\{ [(\lambda_1 + \lambda_2 + \mu)xy - \mu x - \lambda_2 xy^2] P_y(y) \right. \\
&\quad \left. + xy [\gamma_2 y + \theta_2(y-1)] \frac{d}{dy} P_y(y) + \mu x(1-y) \pi(0,0) \right\} \\
& [(\lambda_1 + \lambda_2 + \mu)xy - \mu y - \lambda_1 x^2 y - \lambda_2 xy^2] P(x, y) \\
&+ xy [\gamma_1(x-y) + \theta_1(x-1)] \frac{\partial}{\partial x} P(x, y) \\
&+ xy [\gamma_2(y-x) + \theta_2(y-1)] \frac{\partial}{\partial y} P(x, y) \\
&= \mu(x-y) P_y(y) - \mu x(1-y) \pi(0,0)
\end{aligned}$$

□

4.2.2 Equation for PPGF in state space $\tilde{S}$

Lemma 2 (Fundamental equation for the PPGF on $\tilde{S}$). *Consider the same system on the state space $\tilde{S} = \{(n_2, n) : 0 \leq n_2 \leq n\}$, where n is the total number of customers waiting in the queues and n_2 the number of waiting class-2 customers. Assume positive recurrence, so that the stationary distribution $\{\tilde{\pi}(n_2, n)\}$ exists, and define the partial PGF (transforming only n_2 , keeping n discrete)*

$$\tilde{P}(y, n) = \sum_{n_2=0}^n \tilde{\pi}(n_2, n) y^{n_2},$$

a polynomial of degree n in y . Then, for every $n \geq 1$, the family $\{\tilde{P}(y, n)\}_{n \geq 0}$ satisfies the differential-difference equation

$$\begin{aligned} & [-(\gamma_2 + \gamma_1 y)(1 - y) + (\theta_2 - \theta_1)y] \frac{d}{dy} \tilde{P}(y, n) + [\lambda_1 + \lambda_2 + \mu + \gamma_1 n(1 - y) + \theta_1 n] \tilde{P}(y, n) \\ &= (\lambda_1 + \lambda_2 y) \tilde{P}(y, n - 1) + [\mu + \theta_1(n + 1)] \tilde{P}(y, n + 1) \\ &+ (\theta_2 - \theta_1 y) \frac{d}{dy} \tilde{P}(y, n + 1) + \mu y^n (1 - y) \tilde{\pi}(n + 1, n + 1). \end{aligned}$$

Proof. Multiply the interior equation (4.8) by y^{n_2} and sum over $1 \leq n_2 \leq n - 1$:

$$\begin{aligned} & (\lambda_1 + \lambda_2 + \mu) \sum_{n_2=1}^{n-1} \tilde{\pi}(n_2, n) y^{n_2} \\ &+ \gamma_1 \sum_{n_2=1}^{n-1} (n - n_2) \tilde{\pi}(n_2, n) y^{n_2} \\ &+ \gamma_2 \sum_{n_2=1}^{n-1} n_2 \tilde{\pi}(n_2, n) y^{n_2} \\ &+ \theta_1 \sum_{n_2=1}^{n-1} (n - n_2) \tilde{\pi}(n_2, n) y^{n_2} \\ &+ \theta_2 \sum_{n_2=1}^{n-1} n_2 \tilde{\pi}(n_2, n) y^{n_2} \\ &= \lambda_1 \sum_{n_2=1}^{n-1} \tilde{\pi}(n_2, n - 1) y^{n_2} \\ &+ \lambda_2 \sum_{n_2=1}^{n-1} \tilde{\pi}(n_2 - 1, n - 1) y^{n_2} \\ &+ \mu \sum_{n_2=1}^{n-1} \tilde{\pi}(n_2, n + 1) y^{n_2} \\ &+ \gamma_1 \sum_{n_2=1}^{n-1} (n - n_2 + 1) \tilde{\pi}(n_2 - 1, n) y^{n_2} \\ &+ \gamma_2 \sum_{n_2=1}^{n-1} (n_2 + 1) \tilde{\pi}(n_2 + 1, n) y^{n_2} \\ &+ \theta_1 \sum_{n_2=1}^{n-1} (n + 1 - n_2) \tilde{\pi}(n_2, n + 1) y^{n_2} \\ &+ \theta_2 \sum_{n_2=1}^{n-1} (n_2 + 1) \tilde{\pi}(n_2 + 1, n + 1) y^{n_2} \end{aligned}$$

Adding the diagonal equation (4.9) multiplied by y^n extends most summation limits to n ; the λ_1 - and γ_2 -terms on the RHS remain restricted to $n_2 \leq n - 1$, and a spare $\mu \tilde{\pi}(n + 1, n + 1) y^n$

survives:

$$\begin{aligned}
& (\lambda_1 + \lambda_2 + \mu) \sum_{n_2=1}^n \tilde{\pi}(n_2, n) y^{n_2} \\
& + \gamma_1 \sum_{n_2=1}^n (n - n_2) \tilde{\pi}(n_2, n) y^{n_2} \\
& + \gamma_2 \sum_{n_2=1}^n n_2 \tilde{\pi}(n_2, n) y^{n_2} \\
& + \theta_1 \sum_{n_2=1}^n (n - n_2) \tilde{\pi}(n_2, n) y^{n_2} \\
& + \theta_2 \sum_{n_2=1}^n n_2 \tilde{\pi}(n_2, n) y^{n_2} \\
& = \lambda_1 \sum_{n_2=1}^{n-1} \tilde{\pi}(n_2, n-1) y^{n_2} \\
& + \lambda_2 \sum_{n_2=1}^n \tilde{\pi}(n_2-1, n-1) y^{n_2} \\
& + \mu \sum_{n_2=1}^n \tilde{\pi}(n_2, n+1) y^{n_2} \\
& + \mu \tilde{\pi}(n+1, n+1) y^n \\
& + \gamma_1 \sum_{n_2=1}^n (n - n_2 + 1) \tilde{\pi}(n_2-1, n) y^{n_2} \\
& + \gamma_2 \sum_{n_2=1}^{n-1} (n_2 + 1) \tilde{\pi}(n_2+1, n) y^{n_2} \\
& + \theta_1 \sum_{n_2=1}^n (n+1 - n_2) \tilde{\pi}(n_2, n+1) y^{n_2} \\
& + \theta_2 \sum_{n_2=1}^n (n_2 + 1) \tilde{\pi}(n_2+1, n+1) y^{n_2}.
\end{aligned}$$

Evaluating each summand:

LHS_1

$$\begin{aligned}
& = (\lambda_1 + \lambda_2 + \mu) \sum_{n_2=1}^n \tilde{\pi}(n_2, n) y^{n_2} \\
& = (\lambda_1 + \lambda_2 + \mu) \left[\sum_{n_2=0}^n \tilde{\pi}(n_2, n) y^{n_2} - \tilde{\pi}(0, n) y^0 \right] \\
& = (\lambda_1 + \lambda_2 + \mu) \left[\tilde{P}(y, n) - \tilde{\pi}(0, n) \right]
\end{aligned}$$

LHS_2

$$\begin{aligned}
& = \gamma_1 \sum_{n_2=1}^n (n - n_2) \tilde{\pi}(n_2, n) y^{n_2} \\
& = \gamma_1 \left[n \sum_{n_2=1}^n \tilde{\pi}(n_2, n) y^{n_2} - \sum_{n_2=1}^n n_2 \tilde{\pi}(n_2, n) y^{n_2} \right] \\
& = \gamma_1 \left[n \left(\sum_{n_2=0}^n \tilde{\pi}(n_2, n) y^{n_2} - \tilde{\pi}(0, n) y^0 \right) - y \sum_{n_2=1}^n n_2 \tilde{\pi}(n_2, n) y^{n_2-1} \right] \\
& = \gamma_1 \left[n(\tilde{P}(y, n) - \tilde{\pi}(0, n)) - y \frac{d}{dy} \tilde{P}(y, n) \right]
\end{aligned}$$

LHS_3

$$\begin{aligned}
&= \gamma_2 \sum_{n_2=1}^n n_2 \tilde{\pi}(n_2, n) y^{n_2} \\
&= \gamma_2 \left[y \sum_{n_2=1}^n n_2 \tilde{\pi}(n_2, n) y^{n_2-1} \right] \\
&= \gamma_2 y \frac{d}{dy} \tilde{P}(y, n)
\end{aligned}$$

LHS_4

$$\begin{aligned}
&= \theta_1 \sum_{n_2=1}^n (n - n_2) \tilde{\pi}(n_2, n) y^{n_2} \\
&= \theta_1 \left[n(\tilde{P}(y, n) - \tilde{\pi}(0, n)) - y \frac{d}{dy} \tilde{P}(y, n) \right]
\end{aligned}$$

LHS_5

$$\begin{aligned}
&= \theta_2 \sum_{n_2=1}^n n_2 \tilde{\pi}(n_2, n) y^{n_2} \\
&= \theta_2 y \frac{d}{dy} \tilde{P}(y, n)
\end{aligned}$$

RHS_1

$$\begin{aligned}
&= \lambda_1 \sum_{n_2=1}^{n-1} \tilde{\pi}(n_2, n-1) y^{n_2} \\
&= \lambda_1 \left[\sum_{n_2=1}^n \tilde{\pi}(n_2, n-1) y^{n_2} - \tilde{\pi}(n, n-1) y^n \right] \\
&= \lambda_1 \left[\sum_{n_2=0}^n \tilde{\pi}(n_2, n-1) y^{n_2} - \tilde{\pi}(n, n-1) y^n - \tilde{\pi}(0, n-1) y^0 \right] \\
&= \lambda_1 \left[\tilde{P}(y, n-1) - \tilde{\pi}(n, n-1) y^n - \tilde{\pi}(0, n-1) \right]
\end{aligned}$$

RHS_2

$$\begin{aligned}
&= \lambda_2 \sum_{n_2=1}^n \tilde{\pi}(n_2-1, n-1) y^{n_2} \\
&= \lambda_2 \left[y \sum_{n_2=1}^n \tilde{\pi}(n_2-1, n-1) y^{n_2-1} \right] \\
&= \lambda_2 \left[y \sum_{n_2=0}^{n-1} \tilde{\pi}(n_2, n-1) y^{n_2} \right] \\
&= \lambda_2 \left[y \left(\sum_{n_2=0}^n \tilde{\pi}(n_2, n-1) y^{n_2} - \tilde{\pi}(n, n-1) y^n \right) \right] \\
&= \lambda_2 y \left[\tilde{P}(y, n-1) - \tilde{\pi}(n, n-1) y^n \right]
\end{aligned}$$

RHS_3

$$\begin{aligned}
&= \mu \tilde{\pi}(n+1, n+1) y^n + \mu \sum_{n_2=1}^n \tilde{\pi}(n_2, n+1) y^{n_2} \\
&= \mu \left[\tilde{\pi}(n+1, n+1) y^n + \sum_{n_2=1}^n \tilde{\pi}(n_2, n+1) y^{n_2} \right] \\
&= \mu \left[\tilde{\pi}(n+1, n+1) y^n + \sum_{n_2=0}^{n+1} \tilde{\pi}(n_2, n+1) y^{n_2} - \tilde{\pi}(0, n+1) y^0 - \tilde{\pi}(n+1, n+1) y^{n+1} \right] \\
&= \mu \left[\tilde{P}(y, n+1) - \tilde{\pi}(0, n+1) + \tilde{\pi}(n+1, n+1) y^n (1-y) \right]
\end{aligned}$$

RHS_4

$$\begin{aligned}
&= \gamma_1 \sum_{n_2=1}^n (n - n_2 + 1) \tilde{\pi}(n_2 - 1, n) y^{n_2} \\
&= \gamma_1 \left[y \sum_{n_2=1}^n (n - (n_2 - 1)) \tilde{\pi}(n_2 - 1, n) y^{n_2-1} \right] \\
&= \gamma_1 \left[y \sum_{k=0}^{n-1} (n - k) \tilde{\pi}(k, n) y^k \right] \\
&= \gamma_1 \left[y \sum_{k=0}^n (n - k) \tilde{\pi}(k, n) y^k \right] \quad (\text{since the } k = n \text{ term is } 0) \\
&= \gamma_1 y \left[n \sum_{k=0}^n \tilde{\pi}(k, n) y^k - y \sum_{k=1}^n k \tilde{\pi}(k, n) y^{k-1} \right] \\
&= \gamma_1 y \left[n \tilde{P}(y, n) - y \frac{d}{dy} \tilde{P}(y, n) \right]
\end{aligned}$$

RHS_5

$$\begin{aligned}
&= \gamma_2 \sum_{n_2=1}^{n-1} (n_2 + 1) \tilde{\pi}(n_2 + 1, n) y^{n_2} \\
&= \gamma_2 \sum_{k=2}^n k \tilde{\pi}(k, n) y^{k-1} \quad (\text{letting } k = n_2 + 1) \\
&= \gamma_2 \left[\sum_{k=1}^n k \tilde{\pi}(k, n) y^{k-1} - 1 \cdot \tilde{\pi}(1, n) y^0 \right] \\
&= \gamma_2 \left[\frac{d}{dy} \tilde{P}(y, n) - \tilde{\pi}(1, n) \right]
\end{aligned}$$

RHS_6

$$\begin{aligned}
&= \theta_1 \sum_{n_2=1}^n (n + 1 - n_2) \tilde{\pi}(n_2, n + 1) y^{n_2} \\
&= \theta_1 \left[(n + 1) \sum_{n_2=1}^n \tilde{\pi}(n_2, n + 1) y^{n_2} - \sum_{n_2=1}^n n_2 \tilde{\pi}(n_2, n + 1) y^{n_2} \right] \\
&= \theta_1 \left[(n + 1) \left(\sum_{n_2=0}^{n+1} \tilde{\pi}(n_2, n + 1) y^{n_2} - \tilde{\pi}(0, n + 1) - \tilde{\pi}(n + 1, n + 1) y^{n+1} \right) \right] \\
&\quad - \theta_1 \left[y \left(\sum_{n_2=0}^{n+1} n_2 \tilde{\pi}(n_2, n + 1) y^{n_2-1} - (n + 1) \tilde{\pi}(n + 1, n + 1) y^n \right) \right] \\
&= \theta_1 \left[(n + 1) (\tilde{P}(y, n + 1) - \tilde{\pi}(0, n + 1)) - (n + 1) \tilde{\pi}(n + 1, n + 1) y^{n+1} \right] \\
&\quad - \theta_1 \left[y \frac{d}{dy} \tilde{P}(y, n + 1) - (n + 1) \tilde{\pi}(n + 1, n + 1) y^{n+1} \right] \\
&= \theta_1 \left[(n + 1) (\tilde{P}(y, n + 1) - \tilde{\pi}(0, n + 1)) - y \frac{d}{dy} \tilde{P}(y, n + 1) \right]
\end{aligned}$$

RHS_7

$$\begin{aligned}
&= \theta_2 \sum_{n_2=1}^n (n_2+1) \tilde{\pi}(n_2+1, n+1) y^{n_2} \\
&= \theta_2 \sum_{k=2}^{n+1} k \tilde{\pi}(k, n+1) y^{k-1} \quad (\text{letting } k = n_2 + 1) \\
&= \theta_2 \left[\sum_{k=1}^{n+1} k \tilde{\pi}(k, n+1) y^{k-1} - 1 \cdot \tilde{\pi}(1, n+1) y^0 \right] \\
&= \theta_2 \left[\frac{d}{dy} \tilde{P}(y, n+1) - \tilde{\pi}(1, n+1) \right]
\end{aligned}$$

Collecting:

$$\begin{aligned}
&(\lambda_1 + \lambda_2 + \mu) \left[\tilde{P}(y, n) - \tilde{\pi}(0, n) \right] \\
&+ \gamma_1 \left[n(\tilde{P}(y, n) - \tilde{\pi}(0, n)) - y \frac{d}{dy} \tilde{P}(y, n) \right] \\
&+ \gamma_2 y \frac{d}{dy} \tilde{P}(y, n) \\
&+ \theta_1 \left[n(\tilde{P}(y, n) - \tilde{\pi}(0, n)) - y \frac{d}{dy} \tilde{P}(y, n) \right] \\
&+ \theta_2 y \frac{d}{dy} \tilde{P}(y, n) \\
&= \lambda_1 \left[\tilde{P}(y, n-1) - \tilde{\pi}(n, n-1) y^n - \tilde{\pi}(0, n-1) \right] \\
&+ \lambda_2 y \left[\tilde{P}(y, n-1) - \tilde{\pi}(n, n-1) y^n \right] \\
&+ \mu \left[\tilde{P}(y, n+1) - \tilde{\pi}(0, n+1) + \tilde{\pi}(n+1, n+1) y^n (1-y) \right] \\
&+ \gamma_1 y \left[n \tilde{P}(y, n) - y \frac{d}{dy} \tilde{P}(y, n) \right] \\
&+ \gamma_2 \left[\frac{d}{dy} \tilde{P}(y, n) - \tilde{\pi}(1, n) \right] \\
&+ \theta_1 \left[(n+1) (\tilde{P}(y, n+1) - \tilde{\pi}(0, n+1)) - y \frac{d}{dy} \tilde{P}(y, n+1) \right] \\
&+ \theta_2 \left[\frac{d}{dy} \tilde{P}(y, n+1) - \tilde{\pi}(1, n+1) \right]
\end{aligned}$$

Regrouping:

$$\begin{aligned}
&[-(\gamma_2 + \gamma_1 y)(1-y) + (\theta_2 - \theta_1)y] \frac{d}{dy} \tilde{P}(y, n) + [\lambda_1 + \lambda_2 + \mu + \gamma_1 n(1-y) + \theta_1 n] \tilde{P}(y, n) \\
&= (\lambda_1 + \lambda_2 y) \tilde{P}(y, n-1) + [\mu + \theta_1(n+1)] \tilde{P}(y, n+1) + (\theta_2 - \theta_1 y) \frac{d}{dy} \tilde{P}(y, n+1) \\
&+ \{[\lambda_1 + \lambda_2 + \mu + (\gamma_1 + \theta_1)n] \tilde{\pi}(0, n) - \lambda_1 \tilde{\pi}(0, n-1) - [\mu + \theta_1(n+1)] \tilde{\pi}(0, n+1) \\
&\quad - \gamma_2 \tilde{\pi}(1, n) - \theta_2 \tilde{\pi}(1, n+1)\} \\
&- y^n (\lambda_1 + \lambda_2 y) \tilde{\pi}(n, n-1) + \mu y^n (1-y) \tilde{\pi}(n+1, n+1)
\end{aligned}$$

The y -boundary equation (4.10) makes the bracketed terms vanish; since $\tilde{\pi}(n, n-1) = 0$ (outside $\tilde{S}$):

$$\begin{aligned}
&[-(\gamma_2 + \gamma_1 y)(1-y) + (\theta_2 - \theta_1)y] \frac{d}{dy} \tilde{P}(y, n) + [\lambda_1 + \lambda_2 + \mu + \gamma_1 n(1-y) + \theta_1 n] \tilde{P}(y, n) \\
&= (\lambda_1 + \lambda_2 y) \tilde{P}(y, n-1) + [\mu + \theta_1(n+1)] \tilde{P}(y, n+1) + (\theta_2 - \theta_1 y) \frac{d}{dy} \tilde{P}(y, n+1) \\
&+ \mu y^n (1-y) \tilde{\pi}(n+1, n+1)
\end{aligned}$$

□

5 Analysis of Model-A ($\gamma_1 = \gamma_2 = 0, \theta_1 = \theta_2 = 0$)

Model-A is the baseline of this collection: all jockeying and abandonment parameters are zero, leaving an ordinary two-class $M/M/1$ queue with non-preemptive priority. We analyse it in both state spaces S and $\tilde{S}$: on the former we give two complete proofs, on the latter an approximate solution and its verification.

The balance equations are obtained by plugging $\gamma_1 = \gamma_2 = 0$ and $\theta_1 = \theta_2 = 0$ into Equations (4.3), (4.5), (4.4), (4.6), and (4.7) for state space S , and Equations (4.8), (4.9), (4.10), (4.11), and (4.12) for state space $\tilde{S}$. The corresponding starting points are Lemma 1 for S and Lemma 2 for $\tilde{S}$.

Stability. Since there is no abandonment, the system is positive recurrent if and only if the total traffic intensity satisfies

$$\rho = \frac{\lambda_1 + \lambda_2}{\mu} < 1. \quad (5.1)$$

Figure 4 shows the queueing diagram of Model-A.




Figure 4: Queue layout for Model-A. Only Poisson arrivals and exponential service are present; all jockeying and abandonment rates are zero.

5.1 Analysis of Model-A: state space S

We begin from Equation (4.14) with $\gamma_1 = \gamma_2 = \theta_1 = \theta_2 = 0$, in which the marginal $P_y(y) = P(0, y)$ is still to be determined:

$$[(\lambda_1 + \lambda_2 + \mu)xy - \mu y - \lambda_1 x^2 y - \lambda_2 x y^2] P(x, y) = \mu(x - y)P_y(y) - \mu x(1 - y)\pi(0, 0). \quad (5.2)$$

The following theorem summarises the main result of this subsection; two standalone proofs are given, one analytical and one probabilistic.

Theorem 1. *Consider a queueing system with two priority classes with Poisson arrival rates λ_1 and λ_2 , where class-1 customers have non-preemptive priority over those of class 2, and exponentially distributed service times with rate μ independently of the class. Under the stability condition (5.1), the joint PGF $P(x, y) = \mathbb{E}[x^{N_1} y^{N_2}]$ is*

$$P(x, y) = [(1 + \rho_1 + \rho_2)xy - y - \rho_1 x^2 y - \rho_2 x y^2]^{-1} \cdot y(1 - y) \left[\frac{x - x^*(y)}{x^*(y) - y} \right] \cdot \rho(1 - \rho), \quad (5.3)$$

where

$$x^*(y) = \frac{[\mu + \lambda_1 + \lambda_2(1 - y)] - \sqrt{[\mu + \lambda_1 + \lambda_2(1 - y)]^2 - 4\lambda_1\mu}}{2\lambda_1}. \quad (5.4)$$

Corollary 1. *Under the stability condition (5.1), the empty-state probabilities for Model-A are*

$$\pi_0 = 1 - \rho, \quad \pi(0, 0) = \rho(1 - \rho). \quad (5.5)$$

In particular, $P(1, 1) = 1 - \pi_0 = \rho$ and the diagonal generating function satisfies $P(z, z) = \pi(0, 0)/(1 - \rho z)$.

Proof of Corollary 1. Setting $x = y = z$ in the fundamental equation (5.2), the boundary term $\mu(x - y)P_y(y)$ vanishes:

$$[(\lambda_1 + \lambda_2 + \mu)z^2 - \mu z - \lambda_1 z^3 - \lambda_2 z^3] P(z, z) = -\mu z(1 - z)\pi(0, 0).$$

Dividing both sides by z and factoring the bracket:

$$-(1-z)[\mu - (\lambda_1 + \lambda_2)z]P(z, z) = -\mu(1-z)\pi(0, 0).$$

Cancelling $(1-z) \neq 0$ for $z \in [0, 1)$ and rearranging:

$$P(z, z) = \frac{\mu \pi(0, 0)}{\mu - \lambda z} = \frac{\pi(0, 0)}{1 - \rho z},$$

where $\lambda = \lambda_1 + \lambda_2$ and the second equality follows by dividing numerator and denominator by μ .

Setting $z = 1$ and using the normalisation $P(1, 1) = 1 - \pi_0$:

$$1 - \pi_0 = \frac{\pi(0, 0)}{1 - \rho}. \quad (5.6)$$

The global balance equation for the idle state gives $\lambda \pi_0 = \mu \pi(0, 0)$, i.e. $\pi(0, 0) = \rho \pi_0$. Substituting into (5.6):

$$1 - \pi_0 = \frac{\rho \pi_0}{1 - \rho} \implies (1 - \pi_0)(1 - \rho) = \rho \pi_0 \implies \pi_0 = 1 - \rho,$$

and therefore $\pi(0, 0) = \rho(1 - \rho)$. $\square$

5.1.1 Analytical approach for Model-A in state space S

Proof. We begin from the fundamental equation for Model-A (5.2):

$$[(\lambda_1 + \lambda_2 + \mu)xy - \mu y - \lambda_1 x^2 y - \lambda_2 x y^2] P(x, y) = \mu(x - y)P_y(y) - \mu x(1 - y)\pi(0, 0).$$

We apply the *Kernel Method*: fix $y \in (0, 1]$ and find $x^*(y)$ such that the LHS of the expression above is zero. Since $P(x, y)$ is a PGF, it is bounded. At any such root the LHS vanishes, so the RHS must also vanish:

$$\mu[(x^*(y) - y)P_y(y) - x^*(y)(1 - y)\pi(0, 0)] = 0.$$

To find the roots of the coefficient of $P(x, y)$, divide by $y \neq 0$ and set

$$\begin{aligned} 0 &= (\lambda_1 + \lambda_2 + \mu)x - \mu - \lambda_1 x^2 - \lambda_2 x y \\ 0 &= \lambda_1 x^2 - [\mu + \lambda_1 + \lambda_2(1 - y)]x + \mu, \end{aligned}$$

whose roots are

$$x^\pm(y) = \frac{[\mu + \lambda_1 + \lambda_2(1 - y)] \pm \sqrt{[\mu + \lambda_1 + \lambda_2(1 - y)]^2 - 4\lambda_1\mu}}{2\lambda_1}.$$

We label the root with the negative sign $x^*(y)$ and the root with the positive sign $x^+(y)$. At $y = 1$:

$$x^*(1) = 1 \quad \text{and} \quad x^+(1) = \mu/\lambda_1 > 1.$$

We require our root to lie inside the closed unit disk for $y \in [0, 1)$. Computing the derivative of $x^*(y)$:

$$\frac{d}{dy}x^*(y) = \frac{\lambda_2}{2\lambda_1} \left[-1 + \frac{\mu + \lambda_1 + \lambda_2(1 - y)}{\sqrt{(\mu + \lambda_1 + \lambda_2(1 - y))^2 - 4\lambda_1\mu}} \right] > 0, \quad (5.7)$$

since $(\mu + \lambda_1 + \lambda_2(1 - y))^2 > (\mu + \lambda_1 + \lambda_2(1 - y))^2 - 4\lambda_1\mu$ whenever $\lambda_1\mu > 0$. Therefore $x^*(y)$ is strictly increasing on $[0, 1]$, and since $x^*(1) = 1$ we have $x^*(y) < 1$ for $y \in [0, 1)$. By Vieta's formulas for the kernel polynomial $\lambda_1 x^2 - [\mu + \lambda_1 + \lambda_2(1 - y)]x + \mu = 0$, the product of the two roots satisfies $x^*(y)x^+(y) = \mu/\lambda_1 > 1$, so $x^+(y) = \mu/(\lambda_1 x^*(y)) > 1$ throughout $[0, 1)$. Hence the only root strictly inside the unit disk for $y \in [0, 1)$ is $x^*(y)$:

$$x^*(y) = \frac{(\mu + \lambda_1 + \lambda_2(1 - y)) - \sqrt{(\mu + \lambda_1 + \lambda_2(1 - y))^2 - 4\lambda_1\mu}}{2\lambda_1}. \quad (5.8)$$

Since subsequent steps involve L'Hôpital's rule at $y \rightarrow 1$, we evaluate $\frac{d}{dy}x^*(y)$ from (5.7) at $y = 1$. Under stability, $\rho_1 = \lambda_1/\mu < \rho < 1$ implies $\mu > \lambda_1$, so $\sqrt{(\mu - \lambda_1)^2} = \mu - \lambda_1$:

$$\frac{d}{dy}x^*(1) = \frac{\lambda_2}{2\lambda_1} \left[-1 + \frac{\mu + \lambda_1}{\sqrt{(\mu + \lambda_1)^2 - 4\lambda_1\mu}} \right]$$

$$\begin{aligned}
&= \frac{\lambda_2}{2\lambda_1} \left[-1 + \frac{\mu + \lambda_1}{\sqrt{(\mu - \lambda_1)^2}} \right] \\
&= \frac{\lambda_2}{2\lambda_1} \left[\frac{\mu + \lambda_1 - (\mu - \lambda_1)}{\mu - \lambda_1} \right] \\
&= \frac{\lambda_2}{\mu - \lambda_1}.
\end{aligned}$$

For future reference:

$$\frac{d}{dy}x^*(1) = \frac{\lambda_2}{\mu - \lambda_1} = \frac{\rho_2}{1 - \rho_1}. \quad (5.9)$$

Returning to the fundamental equation (5.2) at $x = x^*(y)$, the RHS also equals zero, giving

$$\mu \cdot [(x^*(y) - y)P_y(y) - x^*(y)(1 - y)\pi(0, 0)] = 0 \implies P_y(y) = \frac{x^*(y)(1 - y)}{x^*(y) - y}\pi(0, 0). \quad (5.10)$$

By Corollary 1, $\pi(0, 0) = \rho(1 - \rho)$. Substituting this and Equation (5.10) into the fundamental equation (5.2):

$$\begin{aligned}
[(\lambda_1 + \lambda_2 + \mu)xy - \mu y - \lambda_1 x^2 y - \lambda_2 x y^2] P(x, y) &= \mu(x - y) \frac{x^*(y)(1 - y)}{x^*(y) - y} \rho(1 - \rho) - \mu x(1 - y) \rho(1 - \rho) \\
[(\lambda_1 + \lambda_2 + \mu)xy - \mu y - \lambda_1 x^2 y - \lambda_2 x y^2] P(x, y) &= \mu(1 - y) \left[\frac{x^*(y)(x - y)}{x^*(y) - y} - x \right] \rho(1 - \rho) \\
[(\lambda_1 + \lambda_2 + \mu)xy - \mu y - \lambda_1 x^2 y - \lambda_2 x y^2] P(x, y) &= \mu(1 - y) \left[\frac{y(x - x^*(y))}{x^*(y) - y} \right] \rho(1 - \rho) \\
[(1 + \rho_1 + \rho_2)xy - y - \rho_1 x^2 y - \rho_2 x y^2] P(x, y) &= y(1 - y) \left[\frac{x - x^*(y)}{x^*(y) - y} \right] \rho(1 - \rho).
\end{aligned}$$

Multiplying both sides by $[(1 + \rho_1 + \rho_2)xy - y - \rho_1 x^2 y - \rho_2 x y^2]^{-1}$ isolates $P(x, y)$ and concludes the proof. $\square$

5.1.2 Probabilistic approach for Model-A in state space S

Proof. This proof mirrors the analytical one, but derives the class-2 marginal PGF $P_y(y)$ by probabilistic arguments. We again begin from the fundamental equation for Model-A (Equation (5.2)):

$$[(\lambda_1 + \lambda_2 + \mu)xy - \mu y - \lambda_1 x^2 y - \lambda_2 x y^2] P(x, y) = \mu[(x - y)P_y(y) - x(1 - y)\pi(0, 0)]. \quad (5.11)$$

Class-2 conditional distribution via $M/G/1$. Since class 1 has non-preemptive priority, class-2 customers are not taken into service whenever $N_1 > 0$; their *effective service time* B is therefore the duration of a complete busy period of the single-class $M/M/1$ sub-queue with arrival rate λ_1 and service rate μ . From class-2's viewpoint the system is thus an $M/G/1$ queue with arrival rate λ_2 and generally distributed service times B of mean $\mathbb{E}[B]$. The LST and mean of B are presented in §3.1 (Equations (3.1)–(3.2)), now with λ_1 instead of λ :

$$\begin{aligned}
\tilde{B}(s) &= \frac{(\mu + \lambda_1 + s) - \sqrt{(\mu + \lambda_1 + s)^2 - 4\mu\lambda_1}}{2\lambda_1}, \\
\mathbb{E}[B] &= \frac{1}{\mu(1 - \rho_1)}.
\end{aligned}$$

Applying the Pollaczek–Khinchine (PK) formula (§3.4, Equation (3.8)) to this $M/G/1$ system, with $\tilde{B}(\lambda_2(1 - y)) = x^*(y)$ (cf. (5.8)) and $\lambda_2 \mathbb{E}[B] = \rho_2/(1 - \rho_1)$:

$$L_2(y) = \frac{(1 - \lambda_2 \mathbb{E}[B])(1 - y) \tilde{B}(\lambda_2(1 - y))}{\tilde{B}(\lambda_2(1 - y)) - y} = \frac{1 - \rho}{1 - \rho_1} \cdot \frac{x^*(y)(1 - y)}{x^*(y) - y}.$$

By the PASTA property, the time-average distribution of N_2 conditional on (busy, $N_1 = 0$) coincides with the $M/G/1$ steady-state distribution, so $\mathbb{E}[y^{N_2} \mid \text{busy}, N_1 = 0] = L_2(y)$.

By the law of total expectation, $P_y(y) = \mathbb{P}(\text{busy}, N_1 = 0) \cdot \mathbb{E}[y^{N_2} \mid \text{busy}, N_1 = 0]$. So

$$P_y(y) = \mathbb{P}(\text{busy}, N_1 = 0) \cdot \mathbb{E}[y^{N_2} \mid \text{busy}, N_1 = 0] = \mathbb{P}(\text{busy}, N_1 = 0) \cdot L_2(y).$$

We now can obtain $\mathbb{P}(\text{busy}, N_1 = 0)$, from $\mathbb{P}(\text{busy}) = \rho$ and $\mathbb{P}(N_1 = 0 \mid \text{busy}) = 1 - \rho_1$. The latter holds because, under non-preemptive priority with a common service rate μ , the class-1 queue-plus-service process is exactly the queue-length process of an autonomous $M/M/1(\lambda_1, \mu)$: a class-1 arrival and a class-1 departure occur at rates λ_1 and μ regardless of the class currently in service, since the residual service time is $\text{Exp}(\mu)$ by memorylessness whether the server holds a class-1 or class-2 customer (cf. §3.2); the class-2 occupancy is thus exchangeable with class-1 occupancy for the purpose of counting class-1 customers, so the count obeys the geometric law and $\mathbb{P}(N_1 = 0 \mid \text{busy}) = 1 - \rho_1$. So

$$\mathbb{P}(\text{busy}, N_1 = 0) = \rho(1 - \rho_1)$$

and therefore

$$P_y(y) = \rho(1 - \rho_1) \cdot L_2(y) = \rho(1 - \rho_1) \cdot \frac{1 - \rho}{1 - \rho_1} \cdot \frac{x^*(y)(1 - y)}{x^*(y) - y} = \frac{x^*(y)(1 - y)}{x^*(y) - y} \cdot (1 - \rho)\rho \quad (5.12)$$

Substituting (5.12) and $\pi(0, 0) = \rho(1 - \rho)$ (Corollary 1) into (5.11):

$$\begin{aligned} [(\lambda_1 + \lambda_2 + \mu)xy - \mu y - \lambda_1 x^2 y - \lambda_2 x y^2] P(x, y) &= \mu(x - y) \frac{x^*(y)(1 - y)}{x^*(y) - y} \rho(1 - \rho) - \mu x(1 - y) \rho(1 - \rho) \\ [(\lambda_1 + \lambda_2 + \mu)xy - \mu y - \lambda_1 x^2 y - \lambda_2 x y^2] P(x, y) &= \mu(1 - y) \left[\frac{x^*(y)(x - y)}{x^*(y) - y} - x \right] \rho(1 - \rho) \\ [(\lambda_1 + \lambda_2 + \mu)xy - \mu y - \lambda_1 x^2 y - \lambda_2 x y^2] P(x, y) &= \mu(1 - y) \left[\frac{y(x - x^*(y))}{x^*(y) - y} \right] \rho(1 - \rho) \\ [(1 + \rho_1 + \rho_2)xy - y - \rho_1 x^2 y - \rho_2 x y^2] P(x, y) &= y(1 - y) \left[\frac{x - x^*(y)}{x^*(y) - y} \right] \rho(1 - \rho). \end{aligned}$$

Multiplying both sides by $[(1 + \rho_1 + \rho_2)xy - y - \rho_1 x^2 y - \rho_2 x y^2]^{-1}$ isolates $P(x, y)$ and concludes the proof. $\square$

5.1.3 Partial-PGF recurrence approach for Model-A in state space $\mathcal{S}$

This proof follows Cohen [Coh56]. We keep the class-1 index n_1 as a discrete parameter and generate only with respect to N_2 , via the PPGF $\tilde{P}(n_1, y) = \sum_{n_2 \geq 0} \pi(n_1, n_2) y^{n_2}$ from (4.13). Multiplying the balance equations by y^{n_2} and summing over n_2 turns them into a second-order linear recurrence in n_1 , whose characteristic roots fix the geometric structure of the solution in the priority-class index.

Proof. We start from the Model-A balance equations, obtained by setting $\gamma_1 = \gamma_2 = \theta_1 = \theta_2 = 0$ in the interior equation (4.3) and the x -boundary equation (4.5):

$$\begin{aligned} (\lambda_1 + \lambda_2 + \mu) \pi(n_1, n_2) &= \mu \pi(n_1 + 1, n_2) + \lambda_1 \pi(n_1 - 1, n_2) + \lambda_2 \pi(n_1, n_2 - 1), & n_1 \geq 1, n_2 \geq 1, \\ (\lambda_1 + \lambda_2 + \mu) \pi(n_1, 0) &= \mu \pi(n_1 + 1, 0) + \lambda_1 \pi(n_1 - 1, 0), & n_1 \geq 1, n_2 = 0. \end{aligned} \quad (5.13)$$

$$\begin{aligned} &(\lambda_1 + \lambda_2 + \mu) \sum_{n_2=1}^{\infty} \pi(n_1, n_2) y^{n_2} \\ &= \mu \sum_{n_2=1}^{\infty} \pi(n_1 + 1, n_2) y^{n_2} + \lambda_1 \sum_{n_2=1}^{\infty} \pi(n_1 - 1, n_2) y^{n_2} + \lambda_2 \sum_{n_2=1}^{\infty} \pi(n_1, n_2 - 1) y^{n_2} \\ LHS_1 &= RHS_1 + RHS_2 + RHS_3 \\ LHS_1 &= (\lambda_1 + \lambda_2 + \mu) \sum_{n_2=1}^{\infty} \pi(n_1, n_2) y^{n_2} \\ &= (\lambda_1 + \lambda_2 + \mu) \left[\sum_{n_2=0}^{\infty} \pi(n_1, n_2) y^{n_2} - \pi(n_1, 0) y^0 \right] \\ &= (\lambda_1 + \lambda_2 + \mu) \left[\tilde{P}(n_1, y) - \pi(n_1, 0) \right] \end{aligned}$$

$$\begin{aligned}
& RHS_1 \\
&= \mu \sum_{n_2=1}^{\infty} \pi(n_1+1, n_2) y^{n_2} \\
&= \mu \left[\sum_{n_2=0}^{\infty} \pi(n_1+1, n_2) y^{n_2} - \pi(n_1+1, 0) y^0 \right] \\
&= \mu \left[\tilde{P}(n_1+1, y) - \pi(n_1+1, 0) \right] \\
& RHS_2 \\
&= \lambda_1 \sum_{n_2=1}^{\infty} \pi(n_1-1, n_2) y^{n_2} \\
&= \lambda_1 \left[\sum_{n_2=0}^{\infty} \pi(n_1-1, n_2) y^{n_2} - \pi(n_1-1, 0) y^0 \right] \\
&= \lambda_1 \left[\tilde{P}(n_1-1, y) - \pi(n_1-1, 0) \right] \\
& RHS_3 \\
&= \lambda_2 \sum_{n_2=1}^{\infty} \pi(n_1, n_2-1) y^{n_2} \\
&= \lambda_2 y \sum_{n_2=1}^{\infty} \pi(n_1, n_2-1) y^{n_2-1} \\
&= \lambda_2 y \sum_{m=0}^{\infty} \pi(n_1, m) y^m \\
&= \lambda_2 y \tilde{P}(n_1, y) \\
& (\lambda_1 + \lambda_2 + \mu) [\tilde{P}(n_1, y) - \pi(n_1, 0)] \\
&= \mu [\tilde{P}(n_1+1, y) - \pi(n_1+1, 0)] \\
&\quad + \lambda_1 [\tilde{P}(n_1-1, y) - \pi(n_1-1, 0)] \\
&\quad + \lambda_2 y \tilde{P}(n_1, y) \\
& [\mu + \lambda_1 + \lambda_2(1-y)] \tilde{P}(n_1, y) - \mu \tilde{P}(n_1+1, y) - \lambda_1 \tilde{P}(n_1-1, y) \\
&= (\lambda_1 + \lambda_2 + \mu) \pi(n_1, 0) - \mu \pi(n_1+1, 0) - \lambda_1 \pi(n_1-1, 0)
\end{aligned} \tag{5.14}$$

By the x -boundary balance equation (5.13), the right-hand side of (5.14) vanishes completely, giving the homogeneous recurrence

$$[\mu + \lambda_1 + \lambda_2(1-y)] \tilde{P}(n_1, y) = \mu \tilde{P}(n_1+1, y) + \lambda_1 \tilde{P}(n_1-1, y), \quad n_1 \geq 1. \tag{5.15}$$

Characteristic roots and mode selection. Equation (5.15) is a *second-order linear recurrence* in n_1 with constant coefficients, so its general solution is a linear combination of two geometric sequences $[\tilde{x}^{\pm}(y)]^{n_1}$, where $\tilde{x}^{\pm}(y)$ solve the characteristic equation

$$\mu \tilde{x}^2 - [\mu + \lambda_1 + \lambda_2(1-y)] \tilde{x} + \lambda_1 = 0, \tag{5.16}$$

that is,

$$\tilde{x}^{\pm}(y) = \frac{(\mu + \lambda_1 + \lambda_2(1-y)) \pm \sqrt{(\mu + \lambda_1 + \lambda_2(1-y))^2 - 4\lambda_1\mu}}{2\mu}.$$

The *characteristic polynomial* (Equation (5.16)) is the reciprocal to that of the bivariate PGF approaches (i.e. the reciprocal of $\lambda_1 x^2 - [\mu + \lambda_1 + \lambda_2(1-y)]x + \mu$), thus its roots are the reciprocals of that kernel's roots: the positive root is the reciprocal of the negative one ($\tilde{x}^+(y) = 1/x^*(y)$), and conversely $\tilde{x}^-(y) = 1/x^+(y)$. By Vieta's formulas, $\tilde{x}^-(y) \tilde{x}^+(y) = \lambda_1/\mu = \rho_1 < 1$. Recall from §5.1.1 that $x^*(y) < 1 < x^+(y)$ for $y \in [0, 1)$. We select the negative-sign root because the sequence $\{\tilde{P}(n_1, y)\}_{n_1 \geq 0}$ must be summable:

$$\sum_{n_1=0}^{\infty} \tilde{P}(n_1, y) = P(1, y) \leq P(1, 1) = 1 - \pi_0 = \rho < \infty.$$

At $y = 1$, we have

$$\tilde{x}^-(1) = \rho_1 \quad \text{and} \quad \tilde{x}^+(1) = 1,$$

the latter making the geometric sequence $[\tilde{x}^+(y)]^{n_1}$ non-summable. We therefore define $\tilde{x}^*(y) := \tilde{x}^-(y)$,

$$\tilde{x}^*(y) := \tilde{x}^-(y) = \frac{(\mu + \lambda_1 + \lambda_2(1-y)) - \sqrt{(\mu + \lambda_1 + \lambda_2(1-y))^2 - 4\lambda_1\mu}}{2\mu} = \rho_1 x^*(y). \quad (5.17)$$

The last identity holds because $\tilde{x}^*$ and x^* share the same numerator and differ only in the denominator (2μ versus $2\lambda_1$). Next,

$$\tilde{P}(n_1, y) = \tilde{P}(0, y) [\tilde{x}^*(y)]^{n_1}, \quad n_1 \geq 0, \quad (5.18)$$

where the relation also holds at $n_1 = 0$ trivially. Recall that at $y = 1$, we have $\tilde{x}^*(1) = \lambda_1/\mu = \rho_1$. Thus, $\tilde{P}(n_1, 1) = \tilde{P}(0, 1) \rho_1^{n_1}$ and $\tilde{P}(0, 1) = P_y(1) = \rho(1 - \rho_1)$. The constant $\tilde{P}(0, y)$ is resolved by its definition

$$\tilde{P}(0, y) = \sum_{n_2=0}^{\infty} \pi(0, n_2) y^{n_2} = P(0, y) = P_y(y),$$

as determined in the bivariate PGF approaches. Summing the geometric law (Equation (5.18)) over $n_1 \geq 0$ for $x \in (0, 1]$ recovers

$$P(x, y) = \sum_{n_1=0}^{\infty} \tilde{P}(n_1, y) x^{n_1} = P_y(y) \sum_{n_1=0}^{\infty} [\tilde{x}^*(y) x]^{n_1} = \frac{P_y(y)}{1 - \tilde{x}^*(y) x}. \quad (5.19)$$

To bring this to the form of Theorem 1, we define the kernel of the fundamental equation, divided by μ :

$$K(x, y) := (1 + \rho_1 + \rho_2)xy - y - \rho_1 x^2 y - \rho_2 x y^2 = -y\rho_1(x - x^*(y))(x - x^+(y)), \quad (5.20)$$

where $x^*(y)$ and $x^+(y)$ are the roots found in §5.1.1. By Vieta's product formula, $x^*(y)x^+(y) = \mu/\lambda_1 = 1/\rho_1$. This, together with $\tilde{x}^*(y)$ (Equation (5.17)), gives $\tilde{x}^*(y) = \rho_1 x^*(y) = 1/x^+(y)$. Using the factorisation of $K(x, y)$ (5.20), the denominator of the reconstructed PGF (5.19) becomes

$$1 - \tilde{x}^*(y) x = 1 - \frac{x}{x^+(y)} = \frac{x^+(y) - x}{x^+(y)} = \frac{K(x, y)}{y\rho_1 x^+(y)(x - x^*(y))}.$$

Substituting this into our reconstructed $P(x, y)$ (Equation (5.19)), using $P_y(y) = \frac{x^*(y)(1-y)}{x^*(y)-y} \rho(1-\rho)$ and also noting that $\rho_1 x^*(y) x^+(y) = 1$, we obtain

$$\begin{aligned} P(x, y) &= \frac{x^*(y)(1-y)}{x^*(y)-y} \rho(1-\rho) \cdot \frac{y\rho_1 x^+(y)(x - x^*(y))}{K(x, y)} \\ &= \underbrace{\rho_1 x^*(y) x^+(y)}_{=1} \frac{1}{K(x, y)} y(1-y) \frac{x - x^*(y)}{x^*(y)-y} \rho(1-\rho) \\ &= [(1 + \rho_1 + \rho_2)xy - y - \rho_1 x^2 y - \rho_2 x y^2]^{-1} \cdot y(1-y) \frac{x - x^*(y)}{x^*(y)-y} \rho(1-\rho), \end{aligned}$$

which recovers Equation (5.3) from Theorem 1 and concludes the proof. $\square$

5.2 Analysis of Model-A on $\tilde{S}$: scope and limitations of the partial-PGF recursion

Setting $\gamma_1 = \gamma_2 = \theta_1 = \theta_2 = 0$ in the $\tilde{S}$ balance equations (4.8)–(4.12) and forming the PPGF $\tilde{P}(y, n) = \sum_{n_2=0}^n \tilde{\pi}(n_2, n) y^{n_2}$ yields the inhomogeneous recurrence:

$$\begin{aligned} &[\lambda_1 + \lambda_2 + \mu] \tilde{P}(y, n) \\ &= (\lambda_1 + \lambda_2 y) \tilde{P}(y, n-1) + \mu \tilde{P}(y, n+1) + \mu y^n (1-y) \tilde{\pi}(n+1, n+1) \end{aligned} \quad (5.21)$$

The forcing term in (5.21) is controlled by the diagonal probabilities $\tilde{\pi}(n+1, n+1) = \pi(0, n+1)$, whose decay we first establish from the exact S -solution.

Lemma 3 (Geometric decay of the boundary probabilities). *For Model-A the boundary probabilities $\pi(0, n)$ satisfy $\pi(0, n) = O(r^{-n})$ for every $r < 1/\rho$; equivalently they decay geometrically with rate ρ .*

Proof. By definition $P_y(y) = P(0, y) = \sum_{n \geq 0} \pi(0, n) y^n$ is the generating function of $\{\pi(0, n)\}$, and from (5.10) with $\pi(0, 0) = \rho(1 - \rho)$,

$$P_y(y) = \rho(1 - \rho) \frac{x^*(y)(1 - y)}{x^*(y) - y}.$$

Evaluating the kernel quadratic (5.4) at $x = y$ gives $\lambda_1 y^2 - [\mu + \lambda_1 + \lambda_2(1 - y)]y + \mu = (\lambda y - \mu)(y - 1)$ with $\lambda = \lambda_1 + \lambda_2$, so $x^*(y) = y$ holds only at $y = 1$ and $y = \mu/\lambda = 1/\rho$. At $y = 1$ the simple zeros of $(1 - y)$ and $(x^*(y) - y)$ cancel and P_y is analytic with $P_y(1) = \rho$; the branch point of x^* lies beyond $1/\rho$. Hence the singularity of P_y nearest the origin is the simple pole at $y = 1/\rho > 1$, so P_y is analytic on the disc $|y| < 1/\rho$. By the Cauchy coefficient estimates, for every $r < 1/\rho$ there is $C_r < \infty$ with $\pi(0, n) \leq C_r r^{-n}$, i.e. geometric decay at rate ρ . $\square$

Approximation 1 (Heuristic PPGF on $\tilde{S}$). *Neglecting the boundary-diagonal forcing term $\mu y^n(1 - y)\tilde{\pi}(n + 1, n + 1)$ in (5.21) — which is exponentially small by Lemma 3 — the homogeneous recurrence admits the geometric solution*

$$\tilde{P}(y, n) \approx \rho(1 - \rho) [\tilde{y}^*(y)]^n,$$

where

$$\tilde{y}^*(y) = \frac{(\lambda + \mu) - \sqrt{(\lambda + \mu)^2 - 4\mu(\lambda_1 + \lambda_2 y)}}{2\mu}, \quad \lambda = \lambda_1 + \lambda_2.$$

At $y = 1$ this recovers $\tilde{P}(1, n) = (1 - \rho)\rho^{n+1}$, consistent with the $M/M/1(\lambda, \mu)$ steady-state marginal for n customers waiting.

Justification of Approximation 1. The inhomogeneous term $\mu y^n(1 - y)\tilde{\pi}(n + 1, n + 1)$ prevents the direct application of a geometric-sequence ansatz. The diagonal probabilities $\tilde{\pi}(n, n) = \pi(0, n)$ correspond to states where all n waiting customers are class 2, and by Lemma 3 they decay geometrically in n at rate ρ ; the forcing term is therefore exponentially small, and neglecting it yields an approximation whose accuracy is quantified below. The resulting homogeneous recurrence is

$$(\lambda_1 + \lambda_2 + \mu)\tilde{P}(y, n) = (\lambda_1 + \lambda_2 y)\tilde{P}(y, n - 1) + \mu\tilde{P}(y, n + 1).$$

The characteristic polynomial of this recurrence is

$$\mu[\tilde{y}(y)]^2 - (\lambda_1 + \lambda_2 + \mu)\tilde{y}(y) + (\lambda_1 + \lambda_2 y) = 0,$$

with roots

$$\tilde{y}^\pm(y) = \frac{(\lambda_1 + \lambda_2 + \mu) \pm \sqrt{(\lambda_1 + \lambda_2 + \mu)^2 - 4\mu(\lambda_1 + \lambda_2 y)}}{2\mu}.$$

At $y = 1$ the roots are $\tilde{y}_+(1) = 1$ and $\tilde{y}_-(1) = \rho$. The root $\tilde{y}_+ = 1$ would produce a non-summable geometric series as $n \rightarrow \infty$, so only the negative-sign root is admissible:

$$\tilde{y}^*(y) = \frac{(\lambda_1 + \lambda_2 + \mu) - \sqrt{(\lambda_1 + \lambda_2 + \mu)^2 - 4\mu(\lambda_1 + \lambda_2 y)}}{2\mu}.$$

The geometric ansatz gives $\tilde{P}(y, n) = \tilde{P}(y, 0) \cdot [\tilde{y}^*(y)]^n$. Using $\tilde{P}(y, 0) = \tilde{\pi}(0, 0) = \pi(0, 0) = \rho(1 - \rho)$ (Corollary 1):

$$\tilde{P}(y, n) = \rho(1 - \rho) \cdot [\tilde{y}^*(y)]^n.$$

Since n counts customers *waiting* (the customer in service is not counted), $n + 1$ is the total number in the system, so evaluating at $y = 1$:

$$\tilde{P}(1, n) = \rho(1 - \rho) \cdot \rho^n = (1 - \rho)\rho^{n+1}, \quad \square$$

which matches the $M/M/1(\lambda, \mu)$ stationary distribution for $n + 1$ customers in the system [AR02].

Failure mechanism. The neglected forcing $\mu y^n(1-y)\tilde{\pi}(n+1, n+1)$ feeds in the diagonal states, in which all waiting customers are class 2. Using $\tilde{\pi}(n, n) = \pi(0, n)$ and the rate proved in Lemma 3, this term decays in n like ρ^n (modulated by y^n), whereas the homogeneous solution decays like $[\tilde{y}^*(y)]^n$. The approximation is therefore accurate exactly when the forcing decays faster, i.e. $\rho < \tilde{y}^*(y)$ on the relevant range of y ; this holds when the priority class carries most of the load (ρ_1/ρ large, ρ moderate), precisely the empirical validity region mapped in Figure 22.

The obstruction is structural and recurs throughout this thesis: an unknown boundary/diagonal trace — here the sequence $\pi(0, n)$ — couples into an otherwise solvable recursion, exactly as the unknown diagonal $P(z, z)$ enters the Volterra forcing of Models B and X (§6 and Remark 1). A rigorous closure would retain the forcing and solve (5.21) by variation of constants in n , which reintroduces the very diagonal sequence $\{\pi(0, n)\}$ as an unknown to be determined self-consistently; see §13.5.

5.3 Mean queue lengths and mean waiting times

5.3.1 Class-1 marginal PGF and mean queue length

Setting $y = 1$ in the fundamental equation (5.2) gives

$$[-\lambda_1(x-1)(x-1/\rho_1)]P(x, 1) = \mu(x-1)P_y(1).$$

Cancelling the common factor $(x-1)$ for $x \neq 1$ and rearranging:

$$P(x, 1) = \frac{\mu P_y(1)}{\mu - \lambda_1 x}.$$

The constant $P_y(1)$ is pinned by the normalisation $P(1, 1) = 1 - \pi_0 = \rho$:

$$\rho = \frac{\mu P_y(1)}{\mu - \lambda_1} \implies P_y(1) = \rho(1 - \rho_1),$$

so that

$$P(x, 1) = \frac{\rho(1 - \rho_1)}{1 - \rho_1 x}.$$

Differentiating by the quotient rule and evaluating at $x = 1$:

$$\mathbb{E}[N_1] = \left. \frac{d}{dx} P(x, 1) \right|_{x=1} = \left. \frac{\rho(1 - \rho_1) \cdot \rho_1}{(1 - \rho_1 x)^2} \right|_{x=1} = \frac{\rho(1 - \rho_1) \rho_1}{(1 - \rho_1)^2} = \frac{\rho \rho_1}{1 - \rho_1}. \quad (5.22)$$

5.3.2 Total queue length from the diagonal

Setting $x = y = z$ in the kernel polynomial $\lambda_1 x^2 - (\lambda_1 + \lambda_2 + \mu)x + \mu$ shows that $x^*(z) = z$ iff $z \in \{1, 1/\rho\}$; hence $x^*(z) \neq z$ for all $z \in (0, 1)$ with $\rho < 1$. Therefore the ratio $(z - x^*(z))/(x^*(z) - z) = -1$ throughout $(0, 1)$, and the formula of Theorem 1 reduces to:

$$P(z, z) = \frac{z(1-z)\rho(1-\rho)}{D(z, z)} \cdot (-1).$$

The denominator at $x = y = z$ is $D(z, z) = (1 + \rho)z^2 - z - \rho z^3 = -\rho z(z-1)(z-1/\rho)$, so:

$$\begin{aligned} P(z, z) &= \frac{-z(1-z)\rho(1-\rho)}{-\rho z(z-1)(z-1/\rho)} \\ &= \frac{(1-z)(1-\rho)}{(z-1)(z-1/\rho)} = \frac{-(1-\rho)}{z-1/\rho} = \frac{\rho(1-\rho)}{1-\rho z}, \end{aligned} \quad (5.23)$$

confirming the diagonal identity of Corollary 1. Since $\left. \frac{d}{dz} P(z, z) \right|_{z=1} = \frac{\partial P}{\partial x}(1, 1) + \frac{\partial P}{\partial y}(1, 1) = \mathbb{E}[N_1] + \mathbb{E}[N_2]$, differentiating by the quotient rule and evaluating at $z = 1$ gives:

$$\mathbb{E}[N_1] + \mathbb{E}[N_2] = \left. \frac{d}{dz} \frac{\rho(1-\rho)}{1-\rho z} \right|_{z=1} = \left. \frac{\rho(1-\rho) \cdot \rho}{(1-\rho z)^2} \right|_{z=1} = \frac{\rho^2(1-\rho)}{(1-\rho)^2} = \frac{\rho^2}{1-\rho}.$$

This is the $M/M/1(\lambda, \mu)$ queue-length formula for the combined system [AR02], as expected: the two classes together behave as a single $M/M/1$ queue.

5.3.3 Class-2 mean queue length by subtraction

$$\begin{aligned}
\mathbb{E}[N_2] &= (\mathbb{E}[N_1] + \mathbb{E}[N_2]) - \mathbb{E}[N_1] = \frac{\rho^2}{1-\rho} - \frac{\rho\rho_1}{1-\rho_1} \\
&= \frac{\rho^2(1-\rho_1) - \rho\rho_1(1-\rho)}{(1-\rho)(1-\rho_1)} = \frac{\rho[\rho(1-\rho_1) - \rho_1(1-\rho)]}{(1-\rho)(1-\rho_1)} \\
&= \frac{\rho(\rho - \rho_1)}{(1-\rho)(1-\rho_1)} = \frac{\rho\rho_2}{(1-\rho_1)(1-\rho)}.
\end{aligned} \tag{5.24}$$

5.3.4 Mean waiting times via Little's Law

Little's Law applied to the *waiting sub-system* of each class: since N_i counts only customers *waiting* in queue (excluding the customer in service), Little's Law gives $\mathbb{E}[N_i] = \lambda_i \mathbb{E}[W_i]$, where $\mathbb{E}[W_i]$ is the mean waiting time in queue (time from arrival until service start, not including the service duration):

$$\mathbb{E}[W_1] = \frac{\mathbb{E}[N_1]}{\lambda_1} = \frac{\rho\rho_1}{\lambda_1(1-\rho_1)} = \frac{\rho \cdot (\lambda_1/\mu)}{\lambda_1(1-\rho_1)} = \frac{\rho}{\mu(1-\rho_1)}, \tag{5.25}$$

$$\mathbb{E}[W_2] = \frac{\mathbb{E}[N_2]}{\lambda_2} = \frac{\rho\rho_2}{\lambda_2(1-\rho_1)(1-\rho)} = \frac{\rho \cdot (\lambda_2/\mu)}{\lambda_2(1-\rho_1)(1-\rho)} = \frac{\rho}{\mu(1-\rho_1)(1-\rho)}. \tag{5.26}$$

Class-1 benefits from priority: $\mathbb{E}[W_2]$ carries the additional penalty factor $1/(1-\rho)$ relative to $\mathbb{E}[W_1]$, reflecting the delay that class-2 customers incur while class-1 customers in the queue are served first. Note that both $\mathbb{E}[W_1]$ and $\mathbb{E}[W_2]$ depend on ρ_2 through the total load $\rho = \rho_1 + \rho_2$; class-2 affects class-1 waiting times only via server utilisation, not through a direct queueing penalty.

6 Analysis of Model-B ($\gamma_1, \gamma_2 > 0, \theta_1 = \theta_2 = 0$)

Model-B introduces bidirectional jockeying ($\gamma_1, \gamma_2 > 0$) with no abandonment ($\theta_1 = \theta_2 = 0$). Jockeying conserves the total number of customers, so the stationary distribution exists under the same condition as Model-A. But it breaks the per-class Poisson/renewal structure: a class-2 customer can no longer identify its effective service time with a class-1 busy period, so the Pollaczek–Khinchine route is unavailable.

The method of characteristics yields a general solution of the fundamental equation, but closing it — determining the unknown boundary function $P_y(y)$ — requires solving a Volterra-type integral equation with no closed form. Model-B is therefore left open, in structural parallel with Model-X; its role is to motivate the tractable specialisation Model-B₂.

Stability. $\rho = (\lambda_1 + \lambda_2)/\mu < 1$. Jockeying conserves the total customer count and does not affect this condition. Figure 5 shows the queue layout.




Figure 5: Queue layout for Model-B. Dashed arrows indicate bidirectional jockeying: γ_1 moves a class-1 customer down to Queue 2 and γ_2 moves a class-2 customer up to Queue 1. Both transfers conserve the total n .

Corollary 2. *The empty-system probability π_0 (server idle, no one waiting) and the both-queues-empty probability $\pi(0,0)$ (server busy, $N_1 = N_2 = 0$) of Model-B coincide with the Model-A baseline,*

$$\pi_0 = 1 - \rho, \quad \pi(0,0) = \rho(1 - \rho), \tag{6.1}$$

and the diagonal of the PGF is $P(z, z) = \pi(0, 0)/(1 - \rho z)$.

Proof. Setting $x = y$ in the fundamental equation (4.14) with $\theta_1 = \theta_2 = 0$, both jockeying terms carry the factor $(x - y)$ and vanish, as does the boundary term $\mu(x - y)P_y(y)$; dividing the surviving terms by the common factor x leaves

$$[(\lambda_1 + \lambda_2 + \mu)x - \mu - (\lambda_1 + \lambda_2)x^2]P(x, x) = -\mu(1 - x)\pi(0, 0).$$

Writing $\lambda = \lambda_1 + \lambda_2$, the bracket factors,

$$(\lambda + \mu)x - \mu - \lambda x^2 = -(\lambda x^2 - (\lambda + \mu)x + \mu) = -(\lambda x - \mu)(x - 1) = (\mu - \lambda x)(x - 1);$$

cancelling $(x - 1)$ against $-(1 - x) = (x - 1)$ on the right gives

$$P(x, x) = \frac{\mu \pi(0, 0)}{\mu - \lambda x} = \frac{\pi(0, 0)}{1 - \rho x}. \quad (6.2)$$

The idle state (0) exchanges probability flux only with the both-queues-empty busy state, so the global balance across the idle-busy cut reads $\lambda\pi_0 = \mu\pi(0, 0)$ (an arrival at total rate λ leaves (0); a service completion at rate μ returns to it), whence $\pi(0, 0) = \rho\pi_0$. Evaluating (6.2) at $x = 1$ with $P(1, 1) = 1 - \pi_0$ gives $1 - \pi_0 = \rho\pi_0/(1 - \rho)$, so $(1 - \pi_0)(1 - \rho) = \rho\pi_0$, i.e. $\pi_0 = 1 - \rho$ and $\pi(0, 0) = \rho(1 - \rho)$. This matches Model-A (5.5): jockeying conserves the total-customer process and cannot change the marginal idle probability. The same identities therefore hold in the other no-abandonment model, Model-B₂; Model-C₂ retains the universal idle balance $\pi(0, 0) = \rho\pi_0$ (4.2) but, with abandonment present, attains $\pi_0 = 1 - \rho$ only in the limit $\theta_1 \rightarrow 0^+$. $\square$

Claim 1. *Under stability, $\rho < 1$, the bivariate PGF of Model-B satisfies, on each characteristic curve $\gamma_2 x + \gamma_1 y = C_1$ of the fundamental equation, the integral representation*

$$P(x, y(x)) = P_h(x, y(x)) \int_0^x \frac{h(t; C_1)}{P_h(t, y(t))} dt, \quad (6.3)$$

where P_h is the homogeneous solution (6.9), $y(t) = (C_1 - \gamma_2 t)/\gamma_1$ is the characteristic parametrisation, and $h(t; C_1)$ is an explicit forcing function depending on the unknown boundary trace $P_y(y) = P(0, y)$. Evaluating (6.3) along the diagonal $x = y = z$ and using the conservation identity $P(z, z) = \pi(0, 0)/(1 - \rho z)$ reduces the problem to a Volterra integral equation of the first kind for P_y :

$$\mathcal{F}(z) = \int_0^z \mathcal{K}(z, t) P_y(y(t; z)) dt, \quad (6.4)$$

where $\mathcal{K}$ and $\mathcal{F}$ are fully explicit (defined in (6.20)–(6.21) below). The kernel $\mathcal{K}$ does not factorise for general (γ_1, γ_2) , so (6.4) does not reduce to a closed-form solution for P_y . Model-B is therefore left open.

Proof of Claim 1. The proof proceeds in five steps: (i) homogeneous solution via characteristics, (ii) rigorous sign analysis of the exponent E , (iii) particular solution via variation of parameters, (iv) pinning the free function from the boundary condition, and (v) diagonal evaluation yielding the Volterra equation.

Step 1: Homogeneous solution via the method of characteristics.

The fundamental equation (4.14) with $\theta_1 = \theta_2 = 0$ is a first-order linear PDE in $P(x, y)$. Setting the right-hand side to zero and dividing through by y yields the homogeneous PDE

$$\gamma_1 x(x - y) \frac{\partial P}{\partial x} - \gamma_2 x(x - y) \frac{\partial P}{\partial y} = -[(\lambda_1 + \lambda_2 + \mu)x - \mu - \lambda_1 x^2 - \lambda_2 xy]P.$$

Following the method of characteristics (§3.6.3), we associate to this PDE the characteristic system

$$\frac{dx}{\gamma_1 x(x - y)} = \frac{dy}{-\gamma_2 x(x - y)} = \frac{dP}{-[(\lambda_1 + \lambda_2 + \mu)x - \mu - \lambda_1 x^2 - \lambda_2 xy]P}.$$

Taking the ratio of the first two differentials, the $x(x - y)$ factors cancel, giving $dx/\gamma_1 = dy/(-\gamma_2)$. Integrating yields the first integral

$$C_1 = \gamma_2 x + \gamma_1 y \quad (\text{constant along each characteristic}).$$

The characteristics are therefore the straight lines of slope $-\gamma_2/\gamma_1$ in the (x, y) -plane. Fix one such line and parametrise y by $y(x) = (C_1 - \gamma_2 x)/\gamma_1$. To extract the equation for P along this curve we need the coefficient $A(x, y)$ and the denominator $x(x - y)$ after substitution. Setting

$$A(x, y) := (\lambda_1 + \lambda_2 + \mu)x - \mu - \lambda_1 x^2 - \lambda_2 xy, \quad (6.5)$$

and substituting $y = (C_1 - \gamma_2 x)/\gamma_1$ so that $\lambda_2 xy = (\lambda_2 C_1/\gamma_1)x - (\lambda_2 \gamma_2/\gamma_1)x^2$, we obtain, collecting by powers of x ,

$$\begin{aligned} A(x, y(x)) &= (\lambda_1 + \lambda_2 + \mu)x - \mu - \lambda_1 x^2 - \frac{\lambda_2 C_1}{\gamma_1}x + \frac{\lambda_2 \gamma_2}{\gamma_1}x^2 \\ &= -\frac{\gamma_1 \lambda_1 - \lambda_2 \gamma_2}{\gamma_1}x^2 + \frac{\gamma_1(\lambda_1 + \lambda_2 + \mu) - \lambda_2 C_1}{\gamma_1}x - \mu. \end{aligned}$$

Similarly, $x - y(x) = [(\gamma_1 + \gamma_2)x - C_1]/\gamma_1$, so $\gamma_1 x(x - y(x)) = x[(\gamma_1 + \gamma_2)x - C_1]$. The dP/P equation from the characteristic system therefore becomes

$$\frac{dP}{P} = \frac{-A(x, y(x))}{\gamma_1 x(x - y(x))} dx = \frac{(\gamma_1 \lambda_1 - \lambda_2 \gamma_2)x^2 - [\gamma_1(\lambda_1 + \lambda_2 + \mu) - \lambda_2 C_1]x + \gamma_1 \mu}{\gamma_1 x[(\gamma_1 + \gamma_2)x - C_1]} dx.$$

To integrate this, we decompose the integrand via partial fractions. Writing the integrand as $(1/\gamma_1)I(x)dx$ with $I(x) = K + M/x + N(C_1)/[(\gamma_1 + \gamma_2)x - C_1]$, the partial fraction identity requires

$$K(\gamma_1 + \gamma_2)x^2 + [M(\gamma_1 + \gamma_2) + N - KC_1]x - MC_1 = (\gamma_1 \lambda_1 - \lambda_2 \gamma_2)x^2 - [\gamma_1(\lambda_1 + \lambda_2 + \mu) - \lambda_2 C_1]x + \gamma_1 \mu.$$

Matching the x^2 , x^0 , and x^1 coefficients in turn:

$$K = \frac{\gamma_1 \lambda_1 - \lambda_2 \gamma_2}{\gamma_1 + \gamma_2}, \quad M = -\frac{\gamma_1 \mu}{C_1}, \quad N(C_1) = \gamma_1 \left[\frac{\lambda_1 + \lambda_2}{\gamma_1 + \gamma_2} C_1 - (\lambda_1 + \lambda_2 + \mu) + \frac{\mu(\gamma_1 + \gamma_2)}{C_1} \right]. \quad (6.6)$$

For N , the x^1 match gives $N = -\gamma_1(\lambda_1 + \lambda_2 + \mu) + \lambda_2 C_1 + KC_1 - M(\gamma_1 + \gamma_2)$; substituting K and M , writing $\lambda = \lambda_1 + \lambda_2$, $\Gamma = \gamma_1 + \gamma_2$, and collecting the two C_1 -proportional terms over the common denominator Γ ,

$$\begin{aligned} N &= -\gamma_1(\lambda + \mu) + C_1 \frac{\lambda_2 \Gamma + \gamma_1 \lambda_1 - \lambda_2 \gamma_2}{\Gamma} + \frac{\gamma_1 \mu \Gamma}{C_1} \\ &= -\gamma_1(\lambda + \mu) + \frac{\gamma_1 \lambda}{\Gamma} C_1 + \frac{\gamma_1 \mu \Gamma}{C_1}, \end{aligned}$$

using $\lambda_2 \Gamma + \gamma_1 \lambda_1 - \lambda_2 \gamma_2 = \lambda_2 \gamma_1 + \gamma_1 \lambda_1 = \gamma_1 \lambda$ in the middle term; factoring out γ_1 yields (6.6). Integrating $dP/P = (1/\gamma_1)I(x)dx$ term by term gives

$$\ln P = \frac{K}{\gamma_1}x - \frac{\mu}{C_1} \ln x + \frac{N(C_1)}{\gamma_1(\gamma_1 + \gamma_2)} \ln[(\gamma_1 + \gamma_2)x - C_1] + \ln f(C_1), \quad (6.7)$$

where $f(C_1)$ is an arbitrary function of the characteristic constant. Now observe that

$$(\gamma_1 + \gamma_2)x - C_1 = (\gamma_1 + \gamma_2)x - (\gamma_2 x + \gamma_1 y) = \gamma_1(x - y),$$

so $\ln[(\gamma_1 + \gamma_2)x - C_1] = \ln \gamma_1 + \ln(x - y)$. Exponentiating (6.7) term by term,

$$P = f(C_1) \exp\left(\frac{K}{\gamma_1}x\right) \cdot x^{-\mu/C_1} \cdot \gamma_1^{N(C_1)/(\gamma_1 \Gamma)} (x - y)^{N(C_1)/(\gamma_1 \Gamma)};$$

the factor $\gamma_1^{N(C_1)/(\gamma_1 \Gamma)}$ depends on (x, y) only through C_1 and is absorbed, together with f , into a redefined arbitrary function $F(C_1)$. Reading off the exponent of $(x - y)$:

$$E(x, y) := \frac{N(C_1)}{\gamma_1(\gamma_1 + \gamma_2)} \Big|_{C_1 = \gamma_2 x + \gamma_1 y} = \frac{\lambda_1 + \lambda_2}{(\gamma_1 + \gamma_2)^2} (\gamma_2 x + \gamma_1 y) - \frac{\lambda_1 + \lambda_2 + \mu}{\gamma_1 + \gamma_2} + \frac{\mu}{\gamma_2 x + \gamma_1 y}. \quad (6.8)$$

Exponentiating (6.7) yields the general homogeneous solution

$$P_h(x, y) = F(\gamma_2 x + \gamma_1 y) \cdot \exp\left(\frac{\gamma_1 \lambda_1 - \lambda_2 \gamma_2}{\gamma_1(\gamma_1 + \gamma_2)}x\right) \cdot x^{-\mu/(\gamma_2 x + \gamma_1 y)} \cdot (x - y)^{E(x, y)}, \quad (6.9)$$

where F is an arbitrary differentiable function to be determined.

Step 2: The exponent $E(x, y)$ is strictly positive on $(0, 1)^2$.

Recall $\lambda = \lambda_1 + \lambda_2$ and $\Gamma = \gamma_1 + \gamma_2$. Along any fixed characteristic, $C_1 = \gamma_2 x + \gamma_1 y$ is constant, so the exponent (6.8) depends only on $u := C_1 \in (0, \Gamma)$:

$$E(u) = \frac{\lambda}{\Gamma^2} u - \frac{\lambda + \mu}{\Gamma} + \frac{\mu}{u}.$$

(For $(x, y) \in (0, 1)^2$ we have $u = \gamma_2 x + \gamma_1 y \in (0, \Gamma)$; $E(u)$ is continuous there, with $E(u) \rightarrow +\infty$ as $u \rightarrow 0^+$ and $E(u) \rightarrow 0^+$ as $u \rightarrow \Gamma^-$, so it is finite at every interior point although not uniformly bounded.) Multiplying through by $u\Gamma^2 > 0$ gives the equivalent quadratic

$$q(u) := \lambda u^2 - (\lambda + \mu)\Gamma u + \mu\Gamma^2.$$

The sign of E on $(0, \Gamma)$ equals the sign of $q(u)$ there. We factor q by finding its roots. The discriminant is $\Delta = (\lambda + \mu)^2\Gamma^2 - 4\lambda\mu\Gamma^2 = \Gamma^2(\mu - \lambda)^2$, so (using $\rho < 1 \Rightarrow \mu > \lambda \Rightarrow \mu - \lambda > 0$)

$$u_{1,2} = \frac{(\lambda + \mu)\Gamma \pm \Gamma(\mu - \lambda)}{2\lambda} \implies u_1 = \frac{2\lambda\Gamma}{2\lambda} = \Gamma, \quad u_2 = \frac{2\mu\Gamma}{2\lambda} = \frac{\Gamma}{\rho} > \Gamma.$$

Hence

$$q(u) = \lambda(u - \Gamma)\left(u - \frac{\Gamma}{\rho}\right).$$

For every $u \in (0, \Gamma)$ both factors are strictly negative ($(u - \Gamma) < 0$ and $(u - \Gamma/\rho) < 0$ since $\Gamma/\rho > \Gamma$), so $q(u) = \lambda \cdot (\text{neg}) \cdot (\text{neg}) > 0$, and therefore

$$E(x, y) > 0 \quad \text{for all } (x, y) \in (0, 1)^2.$$

Note that $E \rightarrow 0$ as $u \rightarrow \Gamma^-$ (i.e. as $(x, y) \rightarrow (1, 1)$), while E is bounded and bounded away from zero in any compact subset of $(0, 1)^2$.

Sign convention and diagonal limit. Because $E > 0$, the factor $(x - y)^E$ in P_h carries the complex phase $e^{i\pi E}$ whenever $x < y$. This phase is absorbed into the arbitrary function $F(C_1)$, and the ratio $P_h(x, y(x))/P_h(t, y(t))$ that appears in the integral representation is always real and positive: both numerator and denominator lie on the same characteristic (same C_1 , same constant value of E), so the phases cancel exactly. Crucially, since $E > 0$ the factor $(y(t; z) - t)^E \rightarrow 0$ as $t \rightarrow z^-$, meaning P_h formally vanishes on the diagonal $x = y$. The diagonal equation that appears in Step 5 below is therefore understood as the $\lim_{x \rightarrow z^-}$ of (6.14); the limit-by-limit analysis carried out there, using the exact linear relation $y(t; z) - t = (\Gamma/\gamma_1)(z - t)$, confirms that the $0 \cdot \infty$ limit is finite and equals $\pi(0, 0)/(1 - \rho z)$.

Step 3: Particular solution via variation of parameters.

We now work with the full (inhomogeneous) fundamental equation restricted to a fixed characteristic. Along the curve $y(x) = (C_1 - \gamma_2 x)/\gamma_1$, the total derivative satisfies $\frac{d}{dx}P(x, y(x)) = \partial_x P + \frac{dy}{dx}\partial_y P = \partial_x P - \frac{\gamma_2}{\gamma_1}\partial_y P$, so the jockeying terms in (4.14) become

$$\gamma_1 x y(x - y) \partial_x P + \gamma_2 x y(y - x) \partial_y P = \gamma_1 x y(x - y) \left(\partial_x P - \frac{\gamma_2}{\gamma_1} \partial_y P \right) = \gamma_1 x y(x - y) \frac{d}{dx} P(x, y(x)).$$

The fundamental equation along the characteristic therefore reads, with its algebraic coefficient factorised as $y A(x, y)$ (cf. (6.5)),

$$y A(x, y) P + \gamma_1 x y(x - y) \frac{dP}{dx} = G(x, y),$$

and dividing through by $\gamma_1 x y(x - y)$ reduces the PDE to the first-order linear ODE (for fixed C_1)

$$\frac{dP}{dx} = f(x; C_1) P + h(x; C_1), \tag{6.10}$$

where the coefficient f and the forcing term h are

$$\begin{aligned} f(x; C_1) &= -\frac{A(x, y(x))}{\gamma_1 x(x - y(x))}, \\ h(x; C_1) &= \frac{G(x, y(x))}{\gamma_1 x y(x)(x - y(x))}, \end{aligned} \tag{6.11}$$

with A as in (6.5) and

$$G(x, y) := \mu(x - y) P_y(y) - \mu x(1 - y) \pi(0, 0). \quad (6.12)$$

Since $P_h(x, y(x))$ satisfies the homogeneous equation $\frac{d}{dx} P_h = f(x; C_1) P_h$, it is the fundamental solution of ODE (6.10). Applying variation of parameters (§3.6.2), the general solution of the inhomogeneous ODE is

$$P(x, y(x)) = P_h(x, y(x)) \left[F(C_1) + \int_{x_0}^x \frac{h(t; C_1)}{P_h(t, y(t))} dt \right]. \quad (6.13)$$

Step 4: Pinning the free function from the boundary condition at $x = 0$.

Along any characteristic with constant $C_1 > 0$, the endpoint at $x = 0$ is $y(0) = C_1/\gamma_1$, giving

$$P\left(0, \frac{C_1}{\gamma_1}\right) = P_y\left(\frac{C_1}{\gamma_1}\right),$$

which is finite. However, P_h contains the factor $x^{-\mu/C_1} \rightarrow +\infty$ as $x \rightarrow 0^+$ (since $\mu/C_1 > 0$ for $C_1 > 0$), while the integral in (6.13) vanishes as $x_0 \rightarrow 0$ (the integration interval collapses). Therefore

$$\lim_{x \rightarrow 0^+} P(x, y(x)) = \lim_{x \rightarrow 0^+} P_h(x, y(x)) \cdot F(C_1) = +\infty \cdot F(C_1),$$

which can only be finite if $F(C_1) = 0$. Setting $x_0 = 0$ and $F(C_1) = 0$ in (6.13) gives the integral representation of Claim 1:

$$P(x, y(x)) = P_h(x, y(x)) \int_0^x \frac{h(t; C_1)}{P_h(t, y(t))} dt. \quad (6.14)$$

Step 5: Reduction to the Volterra equation.

By Corollary 2 the empty-state probabilities are $\pi_0 = 1 - \rho$ and $\pi(0, 0) = \rho(1 - \rho)$, with diagonal value $P(z, z) = \pi(0, 0)/(1 - \rho z)$ by (6.2); these are the inputs to the diagonal evaluation below.

We simplify the forcing function h on the characteristic. Substituting $G(t, y(t))$ from (6.12) into (6.11):

$$\begin{aligned} h(t; C_1) &= \frac{\mu(t - y(t)) P_y(y(t)) - \mu t(1 - y(t)) \pi(0, 0)}{\gamma_1 t y(t) (t - y(t))} \\ &= \frac{\mu P_y(y(t))}{\gamma_1 t y(t)} - \frac{\mu(1 - y(t)) \pi(0, 0)}{\gamma_1 y(t) (t - y(t))}. \end{aligned} \quad (6.15)$$

We will evaluate (6.14) along the diagonal characteristic, which passes through (z, z) and has $C_1 = (\gamma_1 + \gamma_2)z$. The parametrisation along that curve is $y(t; z) = ((\gamma_1 + \gamma_2)z - \gamma_2 t)/\gamma_1$, so $y(0; z) = (\gamma_1 + \gamma_2)z/\gamma_1 > z$ and $y(z; z) = z$. Therefore $y(t; z) > t$ for all $t \in [0, z]$, which means $t - y(t; z) < 0$ throughout the integration interval. Writing $t - y(t; z) = -(y(t; z) - t)$ in the denominator of (6.15):

$$h(t; (\gamma_1 + \gamma_2)z) = \frac{\mu}{\gamma_1 t y(t; z)} P_y(y(t; z)) + \frac{\mu(1 - y(t; z))}{\gamma_1 y(t; z) (y(t; z) - t)} \pi(0, 0). \quad (6.16)$$

As established in Step 2, $E > 0$, so the homogeneous factor $(x - y)^E$ vanishes on the diagonal and $P_h(x, y(x; z)) \rightarrow 0$ as $x \rightarrow z^-$; simultaneously the integrand in (6.14) develops a non-integrable singularity $\sim (y(t; z) - t)^{-1-E}$ as $t \rightarrow z^-$, so the integral diverges. Their product is the finite $0 \cdot \infty$ limit

$$\frac{\pi(0, 0)}{1 - \rho z} = \lim_{x \rightarrow z^-} P_h(x, y(x; z)) \int_0^x \frac{h(t; (\gamma_1 + \gamma_2)z)}{P_h(t, y(t; z))} dt,$$

whose value is $P(z, z) = \pi(0, 0)/(1 - \rho z)$ by Corollary 2. Substituting (6.16) and separating the term containing the unknown P_y from the known remainder, this limit splits as

$$\begin{aligned} \frac{\pi(0, 0)}{1 - \rho z} &= \lim_{x \rightarrow z^-} \frac{\mu P_h(x, y(x; z))}{\gamma_1} \int_0^x \frac{P_y(y(t; z))}{t \cdot y(t; z) \cdot P_h(t, y(t; z))} dt \\ &\quad + \lim_{x \rightarrow z^-} \frac{\mu P_h(x, y(x; z)) \pi(0, 0)}{\gamma_1} \int_0^x \frac{1 - y(t; z)}{y(t; z) \cdot (y(t; z) - t) \cdot P_h(t, y(t; z))} dt. \end{aligned} \quad (6.17)$$

We now evaluate the two limits in (6.17) one by one; this both clarifies in what sense the kernel pairing below is to be read and provides an independent consistency check of the representation. The geometry on the diagonal characteristic is exactly linear,

$$y(t; z) - t = \frac{(\gamma_1 + \gamma_2)z - \gamma_2 t}{\gamma_1} - t = \frac{\Gamma}{\gamma_1} (z - t), \quad (6.18)$$

and along this characteristic ($C_1 = \Gamma z$) the exponent of Step 2 takes the closed form

$$E_z := E(\Gamma z) = \frac{\lambda z^2 - (\lambda + \mu)z + \mu}{\Gamma z} = \frac{(\mu - \lambda z)(1 - z)}{\Gamma z} > 0, \quad z \in (0, 1).$$

Every factor of P_h in (6.9) other than the exponential, the power of x , and $(x - y)^E$ depends on (x, y) only through C_1 , so along the characteristic the homogeneous ratio is *exact*:

$$\frac{P_h(x, y(x; z))}{P_h(t, y(t; z))} = R(x, t) \left(\frac{z - x}{z - t} \right)^{E_z}, \quad R(x, t) := e^{\frac{\kappa}{\gamma_1}(x-t)} \left(\frac{t}{x} \right)^{\mu/(\Gamma z)}, \quad (6.19)$$

with R smooth and $R(z, z) = 1$. Substituting (6.19) into the *second* limit of (6.17) and using (6.18), its integrand behaves as $w(z)(z - t)^{-1-E_z}$ near $t = z$, with $w(z) = \mu \pi(0, 0)(1 - z)/(\Gamma z)$; the truncated integral diverges like $(z - x)^{-E_z}/E_z$, the prefactor vanishes like $(z - x)^{E_z}$, and the product converges:

$$\begin{aligned} \lim_{x \rightarrow z^-} (z - x)^{E_z} \int^x w(z)(z - t)^{-1-E_z} dt &= \frac{w(z)}{E_z} = \frac{\mu \pi(0, 0)(1 - z)}{\Gamma z} \cdot \frac{\Gamma z}{(\mu - \lambda z)(1 - z)} \\ &= \frac{\mu \pi(0, 0)}{\mu - \lambda z} = \frac{\pi(0, 0)}{1 - \rho z}. \end{aligned}$$

The second limit alone thus reproduces the full diagonal value — an independent confirmation that the representation (6.14) is consistent with Corollary 2. The *first* limit, by contrast, vanishes: its integrand behaves as $(z - t)^{-E_z}$ near $t = z$, so for $E_z < 1$ the integral converges and is annihilated by the vanishing prefactor $(z - x)^{E_z}$, while for $E_z \geq 1$ the truncated integral grows only like $(z - x)^{1-E_z}$ (logarithmically at $E_z = 1$) and the product is still $O((z - x)^{\min(E_z, 1)} \ln(z - x)) \rightarrow 0$. Hence *at leading order the limit (6.17) reduces to the identity* $\pi(0, 0)/(1 - \rho z) = \pi(0, 0)/(1 - \rho z)$ *and carries no information about* P_y : the dependence on the boundary function enters only at the next order in $(z - x)$. Each limit is accordingly a *Hadamard finite part*: writing $P_h^{\text{diag}}(z) := \lim_{x \rightarrow z^-} P_h(x, y(x; z))$ (which vanishes pointwise), the kernel pairing below is the extraction of the truncation-independent, subleading part of the divergent integrals, and neither factor is meaningful in isolation. Writing the kernel and known term *formally* as

$$\mathcal{K}(z, t) := \frac{\mu \cdot P_h^{\text{diag}}(z)}{\gamma_1 \cdot t \cdot y(t; z) \cdot P_h(t, y(t; z))}, \quad (6.20)$$

$$\mathcal{B}(z) := \frac{\mu \cdot P_h^{\text{diag}}(z) \cdot \pi(0, 0)}{\gamma_1} \int_0^z \frac{1 - y(t; z)}{y(t; z) \cdot (y(t; z) - t) \cdot P_h(t, y(t; z))} dt,$$

to be read through the regularised limit (6.17) rather than pointwise, and setting $\mathcal{F}(z) = \pi(0, 0)/(1 - \rho z) - \mathcal{B}(z)$, equation (6.17) becomes the Volterra integral equation of the first kind

$$\mathcal{F}(z) = \int_0^z \mathcal{K}(z, t) P_y(y(t; z)) dt, \quad (6.21)$$

confirming Claim 1. The finite-part reading is not a decoration. As the limit-by-limit analysis shows, the displayed $\mathcal{K}$ and $\mathcal{B}$ are formal — the prefactor $P_h^{\text{diag}}(z)$ vanishes pointwise — and the equation for P_y lives entirely at subleading order in $(z - x)$. Giving that extraction a rigorous meaning, whether by matched asymptotics near $t = z$ or by a change of variable removing the $(z - t)^{-1-E_z}$ singularity before passing to the limit, is itself the analytical obstruction: it is this step, together with the fact that the kernel $\mathcal{K}$ depends on P_h through both arguments and does not simplify for general (γ_1, γ_2) , that places Model-B beyond the scope of this thesis. Equation (6.21) is recorded as the precise form of the open problem, not solved here; resolving it would require Volterra inversion methods together with the regularisation just described. $\square$

Model- B is therefore left open. Its role in the hierarchy is structural: it identifies the precise obstruction — the Volterra equation for P_y — that the tractable special cases B_2 and C_2 circumvent, each by a different mechanism. Setting $\gamma_2 = 0$ collapses the characteristic family from a two-parameter set of oblique lines to a one-parameter set of vertical lines, degenerating the PDE to an ODE in x alone; the boundary function $P_y(y)$ is then pinned by an analyticity argument at the interior singular point $x = y$. Model- C_2 removes jockeying entirely and pins P_y by a boundedness argument at the boundary of the domain. Both specialisations are analysed in the following sections.

Remark 1 (The full model X is *a fortiori* intractable). Model- B already exhibits the central obstruction of this work: even with abandonment switched off, pinning the boundary function $P_y(y)$ — and hence the joint PGF $P(x, y)$ — still reduces to the first-kind Volterra equation (6.4), which has no closed form. Reinstating abandonment ($\theta_1, \theta_2 > 0$) to recover the full parent model, Model- X , can only deepen the difficulty, and it does so in two ways.

First, the characteristics cease to be straight lines. For $\theta_1 = \theta_2 = 0$ the characteristic slope is the constant $-\gamma_2/\gamma_1$, giving the lines $\gamma_2 x + \gamma_1 y = C_1$ of Claim 1. Once $\theta_i > 0$ it becomes a genuine ratio of affine forms,

$$\frac{dy}{dx} = \frac{-\gamma_2 x + (\gamma_2 + \theta_2) y - \theta_2}{(\gamma_1 + \theta_1) x - \gamma_1 y - \theta_1},$$

whose integral curves are level sets of a product of two *non-integer* powers of linear forms in (x, y) — transcendental curves, not straight lines. Along them $y(x)$ is only implicitly defined, so the homogeneous solution P_h becomes a non-elementary quadrature and the Volterra kernel loses the explicit form it keeps for Model- B .

Second, the diagonal anchor is lost. For Model- B the diagonal value $P(z, z) = \pi(0, 0)/(1 - \rho z)$ is known outright (Corollary 2) and supplies the forcing of (6.4); under abandonment, setting $x = y = z$ in the fundamental equation no longer isolates $P(z, z)$, leaving a surviving diagonal-gradient term $[\theta_1 \partial_x P + \theta_2 \partial_y P]_{(z, z)}$, so the forcing is itself coupled to the unknown solution. Model- X is therefore left open. Model- B is the simplest model in which the obstruction already appears in full, and the sections that follow retreat to specialisations that remove it.

7 Analysis of Model- B_2 ($\gamma_1 > 0$, $\gamma_2 = 0$, $\theta_1 = \theta_2 = 0$)

Model- B_2 restricts jockeying to the direction $1 \rightarrow 2$ only ($\gamma_1 > 0$, $\gamma_2 = 0$) and retains no abandonment ($\theta_1 = \theta_2 = 0$). Setting $\gamma_2 = 0$ eliminates the $\partial P/\partial y$ term from the fundamental equation, reducing it to a first-order linear ODE in x for each fixed $y \in [0, 1]$. An integrating factor then yields a formal integral expression for $P(x, y)$; the boundary function $P_y(y)$ is pinned by requiring $P(\cdot, y)$ to be holomorphic at the interior singular point $x = y$, where the integrating factor carries a branch singularity with negative exponent. This analyticity condition replaces the boundedness argument used for Model- C_2 .

Stability. As in Models A and B , $\rho = (\lambda_1 + \lambda_2)/\mu < 1$. Jockeying conserves the total customer count and does not affect the stability condition. Figure 6 shows the queue layout.




Figure 6: Queue layout for Model- B_2 . Only the $1 \rightarrow 2$ jockeying direction is active (dashed, downward): class-1 customers may defect to Queue 2, but no customer can move in the reverse direction ($\gamma_2 = 0$).

7.1 Reduced fundamental equation

Setting $\gamma_2 = 0$ and $\theta_1 = \theta_2 = 0$ in the general fundamental equation (4.14) (Lemma 1) annihilates the $\partial P/\partial y$ term, whose coefficient $xy[\gamma_2(y - x) + \theta_2(y - 1)]$ vanishes identically, and leaves only

the γ_1 contribution to the $\partial P/\partial x$ term:

$$\begin{aligned} & [(\lambda_1 + \lambda_2 + \mu)xy - \mu y - \lambda_1 x^2 y - \lambda_2 x y^2] P + \gamma_1 x y (x - y) \frac{\partial P}{\partial x} \\ &= \mu(x - y) P_y(y) - \mu x(1 - y) \pi(0, 0). \end{aligned} \quad (7.1)$$

For each fixed $y \in (0, 1]$ this is a first-order linear ODE in x ; the entire analysis of this section is built on it.

Theorem 2. *Consider the two-class non-preemptive priority queueing system of Model-B₂ with per-customer jockeying rate $\gamma_1 > 0$ from Queue 1 to Queue 2 ($\gamma_2 = 0$, $\theta_1 = \theta_2 = 0$). Under the stability condition $\rho = \lambda/\mu < 1$, $\lambda = \lambda_1 + \lambda_2$, the boundary generating function is*

$$P_y(y) = \frac{\mu \pi(0, 0)}{\mu - \lambda y} \cdot \frac{{}_1F_1(\alpha(y) + 1; b^*(y); -\lambda_1 y/\gamma_1)}{{}_1F_1(\alpha(y); b^*(y); -\lambda_1 y/\gamma_1)}, \quad (7.2)$$

and the bivariate PGF, for $0 \leq x \leq y$, is

$$P(x, y) = \frac{\mu e^{\lambda_1 x/\gamma_1} x^{-\alpha} (y - x)^{-\beta}}{\gamma_1 y} \left[P_y(y) \mathcal{H}_1(x, y) + (1 - y) \pi(0, 0) \mathcal{H}_2(x, y) \right], \quad (7.3)$$

where ${}_1F_1$ is Kummer's confluent hypergeometric function, the exponents are

$$\alpha(y) = \frac{\mu}{\gamma_1 y}, \quad \beta(y) = \frac{(1 - y)(\lambda y - \mu)}{\gamma_1 y}, \quad b^*(y) = \alpha + \beta + 1 = \frac{\mu + \lambda(1 - y) + \gamma_1}{\gamma_1}, \quad (7.4)$$

the auxiliary integrals are

$$\mathcal{H}_1(x, y) = \int_0^x e^{-\lambda_1 t/\gamma_1} t^{\alpha-1} (y - t)^\beta dt, \quad (7.5)$$

$$\mathcal{H}_2(x, y) = \int_0^x e^{-\lambda_1 t/\gamma_1} t^\alpha (y - t)^{\beta-1} dt, \quad (7.6)$$

and the empty-state probability $\pi(0, 0) = \rho(1 - \rho)$ is given by Corollary 3 below. On $y < x \leq 1$ the same expression (7.3) holds with $\mathcal{H}_1, \mathcal{H}_2$ replaced by their finite-part continuations.

Corollary 3. *The empty-system probability π_0 (server idle, no one waiting) and the both-queues-empty probability $\pi(0, 0)$ (server busy, $N_1 = N_2 = 0$) of Model-B₂ are*

$$\pi_0 = 1 - \rho, \quad \pi(0, 0) = \rho(1 - \rho),$$

identical to the Model-A baseline.

Proof. Because one-way jockeying conserves the total customer count and there is no abandonment, the total-count process L (equal to the number in queue N plus the single customer in service whenever the server is busy) is itself a birth-death chain: it increases at rate $\lambda = \lambda_1 + \lambda_2$ on any arrival and decreases at rate μ on any service completion, irrespective of class. Thus L is an $M/M/1$ queue, and by the memorylessness of the exponential clocks its stationary law is geometric, $\mathbb{P}(L = k) = (1 - \rho)\rho^k$; in particular $\pi_0 = \mathbb{P}(L = 0) = 1 - \rho$.

For $\pi(0, 0)$, set $x = y$ in the reduced fundamental equation (7.1). The convection term $\gamma_1 x y (x - y) \partial_x P$ and the boundary term $\mu(x - y) P_y(y)$ both carry the factor $(x - y)$ and vanish on the diagonal; dividing the surviving terms by the common factor x leaves

$$[(\lambda_1 + \lambda_2 + \mu)x - \mu - (\lambda_1 + \lambda_2)x^2] P(x, x) = -\mu(1 - x) \pi(0, 0).$$

Factoring the bracket and cancelling $(x - 1)$ against $-(1 - x) = (x - 1)$ on the right:

$$\begin{aligned} [-\lambda x^2 + (\lambda + \mu)x - \mu] P(x, x) &= (\mu - \lambda x)(x - 1) P(x, x) \\ &= \mu(x - 1) \pi(0, 0), \end{aligned}$$

so that

$$P(x, x) = \frac{\mu \pi(0, 0)}{\mu - \lambda x} = \frac{\pi(0, 0)}{1 - \rho x}. \quad (7.7)$$

Evaluating at $x = 1$ with $P(1, 1) = 1 - \pi_0 = \rho$ (using $\pi_0 = 1 - \rho$ just established) yields $\pi(0, 0) = \rho(1 - \rho)$. $\square$

Proof of Theorem 2. Stage 1: Standard form and the integrating factor.

For any fixed $y \neq 0$, divide the reduced fundamental equation (7.1) through by the coefficient $\gamma_1 xy(x-y)$ of $\partial P / \partial x$. The factor y cancels against the overall factor y in the bracket multiplying P , putting the equation in standard first-order form $\frac{dP}{dx} + p(x; y) P = q(x; y)$ with

$$p(x; y) = \frac{(\lambda_1 + \lambda_2 + \mu)x - \mu - \lambda_1 x^2 - \lambda_2 xy}{\gamma_1 x(x-y)}, \quad (7.8)$$

$$q(x; y) = \frac{\mu P_y(y)}{\gamma_1 xy} - \frac{\mu(1-y)\pi(0,0)}{\gamma_1 y(x-y)}. \quad (7.9)$$

Write the numerator of p as a polynomial in x ,

$$N(x) = -\lambda_1 x^2 + [(\lambda_1 + \lambda_2 + \mu) - \lambda_2 y]x - \mu = -\lambda_1 (x - x^*(y))(x - x^+(y)),$$

which factors through the roots $x^*(y), x^+(y)$ of the Model-A kernel (Equation (5.8)), with $\lambda_1 x^* x^+ = \mu$. The partial-fraction coefficients of $p = N(x)/[\gamma_1 x(x-y)]$ follow by residues:

$$\begin{aligned} A &= -\lambda_1 && \text{(ratio of leading } x^2 \text{ terms),} \\ B &= \frac{N(0)}{0-y} = \frac{-\mu}{-y} = \frac{\mu}{y}, \\ C &= \frac{N(y)}{y} = \frac{-(\lambda y - \mu)(y-1)}{y} = \frac{(1-y)(\lambda y - \mu)}{y}, \end{aligned}$$

where $N(y) = -\lambda y^2 + (\lambda + \mu)y - \mu = -(\lambda y - \mu)(y-1)$. Hence

$$p(x; y) = \frac{1}{\gamma_1} \left[-\lambda_1 + \frac{\mu/y}{x} + \frac{(1-y)(\lambda y - \mu)/y}{x-y} \right]. \quad (7.10)$$

We henceforth work on the principal triangle $0 \leq x \leq y$, on which $x-y \leq 0$; the complementary region $y < x \leq 1$ is recovered at the end by analytic continuation. There the real antiderivative of the simple-pole term is $\beta \ln|x-y| = \beta \ln(y-x)$, so integrating the partial-fraction form (7.10) term by term gives, up to an additive constant,

$$\int^x p(t; y) dt = \frac{1}{\gamma_1} \left[-\lambda_1 x + \frac{\mu}{y} \ln x + \frac{(1-y)(\lambda y - \mu)}{y} \ln(y-x) \right].$$

Exponentiating (an integrating factor is determined only up to a multiplicative constant, which we drop) yields the *real* integrating factor

$$\mu_I(x; y) = e^{-\lambda_1 x / \gamma_1} x^{\alpha(y)} (y-x)^{\beta(y)},$$

with exponents $\alpha(y) = \mu/(\gamma_1 y) > 0$ and $\beta(y) = (1-y)(\lambda y - \mu)/(\gamma_1 y)$ as in (7.4); one verifies directly that $\mu'_I/\mu_I = -\lambda_1/\gamma_1 + \alpha/x + \beta/(x-y) = p(x; y)$, using $\frac{d}{dx} \ln(y-x)^\beta = \beta/(x-y)$. Working with $(y-x) > 0$ in place of the complex base $(x-y)$ keeps every quantity manifestly real throughout the principal triangle. Since $\alpha > 0$ we have $\mu_I \rightarrow 0$ as $x \rightarrow 0^+$, so the boundary term $P(0, y) \mu_I(0; y)$ vanishes when we integrate $\frac{d}{dx}[P \cdot \mu_I] = q \cdot \mu_I$ over $(0, x]$:

$$P(x, y) \mu_I(x; y) = \int_0^x q(t; y) \mu_I(t; y) dt.$$

It remains to evaluate the right-hand side. Multiplying q from (7.9) by μ_I and rewriting the second term with the elementary identity $(y-t)^\beta/(t-y) = -(y-t)^{\beta-1}$, the integrand separates into two manifestly real pieces,

$$q(t; y) \mu_I(t; y) = \frac{\mu}{\gamma_1 y} \left[P_y(y) e^{-\lambda_1 t / \gamma_1} t^{\alpha-1} (y-t)^\beta + (1-y)\pi(0,0) e^{-\lambda_1 t / \gamma_1} t^\alpha (y-t)^{\beta-1} \right],$$

the minus sign of (7.9) having been absorbed by that identity. With the auxiliary integrals $\mathcal{H}_1, \mathcal{H}_2$ of (7.5)–(7.6) and dividing through by $\mu_I(x; y) = e^{-\lambda_1 x / \gamma_1} x^\alpha (y-x)^\beta$, the solution on the principal triangle $0 \leq x \leq y$ is

$$P(x, y) = \frac{\mu e^{\lambda_1 x / \gamma_1} x^{-\alpha} (y-x)^{-\beta}}{\gamma_1 y} \left[P_y(y) \mathcal{H}_1(x, y) + (1-y)\pi(0,0) \mathcal{H}_2(x, y) \right], \quad (7.11)$$

which still contains the unknown boundary function $P_y(y)$; Stage 2 determines it.

Stage 2: Determination of $P_y(y)$ by analyticity at $x = y$.

The factor $(x - y)$ surviving in the x -convection coefficient places the operative singularity at the interior point $x = y$, exactly the analyticity-pinning regime anticipated in the remark following Lemma 1. Unlike Models A and C_2 , no probabilistic (Pollaczek–Khinchine) route is available here. One-way jockeying makes the effective service experienced by a class-2 customer depend on the class-1 queue contents at the (random) jockeying instants, so the class-2 input is no longer a Poisson stream feeding an $M/G/1$ subsystem and no busy-period reduction applies (cf. the discussion of Model- B in §6). We therefore determine $P_y(y)$ *purely analytically*, by imposing holomorphy of $P(\cdot, y)$ at the interior point $x = y$, where the solution (7.11) is *a priori* singular.

Under stability $\mu - \lambda y > 0$ for all $y \in (0, 1)$, so the second exponent is negative,

$$\beta(y) = \frac{(1 - y)(\lambda y - \mu)}{\gamma_1 y} < 0, \quad 0 < y < 1. \quad (7.12)$$

The prefactor $(y - x)^{-\beta} = (y - x)^{|\beta|}$ in (7.11) therefore *vanishes* as $x \rightarrow y^-$, while the integrals $\mathcal{H}_1, \mathcal{H}_2$ may *diverge* there; the product is an indeterminate $0 \cdot \infty$ form. Because $P(\cdot, y)$ is a probability generating function it must be holomorphic on the open unit disc, and in particular analytic at the interior point $x = y$. As we show next, the general solution carries a non-integer power $(y - x)^{|\beta|}$ — a branch term that vanishes at $x = y$ yet is not analytic there — and holomorphy forces its coefficient to zero; this single condition pins down $P_y(y)$.

Local analysis at $x = y$. To exhibit the two modes explicitly, read off the local behaviour of the ODE (7.8)–(7.9) at $x = y$. By the partial fraction (7.10), p has a simple pole at $x = y$ with residue β ; by (7.9), q has the leading singularity $-D/(x - y)$ with $D = \mu(1 - y)\pi(0, 0)/(\gamma_1 y)$. The local model equation is therefore

$$P'(x, y) + \frac{\beta}{x - y} P(x, y) = \frac{-D}{x - y} + (\text{analytic near } x = y).$$

A constant particular solution $P \equiv P_*$ requires $\beta P_* = -D$, so, cancelling the common factor $(1 - y)/(\gamma_1 y)$ shared by D and β ,

$$P(y, y) = P_* = -\frac{D}{\beta} = -\frac{\mu \pi(0, 0)}{\lambda y - \mu} = \frac{\mu \pi(0, 0)}{\mu - \lambda y} = \frac{\rho(1 - \rho)}{1 - \rho y}, \quad (7.13)$$

consistent with the diagonal value (7.7) (using $\pi(0, 0) = \rho(1 - \rho)$). The general local solution adds a multiple of the homogeneous mode $(x - y)^{-\beta} = (x - y)^{|\beta|}$. Since $|\beta| \notin \mathbb{Z}$ generically this is a branch term: it vanishes at $x = y$ but is not analytic there, so holomorphy forces its coefficient to zero. The remaining task is to locate that coefficient inside the explicit solution (7.11).

No-branch condition and $P_y(y)$. Write the analytic, non-vanishing prefactor of (7.11) as $\Psi(x) = \mu e^{\lambda_1 x / \gamma_1} x^{-\alpha} / (\gamma_1 y)$, so that $P(x, y) = \Psi(x) (y - x)^{-\beta} [P_y(y) \mathcal{H}_1 + (1 - y)\pi(0, 0) \mathcal{H}_2]$. As $x \rightarrow y^-$ each integral admits a local expansion, obtained by Taylor-expanding its smooth integrand factor — $g_1(t) = e^{-\lambda_1 t / \gamma_1} t^{\alpha-1}$ for $\mathcal{H}_1$, $g_2(t) = e^{-\lambda_1 t / \gamma_1} t^\alpha$ for $\mathcal{H}_2$ — about $t = y$ and integrating the leading power of $(y - t)$:

$$\begin{aligned} \mathcal{H}_1(x, y) &= \text{f.p. } \mathcal{H}_1(y, y) - \frac{g_1(y)}{\beta + 1} (y - x)^{\beta+1} + \dots, \\ \mathcal{H}_2(x, y) &= \text{f.p. } \mathcal{H}_2(y, y) - \frac{g_2(y)}{\beta} (y - x)^\beta + \dots, \end{aligned}$$

where f.p. denotes the Hadamard finite part (the $x \rightarrow y$ limit of the integral once its singular term has been subtracted) and the omitted terms are of higher order. Multiplying by $\Psi(x)(y - x)^{-\beta}$ sorts the contributions to $P(x, y)$ by their power of $(y - x)$:

- the singular part of $\mathcal{H}_2$ gives $-\Psi(y)(1 - y)\pi(0, 0) g_2(y)/\beta = -D/\beta = P_*$, exactly the regular value (7.13) (using $\Psi(y) g_2(y) = \mu/(\gamma_1 y)$);
- the singular part of $\mathcal{H}_1$ gives a term $\propto (y - x)^1$, analytic at $x = y$;
- the finite parts, riding on the prefactor, produce the branch term $\Psi(x)(y - x)^{-\beta} [P_y(y) \text{f.p. } \mathcal{H}_1 + (1 - y)\pi(0, 0) \text{f.p. } \mathcal{H}_2]$.

Only the last is non-analytic $((y-x)^{|\beta|}$ with $|\beta| \notin \mathbb{Z}$), so holomorphy of $P(\cdot, y)$ at $x = y$ requires its bracket to vanish:

$$P_y(y) \text{ f.p. } \mathcal{H}_1(y, y) + (1-y) \pi(0, 0) \text{ f.p. } \mathcal{H}_2(y, y) = 0. \quad (7.14)$$

To evaluate the finite parts, substitute $t = ys$ ($dt = y ds$) in (7.5)–(7.6) and pull out $y^{\alpha+\beta}$ from $t^{\alpha-1}(y-t)^\beta dt = y^{\alpha+\beta} s^{\alpha-1}(1-s)^\beta ds$ (resp. $t^\alpha(y-t)^{\beta-1} dt = y^{\alpha+\beta} s^\alpha(1-s)^{\beta-1} ds$); with $c := \lambda_1 y / \gamma_1$,

$$\mathcal{H}_1(y, y) = y^{\alpha+\beta} \int_0^1 s^{\alpha-1} (1-s)^\beta e^{-cs} ds, \quad \mathcal{H}_2(y, y) = y^{\alpha+\beta} \int_0^1 s^\alpha (1-s)^{\beta-1} e^{-cs} ds.$$

Matching the Euler integral representation (3.10), written as $\int_0^1 t^{a-1} (1-t)^{b-1} e^{zt} dt = B(a, b) {}_1F_1(a; a+b; z)$, with $(a, b, z) = (\alpha, \beta+1, -c)$ for $\mathcal{H}_1$ and $(a, b, z) = (\alpha+1, \beta, -c)$ for $\mathcal{H}_2$ — read as its analytic continuation in the second parameter b beyond the convergence domain $b > 0$, since $\beta < 0$ by (7.12), which is precisely the finite part — and noting $a+b = b^*(y)$ in both cases, gives

$$\text{f.p. } \mathcal{H}_1(y, y) = B(\alpha, \beta+1) y^{\alpha+\beta} {}_1F_1(\alpha; b^*(y); -\lambda_1 y / \gamma_1), \quad (7.15)$$

$$\text{f.p. } \mathcal{H}_2(y, y) = B(\alpha+1, \beta) y^{\alpha+\beta} {}_1F_1(\alpha+1; b^*(y); -\lambda_1 y / \gamma_1). \quad (7.16)$$

Solving the no-branch condition (7.14) for $P_y(y)$ and inserting the finite parts (7.15)–(7.16), the common factor $y^{\alpha+\beta}$ cancels and

$$P_y(y) = -(1-y) \pi(0, 0) \frac{\text{f.p. } \mathcal{H}_2(y, y)}{\text{f.p. } \mathcal{H}_1(y, y)} = -(1-y) \pi(0, 0) \frac{B(\alpha+1, \beta)}{B(\alpha, \beta+1)} \frac{{}_1F_1(\alpha+1; b^*; -\lambda_1 y / \gamma_1)}{{}_1F_1(\alpha; b^*; -\lambda_1 y / \gamma_1)}. \quad (7.17)$$

The Beta ratio reduces, via $\Gamma(\alpha+1) = \alpha\Gamma(\alpha)$ and $\Gamma(\beta+1) = \beta\Gamma(\beta)$, to

$$\begin{aligned} \frac{B(\alpha+1, \beta)}{B(\alpha, \beta+1)} &= \frac{\Gamma(\alpha+1)\Gamma(\beta)}{\Gamma(\alpha+\beta+1)} \cdot \frac{\Gamma(\alpha+\beta+1)}{\Gamma(\alpha)\Gamma(\beta+1)} \\ &= \frac{\Gamma(\alpha+1)}{\Gamma(\alpha)} \cdot \frac{\Gamma(\beta)}{\Gamma(\beta+1)} = \frac{\alpha}{\beta}, \end{aligned} \quad (7.18)$$

and the surviving prefactor simplifies, using $\alpha/\beta = \mu / [(1-y)(\lambda y - \mu)]$, as

$$\begin{aligned} -(1-y) \frac{\alpha}{\beta} &= -(1-y) \cdot \frac{\mu}{(1-y)(\lambda y - \mu)} \\ &= -\frac{\mu}{\lambda y - \mu} = \frac{\mu}{\mu - \lambda y}. \end{aligned} \quad (7.19)$$

Substituting (7.18)–(7.19) into (7.17) gives the boundary generating function

$$P_y(y) = \frac{\mu \pi(0, 0)}{\mu - \lambda y} \cdot \frac{{}_1F_1(\alpha(y) + 1; b^*(y); -\lambda_1 y / \gamma_1)}{{}_1F_1(\alpha(y); b^*(y); -\lambda_1 y / \gamma_1)}, \quad (7.20)$$

confirming (7.2). The prefactor equals the regular value $P(y, y)$ of (7.13), and the Kummer ratio is the boundary correction. As $y \rightarrow 0^+$ the argument $-\lambda_1 y / \gamma_1 \rightarrow 0$, so the ratio tends to 1 and $P_y(0) = \pi(0, 0)$, consistent with $P_y(0) = \sum_{n_2} \pi(0, n_2) \cdot 0^{n_2} = \pi(0, 0)$.

Stage 3: Recovery of $P(x, y)$.

Substituting (7.20) into the real-form expression (7.11) gives, for $0 \leq x \leq y$,

$$\begin{aligned} P(x, y) &= \frac{\mu e^{\lambda_1 x / \gamma_1} x^{-\alpha} (y-x)^{-\beta}}{\gamma_1 y} \left[\frac{\mu \pi(0, 0)}{\mu - \lambda y} \cdot \frac{{}_1F_1(\alpha+1; b^*; -\frac{\lambda_1 y}{\gamma_1})}{{}_1F_1(\alpha; b^*; -\frac{\lambda_1 y}{\gamma_1})} \cdot \mathcal{H}_1(x, y) \right. \\ &\quad \left. + (1-y) \pi(0, 0) \cdot \mathcal{H}_2(x, y) \right], \end{aligned}$$

which is the statement of the theorem (Equation (7.3)). The no-branch condition (7.14) guarantees that this extends analytically across $x = y$, with diagonal value $P(y, y) = \rho(1-\rho)/(1-\rho y)$ as computed in (7.13). On $y < x \leq 1$ the same formula holds with $\mathcal{H}_1, \mathcal{H}_2$ read as their finite-part continuations; $\pi(0, 0) = \rho(1-\rho)$ from Corollary 3 makes the solution fully explicit. The recovery of Model-A in the limit $\gamma_1 \rightarrow 0^+$ is deferred to Remark 2 below, as the limit is singular and merits separate discussion. $\square$

Remark 2 (The singular limit $\gamma_1 \rightarrow 0^+$). The no-jockeying limit is cleanest at the level of the fundamental equation, but $\gamma_1 \rightarrow 0^+$ is a *singular* perturbation — the coefficient $\gamma_1 xy(x-y)$ of the highest derivative $\partial_x P$ vanishes — so one must check that the convection term itself, and not merely its coefficient, disappears; the derivative $\partial_x P$ could *a priori* blow up. The probabilistic bound rules this out. Each $P(\cdot, \cdot)$ is a probability generating function, so $|P| \leq 1$ on the closed unit polydisc *uniformly* in γ_1 ; by Cauchy's estimate $\partial_x P$ is then bounded uniformly in γ_1 on every compact subset of the open polydisc, whence $\gamma_1 xy(x-y) \partial_x P \rightarrow 0$ there. The fundamental equation (7.1) therefore degenerates to the purely algebraic Model-A equation (5.2) (the coefficients of P and the right-hand side are γ_1 -independent, and $\pi(0,0) = \rho(1-\rho)$ by Corollary 3). Since $\{P\}_{\gamma_1}$ is uniformly bounded it is a normal family, so it converges — along subsequences, hence in full by uniqueness of the polydisc-analytic generating-function solution of (5.2) — to the Model-A PGF (5.3). In particular the boundary function reduces to the Model-A value (5.10),

$$P_y(y) \xrightarrow{\gamma_1 \rightarrow 0^+} \pi(0,0) \frac{x^*(y)(1-y)}{x^*(y)-y}, \quad (7.21)$$

with $x^*(y)$ the Model-A kernel root (5.8); this recovery needs no special-function asymptotics.

It is instructive to read the same limit off the closed form (7.20). Its prefactor satisfies $\mu \pi(0,0)/(\mu - \lambda y) = \pi(0,0)/(1 - \rho y)$ and is γ_1 -independent, so (7.21) is equivalent to the Kummer ratio tending to

$$\frac{{}_1F_1(\alpha+1; b^*; -\lambda_1 y/\gamma_1)}{{}_1F_1(\alpha; b^*; -\lambda_1 y/\gamma_1)} \xrightarrow{\gamma_1 \rightarrow 0^+} (1-\rho y) \frac{x^*(y)(1-y)}{x^*(y)-y}. \quad (7.22)$$

Note the factor $(1-\rho y)$: the ratio does *not* tend to the bare rational expression $x^*(y)(1-y)/(x^*(y)-y)$, the discrepancy being precisely the prefactor. This is consistent with the three-term recurrence for ${}_1F_1$ in its first parameter [DLM, §13.3]: as $\gamma_1 \rightarrow 0^+$ the parameters $\alpha = \mu/(\gamma_1 y)$, b^* and $c = \lambda_1 y/\gamma_1$ all grow like $1/\gamma_1$, and to leading order that recurrence reduces to a quadratic in the ratio whose admissible root is the right-hand side of (7.22).

7.2 Mean queue lengths and mean waiting times

Total queue length from the diagonal

Setting $x = y = z$ in the reduced fundamental equation (7.1), the jockeying convection term $\gamma_1 xy(x-y)\partial_x P$ vanishes at $x = y = z$ (factor $(x-y) = 0$), as does the boundary term $\mu(x-y)P_y(y)$. The remaining balance is identical to Model-A's, so the same diagonal cancellation gives:

$$P(z, z) = \frac{\rho(1-\rho)}{1-\rho z},$$

the same formula as (5.23). Since $P(z, z) = \mathbb{E}[z^{N_1+N_2}] = \mathbb{E}[z^N]$ is the PGF of the total in-queue count N , the generating-function identity $\frac{d}{dz} P(z, z)|_{z=1} = \mathbb{E}[N_1] + \mathbb{E}[N_2]$ holds; differentiating by the quotient rule and evaluating at $z = 1$:

$$\mathbb{E}[N_1] + \mathbb{E}[N_2] = \frac{d}{dz} \frac{\rho(1-\rho)}{1-\rho z} \Big|_{z=1} = \frac{\rho(1-\rho) \cdot \rho}{(1-\rho z)^2} \Big|_{z=1} = \frac{\rho^2(1-\rho)}{(1-\rho)^2} = \frac{\rho^2}{1-\rho}.$$

One-way jockeying does not alter the total mean queue length relative to Model-A, since the diagonal PGF is unchanged.

Class-1 mean queue length

We extract $\mathbb{E}[N_1] = \frac{\partial P}{\partial x} \Big|_{(1,1)}$ by first reducing the fundamental equation (7.1) at $y = 1$, and then differentiating the resulting marginal PGF at $x = 1$.

Step 1: marginal PGF $P(x, 1)$. Evaluating (7.1) at $y = 1$, the bracket multiplying P becomes $(\lambda_1 + \mu)x - \mu - \lambda_1 x^2 = -\lambda_1(x-1)(x-1/\rho_1)$, the convection coefficient is $\gamma_1 x(x-1)$, and the right-hand side reduces to $\mu(x-1)P_y(1)$ (the $\pi(0,0)$ term carries the factor $(1-y) = 0$). Every term therefore shares the common factor $(x-1)$; cancelling it leaves the first-order linear ODE in x :

$$-\lambda_1(x-1/\rho_1)P(x, 1) + \gamma_1 x \frac{d}{dx} P(x, 1) = \mu P_y(1).$$

The integrating factor for this ODE is $e^{-\lambda_1 x/\gamma_1} x^{\mu/\gamma_1}$ (noting $\lambda_1/\rho_1 = \mu$), and integrating from 0 to x (the boundary term vanishes since $0^{\mu/\gamma_1} = 0$):

$$P(x, 1) e^{-\lambda_1 x/\gamma_1} x^{\mu/\gamma_1} = \frac{\mu P_y(1)}{\gamma_1} \int_0^x e^{-\lambda_1 t/\gamma_1} t^{\mu/\gamma_1 - 1} dt.$$

Rearranging and writing $\alpha = \mu/\gamma_1$ and $\mathcal{J}(x) = \int_0^x e^{-\lambda_1 t/\gamma_1} t^{\alpha-1} dt$:

$$P(x, 1) = \frac{\mu P_y(1)}{\gamma_1} \cdot e^{\lambda_1 x/\gamma_1} x^{-\alpha} \mathcal{J}(x).$$

The constant $P_y(1)$ follows from $P(1, 1) = 1 - \pi_0 = \rho$, which gives $P_y(1) = \rho \gamma_1 / (\mu e^{\lambda_1/\gamma_1} \mathcal{J}(1))$.

Step 2: differentiate $P(x, 1)$ at $x = 1$. Writing $P(x, 1) = C \cdot f(x) \cdot \mathcal{J}(x)$ with $C = \mu P_y(1)/\gamma_1$ and $f(x) = e^{\lambda_1 x/\gamma_1} x^{-\alpha}$, the product rule together with Leibniz's rule ($d\mathcal{J}/dx = e^{-\lambda_1 x/\gamma_1} x^{\alpha-1}$, the integrand at the upper limit) gives:

$$\frac{d}{dx} P(x, 1) = C \left[f'(x) \mathcal{J}(x) + f(x) \cdot e^{-\lambda_1 x/\gamma_1} x^{\alpha-1} \right].$$

The logarithmic derivative $f'(x)/f(x) = \lambda_1/\gamma_1 - \alpha/x = (\lambda_1 x - \mu)/(\gamma_1 x)$, and the second product term simplifies as $f(x) e^{-\lambda_1 x/\gamma_1} x^{\alpha-1} = x^{-1}$, so:

$$\frac{d}{dx} P(x, 1) = \frac{\mu P_y(1)}{\gamma_1} \left[e^{\lambda_1 x/\gamma_1} x^{-\alpha} \cdot \frac{\lambda_1 x - \mu}{\gamma_1 x} \cdot \mathcal{J}(x) + x^{-1} \right].$$

Step 3: evaluate at $x = 1$.

$$\begin{aligned} \mathbb{E}[N_1] &= \left. \frac{d}{dx} P(x, 1) \right|_{x=1} \\ &= \frac{\mu P_y(1)}{\gamma_1} \left[e^{\lambda_1/\gamma_1} \cdot \frac{\lambda_1 - \mu}{\gamma_1} \cdot \mathcal{J}(1) + 1 \right], \end{aligned} \tag{7.23}$$

where $\mathcal{J}(1) = \int_0^1 e^{-\lambda_1 t/\gamma_1} t^{\mu/\gamma_1 - 1} dt = (\gamma_1/\lambda_1)^{\mu/\gamma_1} \gamma(\mu/\gamma_1, \lambda_1/\gamma_1)$ and $\gamma(\cdot, \cdot)$ is the lower incomplete gamma function. The factor $(\lambda_1 - \mu) < 0$ (since $\rho_1 = \lambda_1/\mu < 1$), confirming that $e^{\lambda_1/\gamma_1}(\lambda_1 - \mu)\mathcal{J}(1)/\gamma_1 + 1 > 0$. No elementary simplification of (7.23) is known; the expression is evaluated numerically. As $\gamma_1 \rightarrow 0^+$ the Model-A value (5.22) is recovered.

Class-2 mean queue length by subtraction

$$\mathbb{E}[N_2] = (\mathbb{E}[N_1] + \mathbb{E}[N_2]) - \mathbb{E}[N_1] = \frac{\rho^2}{1 - \rho} - \mathbb{E}[N_1].$$

Since jockeying defections reduce the class-1 queue, $\mathbb{E}[N_1]$ is smaller than the Model-A value $\rho\rho_1/(1 - \rho_1)$, so $\mathbb{E}[N_2]$ correspondingly exceeds the Model-A value.

Mean waiting times via Little's Law

Little's Law applies here without qualification. It is a rate-conservation identity — for any subsystem in which customers are neither created nor destroyed internally, the time-average number present equals the long-run entry rate times the mean time an entrant spends inside — and it holds for *any* positive-recurrent system with finite mean queue length (here ensured by $\rho < 1$, which gives $\mathbb{E}[N_1] \leq \mathbb{E}[N] = \rho^2/(1 - \rho) < \infty$ and so makes the jockeying rate $\gamma_1 \mathbb{E}[N_1]$ below a genuine finite rate), requiring neither Poisson input nor independence. In particular it remains valid even though one-way jockeying has destroyed the per-class Poisson/renewal structure (the very obstruction that ruled out a probabilistic route in Stage 2). We apply it to each queue (the waiting line, excluding the customer in service), with $\mathbb{E}[W_i]$ understood as the mean time a customer spends in Queue i averaged over all customers that enter it — an arrival routed directly into service contributing a zero-length visit.

Class 1. Every class-1 arrival is counted as an entrant to Queue 1 at rate λ_1 (those finding the server idle enter service at once, contributing zero waiting time), and each leaves Queue 1 either by

entering service or by jockeying to Queue 2. Since no other stream feeds Queue 1, its throughput is exactly λ_1 , and Little's Law gives

$$\mathbb{E}[W_1] = \frac{\mathbb{E}[N_1]}{\lambda_1}. \quad (7.24)$$

Class 2. Queue 2 has *two* input streams: the class-2 Poisson arrivals at rate λ_2 and the jockeying transfers from Queue 1. Because each of the N_1 waiting class-1 customers jockeys independently at rate γ_1 , the long-run transfer rate is $\gamma_1 \mathbb{E}[N_1]$ (a genuine rate by the ergodic theorem, equal to the time average of the instantaneous jockeying intensity $\gamma_1 N_1$). The throughput of Queue 2 is therefore $\lambda_2 + \gamma_1 \mathbb{E}[N_1]$, and Little's Law reads $\mathbb{E}[N_2] = (\lambda_2 + \gamma_1 \mathbb{E}[N_1]) \mathbb{E}[W_2]$, so that

$$\mathbb{E}[W_2] = \frac{\mathbb{E}[N_2]}{\lambda_2 + \gamma_1 \mathbb{E}[N_1]}. \quad (7.25)$$

Here $\mathbb{E}[W_2]$ is the mean time in the class-2 line averaged over *all* Queue-2 entrants of both streams; a jockeyed customer's Queue-2 clock starts at the transfer instant, so its total wait is an initial Queue-1 segment (of mean $\mathbb{E}[W_1]$) followed by this Queue-2 segment. Thus (7.25) is the per-Queue-2-entrant residence time, not the end-to-end delay of a tagged class-2 *arrival*.

The jockeying inflow is what distinguishes (7.25) from the single-stream Models A and C₂: the naive denominator λ_2 would undercount the entrants to Queue 2 and is correct only in the no-jockeying limit. Indeed, as $\gamma_1 \rightarrow 0^+$ the transfer rate $\gamma_1 \mathbb{E}[N_1] \rightarrow 0$, so (7.25) reduces to $\mathbb{E}[N_2]/\lambda_2$ and both (7.24) and (7.25) recover the Model-A values (5.25)–(5.26).

8 Analysis of Model-C₂ ($\theta_1 > 0$, $\gamma_1 = \gamma_2 = \theta_2 = 0$)

In this section we derive $P(x, y)$ for the model in which only class-1 customers abandon, with no class-2 abandonment and no jockeying. Figure 7 shows the queue layout.




Figure 7: Queue layout for Model-C₂. Class-1 customers abandon the system at rate θ_1 per waiting customer (dotted arrow, upward from Queue 1). Class-2 never abandons ($\theta_2 = 0$) and there is no jockeying.

The rationale for this parameter choice is twofold. First, class-2 still sees Poisson arrivals, so the probabilistic arguments used for Model-A carry over. Second, substituting these parameters into the fundamental equation (8.3) yields a *first-order linear ODE* in x in which $P_y(y)$ can be taken outside the integral, sparing us a Volterra-type equation.

Theorem 3. *Consider a queueing system with 2 priority classes with Poisson arrival rates λ_1 and λ_2 , where class-1 customers have non-preemptive priority over those of class-2, exponentially distributed service times of rate μ independently of the class, and per-customer abandonment of class-1 customers from the queue at rate θ_1 (with no class-2 abandonment and no jockeying). Under the stability condition*

$$\rho_2 \cdot {}_1F_1\left(1; \frac{\mu}{\theta_1} + 1; \frac{\lambda_1}{\theta_1}\right) < 1 \iff \lambda_2 \mathbb{E}[B_C] < 1,$$

the probability generating function $P(x, y)$ of such a system is given by

$$P(x, y) = \frac{e^{\lambda_1 x / \theta_1} \mu \pi(0, 0) (1 - y)}{\theta_1 x^{\mu / \theta_1} (1 - x)^{\lambda_2 (1 - y) / \theta_1} (\tilde{B}_C(\lambda_2 (1 - y)) - y)} \\ \times \left[\tilde{B}_C(\lambda_2 (1 - y)) \mathcal{H}_1(x, y) - (1 - \tilde{B}_C(\lambda_2 (1 - y))) \mathcal{H}_2(x, y) \right],$$

where

$$\mathcal{H}_1(x, y) = \int_0^x e^{-\lambda_1 t / \theta_1} t^{\mu / \theta_1 - 1} (1 - t)^{\lambda_2 (1 - y) / \theta_1} dt, \\ \mathcal{H}_2(x, y) = \int_0^x e^{-\lambda_1 t / \theta_1} t^{\mu / \theta_1} (1 - t)^{\lambda_2 (1 - y) / \theta_1 - 1} dt,$$

$\tilde{B}_C(s)$ is the LST of the busy period of the $M/M/1+M$ subsystem of class-1 given by Equation (3.3), and $\pi(0, 0)$ is given by Corollary 4 below.

Proof. The proof is structured in three stages.

Stage 1: Reduction to a first-order ODE in x and the integrating factor. Substituting $\theta_1 > 0$, $\gamma_1 = \gamma_2 = \theta_2 = 0$ into the general fundamental equation eliminates both convection terms in y , leaving a first-order linear ODE in x for each fixed $y \in [0, 1]$ (Equation (8.3)). Partial-fraction decomposition of the ODE coefficient gives the integrating factor

$$\mu_I(x; y) = e^{-\lambda_1 x / \theta_1} x^\alpha (1 - x)^{\beta(y)}, \quad \alpha = \frac{\mu}{\theta_1}, \quad \beta(y) = \frac{\lambda_2 (1 - y)}{\theta_1},$$

which vanishes at $x = 0$ and $x = 1$ because $\alpha, \beta(y) > 0$.

Stage 2: Determination of $P_y(y)$. Two independent routes are given in §§8.2–8.3:

- *Probabilistic (Pollaczek–Khinchine).* Because $\theta_2 = \gamma_i = 0$, the class-2 arrival stream remains Poisson; PASTA applies. Class-2 customers see an $M/G/1$ queue with effective service time B_C (the busy period of the class-1 $M/M/1+M$ subsystem), and the P-K formula yields $P_y(y)$ directly (Equation (8.8)).
- *Analytical (boundedness at $x = 1$).* Since $\mu_I(1; y) = 0$, the only way for $P(x, y)$ to remain bounded as $x \rightarrow 1$ is $\int_0^1 q(t; y) \mu_I(t; y) dt = 0$ (Equation (8.13)). Evaluating this integral via the Kummer–Euler identity (8.15) recovers the same $P_y(y)$.

Stage 3: Recovery of $P(x, y)$. With $P_y(y)$ known, the ODE is integrated against the integrating factor (§8.4), yielding the closed form (8.16) stated in the theorem. $\square$

Corollary 4. *The limiting probability $\pi(0, 0)$ of the joint state $(N_1, N_2) = (0, 0)$ with the server busy, and the empty-system probability π_0 , are given by*

$$\pi(0, 0) = \frac{(\lambda_1 + \lambda_2)(1 - \lambda_2 \mathbb{E}[B_C])}{\mu(1 + \lambda_1 \mathbb{E}[B_C])} = (\rho_1 + \rho_2) \frac{1 - \lambda_2 \mathbb{E}[B_C]}{1 + \lambda_1 \mathbb{E}[B_C]}, \quad (8.1)$$

$$\pi_0 = \frac{1 - \lambda_2 \mathbb{E}[B_C]}{1 + \lambda_1 \mathbb{E}[B_C]}, \quad (8.2)$$

where $\mathbb{E}[B_C]$ is the mean busy period of the $M/M/1 + M$ subsystem of class-1, given by Equation (3.4). In the limit $\theta_1 \rightarrow 0^+$, $\mathbb{E}[B_C] \rightarrow (\mu - \lambda_1)^{-1}$ and the expressions above recover the classical Model-A results $\pi(0, 0) = \rho(1 - \rho)$ and $\pi_0 = 1 - \rho$.

8.1 Setup: reduction to a first-order linear ODE in x

Plugging $\theta_1 > 0$, $\gamma_1 = \gamma_2 = \theta_2 = 0$ into the general fundamental equation (4.14) (Lemma 1) annihilates the $\partial P / \partial y$ term entirely, since its coefficient $xy[\gamma_2(y - x) + \theta_2(y - 1)]$ vanishes identically, and reduces the $\partial P / \partial x$ coefficient $xy[\gamma_1(x - y) + \theta_1(x - 1)]$ to its single surviving abandonment term $\theta_1 xy(x - 1)$. What remains is

$$[(\lambda_1 + \lambda_2 + \mu)xy - \mu y - \lambda_1 x^2 y - \lambda_2 x y^2] P(x, y) + \theta_1 xy(x - 1) \frac{\partial}{\partial x} P(x, y) \\ = \mu(x - y) P_y(y) - \mu x(1 - y) \pi(0, 0), \quad (8.3)$$

where the boundary function $P_y(y) = P(0, y) = \sum_{n_2=0}^{\infty} \pi(0, n_2) y^{n_2}$ is, as in Model-A, the partial PGF restricted to states where the server is busy and the priority queue is empty ($N_1 = 0$).

For any fixed y , Equation (8.3) is a *first-order linear ODE in x* whose unknowns are the function $P(x, y)$ on $[0, 1]^2$ and the boundary trace $P_y(y) = P(0, y)$. The strategy in what follows is in two stages:

1. Identify $P_y(y)$ — first by a probabilistic argument exploiting the $M/G/1$ structure seen by class-2 (§8.2), then re-derived analytically from boundedness of $P(x, y)$ at $x = 1$ (§8.3).
2. Recover $P(x, y)$ by solving the ODE with $P_y(y)$ now known (§8.4).

We close the setup with the conditional-expectation decomposition that will drive the probabilistic derivation. By the law of total expectation, restricting the defining sum of $P_y(y) = \mathbb{E}[\mathbf{1}\{\text{busy}, N_1 = 0\} y^{N_2}]$ to the event $\{\text{busy}, N_1 = 0\}$ factors it into the mass of that event and the conditional PGF of N_2 given it:

$$P_y(y) = \mathbb{P}(\text{busy}, N_1 = 0) \cdot \mathbb{E}[y^{N_2} \mid \text{busy}, N_1 = 0]. \quad (8.4)$$

8.2 Probabilistic derivation of $P_y(y)$

The idea is, again, to adopt the perspective of a class-2 customer waiting in line while the server is busy. Unlike in the case of jockeying, the customers from class-1 truly abandon the system, leaving the Poisson arrival process of class-2 undisturbed. Because class-1 holds non-preemptive priority (§3.2), a waiting class-2 customer is not taken into service while any class-1 customer is present ($N_1 > 0$); its *effective service time* is therefore the duration of one complete busy period B_C of the class-1 subsystem—now an $M/M/1 + M$ queue with arrival rate λ_1 , service rate μ and per-customer abandonment rate θ_1 , rather than the pure $M/M/1$ of Model-A. By the *memorylessness* of the exponential service, arrival and abandonment clocks, these effective service times are i.i.d., so from class-2's viewpoint the system is an $M/G/1$ queue with Poisson arrivals at rate λ_2 and general service times distributed as B_C .

The LST $\tilde{B}_C(s)$ and mean $\mathbb{E}[B_C]$ of this busy period are derived via the first-step analysis in §3.3; see Equations (3.3)–(3.4).

The stability condition required for the class-2 subsystem is that of the $M/G/1$ defined above, i.e. $\lambda_2 \mathbb{E}[B_C] < 1$:

$$\rho_2 \cdot {}_1F_1\left(1; \frac{\mu}{\theta_1} + 1; \frac{\lambda_1}{\theta_1}\right) < 1. \quad (8.5)$$

This condition is also necessary, not merely sufficient. The quantity $\lambda_2 \mathbb{E}[B_C]$ is precisely the utilisation of the embedded $M/G/1$ queue of class-2 customers (Poisson input λ_2 , effective service time B_C), so when $\lambda_2 \mathbb{E}[B_C] \geq 1$ that $M/G/1$ is null recurrent (at equality) or transient (strictly above) and its number-in-system has no proper stationary law. Because the class-2 population is a deterministic function of the joint state (n_1, n_2) , a divergent class-2 queue rules out positive recurrence of the joint chain, so (8.5) is sharp. We argue this at the level of the embedded $M/G/1$, whose marginal dynamics are exact, rather than by constructing a Lyapunov function on the joint generator. Compared to Model-A, this stability condition is strictly weaker than the classical $\rho_1 + \rho_2 < 1$, which aligns with the fact that abandonments lower the load of the system: for all $\theta_1 > 0$ the series is termwise dominated by the geometric series for $1/(1 - \rho_1)$. Indeed, since $1^{(n)} = n!$ the Kummer series (3.9) reduces to

$${}_1F_1\left(1; \frac{\mu}{\theta_1} + 1; \frac{\lambda_1}{\theta_1}\right) = \sum_{n=0}^{\infty} \frac{(\lambda_1/\theta_1)^n}{\left(\frac{\mu}{\theta_1} + 1\right)^{(n)}} = \sum_{n=0}^{\infty} \frac{(\lambda_1/\theta_1)^n}{\prod_{k=1}^n \left(\frac{\mu}{\theta_1} + k\right)},$$

and since $\frac{\mu}{\theta_1} + k > \frac{\mu}{\theta_1}$ for every $k \geq 1$, the n -th term is bounded above by $(\lambda_1/\theta_1)^n / (\mu/\theta_1)^n = \rho_1^n$, with strict inequality for $n \geq 1$. Summing gives ${}_1F_1(1; \mu/\theta_1 + 1; \lambda_1/\theta_1) < \sum_{n=0}^{\infty} \rho_1^n = 1/(1 - \rho_1)$. Abandonments in the priority queue therefore stabilise the non-priority queue.

The class-2 subsystem is therefore an $M/G/1$ queue with Poisson arrival rate λ_2 , service-time LST $\tilde{B}_C(s)$ and utilisation $\rho_B = \lambda_2 \mathbb{E}[B_C]$. Applying the Pollaczek–Khinchine formula (3.8) (§3.4) with these identifications ($\lambda \mapsto \lambda_2$, $\tilde{B} \mapsto \tilde{B}_C$, $\rho_B = \lambda_2 \mathbb{E}[B_C]$) gives the PGF of the stationary number of class-2 customers in that $M/G/1$:

$$L_2(y) = \frac{(1 - \lambda_2 \mathbb{E}[B_C]) (1 - y) \tilde{B}_C(\lambda_2(1 - y))}{\tilde{B}_C(\lambda_2(1 - y)) - y}.$$

Because $\theta_2 = \gamma_i = 0$, the class-2 arrival stream is Poisson and satisfies the lack-of-anticipation assumption, so by the *PASTA* property (§3.2) the time-stationary distribution of the number of class-2 customers present coincides with the $M/G/1$ steady-state distribution. Conditioning on $\{\text{busy}, N_1 = 0\}$ — the event on which a class-2 customer is at the head of its queue — therefore identifies that count with N_2 , giving

$$\mathbb{E}[y^{N_2} \mid \text{busy}, N_1 = 0] = L_2(y). \quad (8.6)$$

Substituting the conditional identification (8.6) into the law-of-total-expectation decomposition (8.4) yields

$$P_y(y) = \mathbb{P}(\text{busy}, N_1 = 0) \left[\frac{(1 - \lambda_2 \mathbb{E}[B_C]) (1 - y) \tilde{B}_C(\lambda_2(1 - y))}{\tilde{B}_C(\lambda_2(1 - y)) - y} \right]. \quad (8.7)$$

By the definition of P_y , $P_y(0) = \pi(0, 0)$. Evaluating (8.7) at $y = 0$,

$$\begin{aligned} P_y(0) &= \pi(0, 0) = \mathbb{P}(\text{busy}, N_1 = 0) \left[\frac{(1 - \lambda_2 \mathbb{E}[B_C]) \cdot 1 \cdot \tilde{B}_C(\lambda_2)}{\tilde{B}_C(\lambda_2)} \right] \\ &= \mathbb{P}(\text{busy}, N_1 = 0) (1 - \lambda_2 \mathbb{E}[B_C]), \end{aligned}$$

so that $\mathbb{P}(\text{busy}, N_1 = 0) = \pi(0, 0)/(1 - \lambda_2 \mathbb{E}[B_C])$. Substituting this back into (8.7), the $(1 - \lambda_2 \mathbb{E}[B_C])$ factor cancels and we are left with

$$P_y(y) = \pi(0, 0) \frac{(1 - y) \tilde{B}_C(\lambda_2(1 - y))}{\tilde{B}_C(\lambda_2(1 - y)) - y}. \quad (8.8)$$

This is structurally the same expression as in the probabilistic argument for Model-A (Equation (5.10)), with the only difference being the effective service time: B_C is now the busy period of an $M/M/1 + M$ subsystem with abandonment rate θ_1 , rather than the busy period of a pure $M/M/1$.

8.3 Analytical re-derivation of $P_y(y)$ via boundedness at $x = 1$

Here the surviving abandonment factor $(x - 1)$ places the operative singularity on the boundary $x = 1$, exactly the boundedness-pinning regime anticipated in the remark following Lemma 1. For an analytic counterpart of the previous derivation, we put (8.3) in standard form:

$$\frac{\partial P}{\partial x} + \underbrace{\frac{(\lambda_1 + \lambda_2 + \mu)x - \mu - \lambda_1 x^2 - \lambda_2 xy}{\theta_1 x(x - 1)}}_{p(x; y)} P = \underbrace{\frac{\mu(x - y) P_y(y) - \mu x(1 - y) \pi(0, 0)}{\theta_1 xy(x - 1)}}_{q(x; y)}.$$

The numerator of $p(x; y)$ is the same polynomial we have seen in every model so far. With $x^*(y)$ and $x^+(y)$ as its roots (Equation (5.8)), we can write

$$(\lambda_1 + \lambda_2 + \mu)x - \mu - \lambda_1 x^2 - \lambda_2 xy = -\lambda_1 (x - x^*(y))(x - x^+(y)),$$

with $x^*(y) x^+(y) = \mu/\lambda_1$. Partial fraction decomposition gives

$$\frac{-\lambda_1 (x - x^*(y))(x - x^+(y))}{x(x - 1)} = A + \frac{B}{x} + \frac{C}{x - 1},$$

where $A = -\lambda_1$ is read off from the ratio of leading coefficients of the quadratic terms in numerator and denominator, $B = \mu$ from setting $x = 0$ and using $\lambda_1 x^* x^+ = \mu$, and $C = \lambda_2(1 - y)$ by the cover-up rule at $x = 1$. Writing the numerator as $N(x) = (\lambda_1 + \lambda_2 + \mu)x - \mu - \lambda_1 x^2 - \lambda_2 xy$, the coefficient of $1/(x - 1)$ is its residue $C = N(1) = (\lambda_1 + \lambda_2 + \mu) - \mu - \lambda_1 - \lambda_2 y = \lambda_2(1 - y)$; equivalently, since $N(x) = -\lambda_1 (x - x^*(y))(x - x^+(y))$, this is the identity $(1 - x^*(y))(1 - x^+(y)) = -\lambda_2(1 - y)/\lambda_1$ evaluated at $x = 1$. Thus

$$p(x; y) = \frac{1}{\theta_1} \left[-\lambda_1 + \frac{\mu}{x} + \frac{\lambda_2(1 - y)}{x - 1} \right]. \quad (8.9)$$

Integrating term by term,

$$\int p(x; y) dx = \frac{1}{\theta_1} [-\lambda_1 x + \mu \ln x + \lambda_2(1 - y) \ln(1 - x)],$$

which yields the integrating factor

$$\mu_I(x; y) = e^{-\lambda_1 x / \theta_1} \cdot x^\alpha \cdot (1 - x)^{\beta(y)}, \quad (8.10)$$

where, in parallel with Model- B_2 , we set

$$\alpha = \frac{\mu}{\theta_1}, \quad \beta(y) = \frac{\lambda_2(1 - y)}{\theta_1}. \quad (8.11)$$

Under our stability assumptions both exponents are positive ($\alpha > 0$ always; $\beta(y) \geq 0$ for $y \in [0, 1]$), so μ_I vanishes at both $x = 0$ and $x = 1$. From $\frac{d}{dx}[P \cdot \mu_I] = q \cdot \mu_I$, integration over $(0, x]$ reads

$$P(x, y) \mu_I(x; y) - P(0, y) \mu_I(0; y) = \int_0^x q(t; y) \mu_I(t; y) dt,$$

and the boundary term vanishes, yielding

$$P(x, y) = \frac{1}{\mu_I(x; y)} \int_0^x q(t; y) \mu_I(t; y) dt. \quad (8.12)$$

Now, in the same fashion as the root-cancellation argument used for Model- A , we exploit boundedness of PGFs to determine $P_y(y)$ by observing the behaviour of $P(x, y)$ as $x \rightarrow 1$, for fixed $y \in [0, 1)$: there $\beta(y) > 0$, so $\mu_I(x; y) \rightarrow 0$ as $x \rightarrow 1$. (At the endpoint $y = 1$ the exponent $\beta(1) = 0$ and $\mu_I(\cdot; 1)$ no longer vanishes at $x = 1$; the value $P_y(1)$ is recovered from the resulting formula by continuity, $y \rightarrow 1^-$.) Since $\mu_I(x; y) \rightarrow 0$, the only way to keep the right-hand side of (8.12) bounded is

$$\int_0^1 q(t; y) \mu_I(t; y) dt = 0. \quad (8.13)$$

Substituting the explicit forms of q and μ_I into (8.13), the integrand is

$$q(t; y) \mu_I(t; y) = \frac{\mu[(t - y) P_y(y) - t(1 - y) \pi(0, 0)]}{\theta_1 y} \cdot \frac{e^{-\lambda_1 t / \theta_1} t^\alpha (1 - t)^{\beta(y)}}{t(t - 1)},$$

where $t^\alpha / t = t^{\alpha-1}$ and, with $b = \beta(y)$, the elementary identity $(1 - t)^{\beta(y)} / (t - 1) = -(1 - t)^{\beta(y)-1}$ flips the sign and lowers the exponent. The constant prefactor $-\mu / (\theta_1 y)$, nonzero for $y \in (0, 1]$, divides out of the vanishing condition (8.13), leaving

$$\int_0^1 e^{-\lambda_1 t / \theta_1} t^{\alpha-1} (1 - t)^{\beta(y)-1} [(t - y) P_y(y) - t(1 - y) \pi(0, 0)] dt = 0,$$

with α and $\beta(y)$ as in (8.11). Introducing the auxiliary integrals

$$\begin{aligned} \mathcal{J}_0(y) &= \int_0^1 e^{-\lambda_1 t / \theta_1} t^{\alpha-1} (1 - t)^{\beta(y)-1} dt, \\ \mathcal{J}_1(y) &= \int_0^1 e^{-\lambda_1 t / \theta_1} t^\alpha (1 - t)^{\beta(y)-1} dt, \end{aligned}$$

and distributing the bracket over $t^{\alpha-1}$, note that $t \cdot t^{\alpha-1} = t^\alpha$ feeds $\mathcal{J}_1$ while $t^{\alpha-1}$ feeds $\mathcal{J}_0$. The two terms regroup as

$$P_y(y) \int_0^1 e^{-\lambda_1 t / \theta_1} (t^\alpha - y t^{\alpha-1}) (1 - t)^{\beta(y)-1} dt - (1 - y) \pi(0, 0) \int_0^1 e^{-\lambda_1 t / \theta_1} t^\alpha (1 - t)^{\beta(y)-1} dt = 0,$$

so the boundedness condition (8.13) becomes

$$P_y(y) [\mathcal{J}_1(y) - y \mathcal{J}_0(y)] = (1 - y) \pi(0, 0) \mathcal{J}_1(y),$$

which gives

$$P_y(y) = \pi(0, 0) (1 - y) \frac{\mathcal{J}_1(y)}{\mathcal{J}_1(y) - y \mathcal{J}_0(y)}. \quad (8.14)$$

The continued-fraction representation (3.3) admits an equivalent integral form. Setting $a = \alpha = \mu / \theta_1$, $b = s / \theta_1$, $c = \lambda_1 / \theta_1$, the Euler integral representation (3.10) identifies each of the integrals below as a Beta-function multiple of a ${}_1F_1$ value:

$$\tilde{B}_C(s) = \frac{\int_0^1 t^a (1 - t)^{b-1} e^{-ct} dt}{\int_0^1 t^{a-1} (1 - t)^{b-1} e^{-ct} dt} = \frac{B(a + 1, b) {}_1F_1(a + 1; a + 1 + b; -c)}{B(a, b) {}_1F_1(a; a + b; -c)}. \quad (8.15)$$

That this ratio equals the continued fraction (3.3) is exactly the closed-form summation (3.7) proved in §3.3: the integrals satisfy the three-term recurrence (3.6), they form its minimal solution, and Pincherle's theorem identifies the convergent continued fraction with the minimal-solution ratio.

Setting $s = \lambda_2(1 - y)$ — so that $b = \beta(y)$ matches the exponent in (8.11) — yields exactly the ratio of our $\mathcal{J}_i$:

$$\tilde{B}_C(\lambda_2(1 - y)) = \frac{\mathcal{J}_1(y)}{\mathcal{J}_0(y)}.$$

Dividing numerator and denominator of (8.14) by $\mathcal{J}_0(y)$ then recovers exactly Equation (8.8). The boundedness route, which uses no probabilistic structure, reaches the same Pollaczek–Khinchine form only through the Beta-integral identity (8.15), which is the analytic shadow of the $M/G/1$ /busy-period decomposition used in §8.2.

8.4 Recovery of $P(x, y)$ from $P_y(y)$

With $P_y(y)$ now in hand, we return to the ODE (8.3) and substitute (8.8) into its right-hand side. Writing $\zeta(y) \equiv \tilde{B}_C(\lambda_2(1 - y))$ for brevity,

$$\begin{aligned} \text{RHS} &= \mu(x - y) \pi(0, 0) \frac{(1 - y) \zeta(y)}{\zeta(y) - y} - \mu x(1 - y) \pi(0, 0) \\ &= \mu \pi(0, 0) (1 - y) \left[\frac{(x - y) \zeta(y)}{\zeta(y) - y} - x \right] \\ &= \mu \pi(0, 0) (1 - y) \frac{(x - y) \zeta(y) - x[\zeta(y) - y]}{\zeta(y) - y} \\ &= \mu \pi(0, 0) y(1 - y) \frac{x - \zeta(y)}{\zeta(y) - y}, \end{aligned}$$

which is structurally identical to the RHS found for Model-A, though now with the more involved effective service time B_C . With $p(x; y)$ and $\mu_I(x; y)$ as in (8.9) and (8.10), the source $q(x; y)$ of the standard-form ODE now reads

$$q(x; y) = \frac{\mu \pi(0, 0) (1 - y) [x - \zeta(y)]}{\theta_1 x(x - 1) [\zeta(y) - y]},$$

and the implicit solution (8.12) becomes

$$P(x, y) = \frac{1}{\mu_I(x; y)} \int_0^x q(t; y) \mu_I(t; y) dt.$$

To evaluate the integral, we use the partial-fraction identity

$$\frac{t - \zeta(y)}{t(t - 1)} = \frac{\zeta(y)(t - 1) + t(1 - \zeta(y))}{t(t - 1)} = \frac{\zeta(y)}{t} + \frac{1 - \zeta(y)}{t - 1},$$

and define the auxiliary integrals

$$\begin{aligned} \mathcal{H}_1(x, y) &= \int_0^x e^{-\lambda_1 t / \theta_1} t^{\alpha-1} (1 - t)^{\beta(y)} dt, \\ \mathcal{H}_2(x, y) &= \int_0^x e^{-\lambda_1 t / \theta_1} t^{\alpha} (1 - t)^{\beta(y)-1} dt, \end{aligned}$$

with α and $\beta(y)$ as in (8.11).

Note on the $\mathcal{J}$ – $\mathcal{H}$ correspondence. The integrals $\mathcal{H}_1, \mathcal{H}_2$ are the partial (upper limit x) counterparts of the full integrals $\mathcal{J}_0, \mathcal{J}_1$ from (8.13). Evaluating at $x = 1$ gives $\mathcal{H}_2(1, y) = \mathcal{J}_1(y)$ exactly. However, $\mathcal{H}_1(1, y)$ is *not* $\mathcal{J}_0(y)$: the former has the exponent $\beta(y)$ on $(1 - t)$ while the latter has $\beta(y) - 1$, which under the Euler representation (3.10) corresponds to incrementing the second Kummer parameter by one. Both families are evaluated by the same representation; the exponent shift is harmless because $\mathcal{H}_1$ and $\mathcal{H}_2$ always appear together in the combination $\zeta(y) \mathcal{H}_1 - (1 - \zeta(y)) \mathcal{H}_2$, not individually.

Multiplying the source $q(t; y)$ by the integrating factor $\mu_I(t; y) = e^{-\lambda_1 t / \theta_1} t^{\alpha} (1 - t)^{\beta(y)}$ and inserting the partial fraction, the integrand splits into the two pieces defining $\mathcal{H}_1, \mathcal{H}_2$:

$$q(t; y) \mu_I(t; y) = \frac{\mu \pi(0, 0) (1 - y)}{\theta_1 (\zeta(y) - y)}$$

$$\begin{aligned}
& \times \left[\frac{\zeta(y)}{t} + \frac{1 - \zeta(y)}{t - 1} \right] e^{-\lambda_1 t / \theta_1} t^\alpha (1 - t)^{\beta(y)} \\
& = \frac{\mu \pi(0, 0) (1 - y)}{\theta_1 (\zeta(y) - y)} \left[\zeta(y) e^{-\lambda_1 t / \theta_1} t^{\alpha-1} (1 - t)^{\beta(y)} - (1 - \zeta(y)) e^{-\lambda_1 t / \theta_1} t^\alpha (1 - t)^{\beta(y)-1} \right],
\end{aligned}$$

where the second term used $(1 - t)^{\beta(y)} / (t - 1) = -(1 - t)^{\beta(y)-1}$, the same identity used to simplify the boundedness condition (8.13); this is the source of both the minus sign and the exponent shift $\beta(y) \rightarrow \beta(y) - 1$ on $\mathcal{H}_2$. Since $\pi(0, 0)$ and $\zeta(y)$ come out of the integral,

$$\int_0^x q(t; y) \mu_I(t; y) dt = \frac{\mu \pi(0, 0) (1 - y)}{\theta_1 (\zeta(y) - y)} \left[\zeta(y) \mathcal{H}_1(x, y) - (1 - \zeta(y)) \mathcal{H}_2(x, y) \right].$$

Dividing by $\mu_I(x; y)$ and restoring $\zeta(y) = \tilde{B}_C(\lambda_2(1 - y))$, we arrive at the closed form

$$\begin{aligned}
P(x, y) &= \frac{e^{\lambda_1 x / \theta_1} \mu \pi(0, 0) (1 - y)}{\theta_1 x^\alpha (1 - x)^{\beta(y)} (\tilde{B}_C(\lambda_2(1 - y)) - y)} \\
&\times \left[\tilde{B}_C(\lambda_2(1 - y)) \mathcal{H}_1(x, y) - (1 - \tilde{B}_C(\lambda_2(1 - y))) \mathcal{H}_2(x, y) \right],
\end{aligned} \tag{8.16}$$

which is the statement of the theorem. This closed form is the product of a probabilistic derivation of $P_y(y)$ via Pollaczek–Khinchine, combined with an analytic integration of the ODE in x ; the same result is reached from the purely analytic route of §8.3.

Proof of Corollary 4. The fully idle state—the server empty with no customer present, of probability π_0 —communicates with the rest of the chain through a single neighbour: the busy-and-empty state $(N_1, N_2) = (0, 0)$ of probability $\pi(0, 0)$, in which exactly one customer is in service and both queues are empty (recall the standing convention $\pi_0 \neq \pi(0, 0)$). The system leaves the idle state only upon an arrival of either class, at combined Poisson rate $\lambda_1 + \lambda_2$ (Poisson superposition), and re-enters it only upon the service completion that empties the busy-and-empty state, at rate μ ; no abandonment can fire from $\pi(0, 0)$, since both queues are empty. Equating the probability flow across the cut separating the idle state from its single neighbour (global balance for the idle state, as in (4.7)) therefore gives

$$(\lambda_1 + \lambda_2) \pi_0 = \mu \pi(0, 0), \tag{8.17}$$

so $\pi(0, 0) = (\rho_1 + \rho_2) \pi_0$.

To fix π_0 , we exploit the $M/G/1$ structure used in §8.2. Class-2 customers form an $M/G/1$ queue with arrival rate λ_2 and effective service time B_C . By the standard $M/G/1$ work-conservation (utilisation) identity, the long-run fraction of time this $M/G/1$ is busy is $\rho_B = \lambda_2 \mathbb{E}[B_C]$, so its idle probability is $1 - \rho_B$. The event “this $M/G/1$ is idle” coincides with the absence of any class-2 customer from the system, so

$$\mathbb{P}(\text{no class-2 present}) = 1 - \lambda_2 \mathbb{E}[B_C].$$

Conditional on no class-2 being present, the priority subsystem evolves freely as the same $M/M/1+M$ queue used to define B_C , with arrival rate λ_1 , service rate μ and abandonment rate θ_1 . This subsystem alternates between idle periods (no class-1 customer present) and busy periods B_C , the successive idle+busy spans forming an alternating renewal process whose regeneration epochs are the instants at which the subsystem empties. An idle period lasts until the next class-1 arrival; by the memorylessness of the exponential interarrival clock this duration is $\text{Exp}(\lambda_1)$, with mean $1/\lambda_1$. The mean regeneration cycle is therefore $1/\lambda_1 + \mathbb{E}[B_C]$, and by the renewal–reward theorem the stationary fraction of time the subsystem is idle equals the ratio of the mean idle period to the mean cycle:

$$\mathbb{P}(\text{class-1 absent} \mid \text{no class-2 present}) = \frac{1/\lambda_1}{1/\lambda_1 + \mathbb{E}[B_C]} = \frac{1}{1 + \lambda_1 \mathbb{E}[B_C]}.$$

The fully idle event $\{\text{server idle}, N_1 = N_2 = 0\}$ is exactly the intersection $A \cap B$ of the events $B = \{\text{no class-2 present}\}$ and $A = \{\text{class-1 absent}\}$, since the server is idle precisely when no customer of either class is in the system. By the multiplication (chain) rule $\mathbb{P}(A \cap B) = \mathbb{P}(B) \mathbb{P}(A \mid B)$ —using the conditional idle fraction just computed for the factor $\mathbb{P}(A \mid B)$, and *not* independence of A and B —we obtain

$$\pi_0 = (1 - \lambda_2 \mathbb{E}[B_C]) \cdot \frac{1}{1 + \lambda_1 \mathbb{E}[B_C]} = \frac{1 - \lambda_2 \mathbb{E}[B_C]}{1 + \lambda_1 \mathbb{E}[B_C]},$$

which is (8.2). Substituting into (8.17) yields (8.1).

For the Model-A limit, $\theta_1 \rightarrow 0^+$ gives $\mathbb{E}[B_C] \rightarrow (\mu - \lambda_1)^{-1}$ (the busy period of the $M/M/1 + M$ subsystem collapses to that of the pure $M/M/1$ priority subsystem). Then

$$\lambda_1 \mathbb{E}[B_C] \rightarrow \frac{\lambda_1}{\mu - \lambda_1} = \frac{\rho_1}{1 - \rho_1}, \quad \lambda_2 \mathbb{E}[B_C] \rightarrow \frac{\lambda_2}{\mu - \lambda_1} = \frac{\rho_2}{1 - \rho_1},$$

so that the denominator $1 + \lambda_1 \mathbb{E}[B_C] \rightarrow 1/(1 - \rho_1)$ and the numerator

$$1 - \lambda_2 \mathbb{E}[B_C] \rightarrow 1 - \frac{\rho_2}{1 - \rho_1} = \frac{1 - \rho_1 - \rho_2}{1 - \rho_1} = \frac{1 - \rho}{1 - \rho_1}.$$

Substituting into (8.2),

$$\pi_0 \rightarrow \frac{(1 - \rho)/(1 - \rho_1)}{1/(1 - \rho_1)} = 1 - \rho,$$

and then $\pi(0, 0) = (\rho_1 + \rho_2) \pi_0 \rightarrow \rho(1 - \rho)$ via (8.17), recovering the classical Model-A baseline values. $\square$

8.5 Mean queue lengths and mean waiting times

Why the diagonal no longer simplifies

Setting $x = y = z$ in the fundamental equation (8.3), the abandonment convection $\theta_1 xy(x - 1) \partial_x P$ does *not* vanish (since $(x - 1)|_{x=z} \neq 0$ for $z \in (0, 1)$). The diagonal equation reads:

$$\begin{aligned} & [(\lambda_1 + \lambda_2 + \mu)z^2 - \mu z - \lambda z^3] P(z, z) + \theta_1 z^2(z - 1) \frac{\partial P}{\partial x}(z, z) \\ & = -\mu z(1 - z) \pi(0, 0), \end{aligned} \quad (8.18)$$

Comparison with Model-A is instructive. There the same diagonal substitution $x = y = z$ also annihilates the boundary term $\mu(x - y) P_y(y)$, but with $\theta_1 = \gamma_1 = \gamma_2 = 0$ there is *no* convection term at all, so the diagonal equation reduced to a purely algebraic relation in the single unknown $P(z, z)$, with closed form $P(z, z) = \pi(0, 0)/(1 - \rho z)$ (Equation (5.23)), from which $\mathbb{E}[N_1] + \mathbb{E}[N_2]$ followed by a single differentiation. Here the surviving convection $\theta_1 z^2(z - 1) \partial_x P(z, z)$ injects the off-diagonal slope $\partial_x P(z, z)$ as a *second* unknown into the lone diagonal relation, and no companion equation closes the system; consequently no closed form for $P(z, z)$ —and hence for the total $\mathbb{E}[N_1] + \mathbb{E}[N_2]$ —is available. Abandonment removes customers and reduces both means relative to Model-A.

Class-2 mean queue length via $P(1, y)$

Setting $x = 1$ in the fundamental equation (8.3), the abandonment term $\theta_1 xy(x - 1) \partial_x P$ vanishes identically, and the coefficient of $P(1, y)$ simplifies to $\lambda_2 y(1 - y)$:

$$\lambda_2 y(1 - y) P(1, y) = \mu(1 - y) [P_y(y) - \pi(0, 0)].$$

Cancelling $(1 - y)$ for $y \neq 1$ and substituting the known $P_y(y)$ (8.8):

$$P(1, y) = \frac{\mu [P_y(y) - \pi(0, 0)]}{\lambda_2 y} = \frac{\mu \pi(0, 0)}{\lambda_2} \cdot \frac{1 - \tilde{B}_C(\lambda_2(1 - y))}{\tilde{B}_C(\lambda_2(1 - y)) - y}.$$

This is the marginal PGF of N_2 . To differentiate at $y = 1$ (where the expression is 0/0), set $\varepsilon = 1 - y \rightarrow 0$ and use the LST expansion $\tilde{B}_C(\lambda_2 \varepsilon) = 1 - \lambda_2 \mathbb{E}[B_C] \varepsilon + \frac{1}{2} \lambda_2^2 \mathbb{E}[B_C^2] \varepsilon^2 + O(\varepsilon^3)$:

$$\begin{aligned} 1 - \tilde{B}_C(\lambda_2 \varepsilon) &= \lambda_2 \mathbb{E}[B_C] \varepsilon - \frac{1}{2} \lambda_2^2 \mathbb{E}[B_C^2] \varepsilon^2 + O(\varepsilon^3), \\ \tilde{B}_C(\lambda_2 \varepsilon) - y &= (1 - \lambda_2 \mathbb{E}[B_C] \varepsilon + \frac{1}{2} \lambda_2^2 \mathbb{E}[B_C^2] \varepsilon^2) - (1 - \varepsilon) \\ &= (1 - \lambda_2 \mathbb{E}[B_C]) \varepsilon + \frac{1}{2} \lambda_2^2 \mathbb{E}[B_C^2] \varepsilon^2 + O(\varepsilon^3), \end{aligned}$$

so that (writing $a = \lambda_2 \mathbb{E}[B_C]$ and $c = \frac{1}{2} \lambda_2^2 \mathbb{E}[B_C^2]$):

$$P(1, 1 - \varepsilon) = \frac{\mu \pi(0, 0)}{\lambda_2} \cdot \frac{a \varepsilon - c \varepsilon^2 + O(\varepsilon^3)}{(1 - a) \varepsilon + c \varepsilon^2 + O(\varepsilon^3)} = \frac{\mu \pi(0, 0)}{\lambda_2} \cdot \frac{a - c \varepsilon}{(1 - a) + c \varepsilon} + O(\varepsilon^2).$$

Differentiating $f(\varepsilon) = (a - c\varepsilon)/((1 - a) + c\varepsilon)$ by the quotient rule:

$$f'(\varepsilon) = \frac{-c((1 - a) + c\varepsilon) - (a - c\varepsilon) \cdot c}{((1 - a) + c\varepsilon)^2} = \frac{-c}{(1 - a + c\varepsilon)^2},$$

and since $d/dy P(1, y)|_{y=1} = -d/d\varepsilon P(1, 1 - \varepsilon)|_{\varepsilon=0}$:

$$\mathbb{E}[N_2] = \frac{d}{dy} P(1, y) \Big|_{y=1} = \frac{\mu \pi(0, 0)}{\lambda_2} \cdot \frac{c}{(1 - a)^2} = \frac{\mu \pi(0, 0) \lambda_2 \mathbb{E}[B_C^2]}{2(1 - \lambda_2 \mathbb{E}[B_C])^2}.$$

As $\theta_1 \rightarrow 0^+$, $\mathbb{E}[B_C] \rightarrow (\mu - \lambda_1)^{-1}$ and $\mathbb{E}[B_C^2] \rightarrow 2\mu/(\mu - \lambda_1)^3$; substituting recovers the Model-A value $\mathbb{E}[N_2] = \rho\rho_2/((1 - \rho_1)(1 - \rho))$.

Class-1 mean queue length

We extract $\mathbb{E}[N_1] = \frac{\partial P}{\partial x} \Big|_{(1,1)}$ by evaluating the theorem formula at $y = 1$ (limit) and then differentiating in x .

Step 1: marginal PGF $P(x, 1)$. As $y \rightarrow 1$, the exponent $\beta(y) = \lambda_2(1 - y)/\theta_1 \rightarrow 0$ so $(1 - x)^{\beta(y)} \rightarrow 1$ for fixed $x \in (0, 1)$, and consequently $\mathcal{H}_1(x, y) \rightarrow \int_0^x e^{-\lambda_1 t/\theta_1} t^{\mu/\theta_1 - 1} dt$ and $\mathcal{H}_2(x, y)$ tends to a finite limit (its integrand $e^{-\lambda_1 t/\theta_1} t^{\mu/\theta_1} (1 - t)^{-1}$ is integrable on $[0, x]$ for $x < 1$). Inside the theorem's bracket the two integrals carry different weights: by the LST expansion of §8.2, $\tilde{B}_C(\lambda_2(1 - y)) \rightarrow \tilde{B}_C(0) = 1$ while $1 - \tilde{B}_C(\lambda_2(1 - y)) = \lambda_2 \mathbb{E}[B_C] (1 - y) + O((1 - y)^2) \rightarrow 0$, so the $\mathcal{H}_2$ term vanishes in the limit and the bracket collapses to $\mathcal{H}_1(x, 1)$ alone. Combining this with the scalar prefactor, whose y -dependent part satisfies $(1 - y)/(\tilde{B}_C(\lambda_2(1 - y)) - y) \rightarrow 1/(1 - \lambda_2 \mathbb{E}[B_C])$ (again from the expansion $\tilde{B}_C(\lambda_2(1 - y)) - y = (1 - \lambda_2 \mathbb{E}[B_C])(1 - y) + O((1 - y)^2)$), the theorem formula gives:

$$P(x, 1) = \frac{\mu \pi(0, 0)}{\theta_1 (1 - \lambda_2 \mathbb{E}[B_C])} \cdot e^{\lambda_1 x/\theta_1} x^{-\mu/\theta_1} \int_0^x e^{-\lambda_1 t/\theta_1} t^{\mu/\theta_1 - 1} dt.$$

Step 2: differentiate $P(x, 1)$ at $x = 1$. Write $P(x, 1) = C \cdot f(x) \cdot \mathcal{K}(x)$ with $C = \mu \pi(0, 0)/(\theta_1 (1 - \lambda_2 \mathbb{E}[B_C]))$, $f(x) = e^{\lambda_1 x/\theta_1} x^{-\mu/\theta_1}$, and $\mathcal{K}(x) = \int_0^x e^{-\lambda_1 t/\theta_1} t^{\mu/\theta_1 - 1} dt$. By the product rule and Leibniz's rule ($d\mathcal{K}/dx = e^{-\lambda_1 x/\theta_1} x^{\mu/\theta_1 - 1}$):

$$\frac{d}{dx} P(x, 1) = C \left[f'(x) \mathcal{K}(x) + f(x) \cdot e^{-\lambda_1 x/\theta_1} x^{\mu/\theta_1 - 1} \right].$$

The logarithmic derivative $f'(x)/f(x) = \lambda_1/\theta_1 - (\mu/\theta_1)/x = (\lambda_1 x - \mu)/(\theta_1 x)$, so:

$$\frac{d}{dx} P(x, 1) = C \left[e^{\lambda_1 x/\theta_1} x^{-\mu/\theta_1} \cdot \frac{\lambda_1 x - \mu}{\theta_1 x} \cdot \mathcal{K}(x) + x^{-1} \right].$$

Step 3: evaluate at $x = 1$.

$$\begin{aligned} \mathbb{E}[N_1] &= \frac{d}{dx} P(x, 1) \Big|_{x=1} \\ &= \frac{\mu \pi(0, 0)}{\theta_1 (1 - \lambda_2 \mathbb{E}[B_C])} \left[e^{\lambda_1/\theta_1} \cdot \frac{\lambda_1 - \mu}{\theta_1} \cdot \mathcal{K}(1) + 1 \right], \end{aligned}$$

where $\mathcal{K}(1) = \int_0^1 e^{-\lambda_1 t/\theta_1} t^{\mu/\theta_1 - 1} dt = (\theta_1/\lambda_1)^{\mu/\theta_1} \gamma(\mu/\theta_1, \lambda_1/\theta_1)$, with $\gamma(\cdot, \cdot)$ the lower incomplete gamma function. The factor $(\lambda_1 - \mu) < 0$ (since $\rho_1 < 1$), so the bracket is positive. No elementary simplification is known; $\mathbb{E}[N_1]$ is evaluated numerically. As $\theta_1 \rightarrow 0^+$, $\mathbb{E}[N_1]$ recovers the Model-A value (5.22).

Little's Law with abandonment

The universal relation $\mathbb{E}[N_i] = \lambda_i \mathbb{E}[W_i]$ holds for any stable system. Here λ_i is the *arrival* rate of all class- i customers, including those that will eventually abandon, so $\mathbb{E}[W_i]$ is the mean time in queue averaged over both served and abandoned customers:

$$\mathbb{E}[W_1] = \frac{\mathbb{E}[N_1]}{\lambda_1}, \quad \mathbb{E}[W_2] = \frac{\mathbb{E}[N_2]}{\lambda_2}.$$

Using the throughput $\mu(1 - \pi_0)$ instead would give the conditional waiting time of customers who complete service only — a different object. As $\theta_1 \rightarrow 0^+$ both means recover the Model-A values (5.22)–(5.26).

9 Analysis of Model- B_2^H ($\gamma_1 > 0$ head-of-line, $\gamma_2 = \theta_1 = \theta_2 = 0$)

Model- B_2^H is a deliberately simplified variant of the one-way jockeying model Model- B_2 . In Model- B_2 *every* waiting class-1 customer may jockey to the class-2 queue, so the aggregate jockeying rate out of a state grows linearly with the queue length, $\gamma_1 n_1$. Here, by contrast, only the *customer at the head of Queue 1* — the one that would be the next class-1 customer to enter service — is allowed to jockey, at the fixed rate γ_1 whenever $n_1 \geq 1$. The aggregate jockeying rate is therefore $\gamma_1 \mathbf{1}_{\{n_1 \geq 1\}}$ rather than $\gamma_1 n_1$.

This single change is decisive for tractability. In Model- B_2 the linear coefficient $\gamma_1 n_1$ produces, upon forming the joint PGF, a term proportional to $xy(x-y)\partial P/\partial x$, turning the fundamental equation into a first-order linear ODE in x that requires the integrating-factor machinery of §7. The flat rate $\gamma_1 \mathbf{1}_{\{n_1 \geq 1\}}$, on the other hand, contributes only the *algebraic* factor $\gamma_1 y(x-y)$ to the coefficient of $P(x, y)$; *no derivative of P appears*. The fundamental equation is then purely algebraic and can be solved by the Kernel Method exactly as in Model- A (§5). The analysis below mirrors the analytical treatment of Model- A step by step, and the resulting PGF turns out to be *structurally identical* to that of Model- A , the only difference being a shift in the characteristic roots.

Stability. Jockeying relocates a waiting customer between the two queues but does not remove anyone from the system; the total number of customers is therefore conserved by the jockeying mechanism, exactly as in Model- A . Since there is no abandonment, the chain is positive recurrent if and only if the total traffic intensity satisfies

$$\rho = \frac{\lambda_1 + \lambda_2}{\mu} < 1, \quad (9.1)$$

with $\rho_1 = \lambda_1/\mu$ governing the priority class.

9.1 Balance equations

Because the jockeying rate is now constant in n_1 rather than linear, the balance equations are *not* obtained by zeroing parameters in the general Lemma 1 (whose jockeying coefficient is $\gamma_1 n_1$); they must be stated directly. The jockeying transition $(n_1, n_2) \rightarrow (n_1 - 1, n_2 + 1)$ occurs at the fixed rate γ_1 whenever $n_1 \geq 1$. On the state space S the global balance equations read

$$\begin{aligned} & [\lambda_1 + \lambda_2 + \mu + \gamma_1] \pi(n_1, n_2) \\ &= \mu \pi(n_1 + 1, n_2) + \lambda_1 \pi(n_1 - 1, n_2) + \lambda_2 \pi(n_1, n_2 - 1) + \gamma_1 \pi(n_1 + 1, n_2 - 1) \end{aligned} \quad n_1 \geq 1, n_2 \geq 1, \quad (9.2)$$

$$[\lambda_1 + \lambda_2 + \mu + \gamma_1] \pi(n_1, 0) = \mu \pi(n_1 + 1, 0) + \lambda_1 \pi(n_1 - 1, 0) \quad n_1 \geq 1, n_2 = 0, \quad (9.3)$$

$$\begin{aligned} & [\lambda_1 + \lambda_2 + \mu] \pi(0, n_2) \\ &= \mu \pi(1, n_2) + \gamma_1 \pi(1, n_2 - 1) + \lambda_2 \pi(0, n_2 - 1) + \mu \pi(0, n_2 + 1) \end{aligned} \quad n_1 = 0, n_2 \geq 1, \quad (9.4)$$

$$\begin{aligned} & [\lambda_1 + \lambda_2 + \mu] \pi(0, 0) = (\lambda_1 + \lambda_2) \pi_0 + \mu \pi(1, 0) + \mu \pi(0, 1), \\ & (\lambda_1 + \lambda_2) \pi_0 = \mu \pi(0, 0). \end{aligned} \quad (9.5)$$

Two structural remarks are in order. First, jockeying is an *out*-transition from any state with $n_1 \geq 1$ (hence the $+\gamma_1$ in the left-hand sides of (9.2) and (9.3)) and an *in*-transition into any state with $n_2 \geq 1$ from the state $(n_1 + 1, n_2 - 1)$ (hence the $\gamma_1 \pi(n_1 + 1, n_2 - 1)$ in (9.2) and the $\gamma_1 \pi(1, n_2 - 1)$ in (9.4)). Second, when $n_2 = 0$ there is no jockeying in-flow, since it would have to originate from the non-existent state $(n_1 + 1, -1)$; and when $n_1 = 0$ there is no jockeying out-flow. The remaining transitions — Poisson arrivals at λ_1, λ_2 and service completions at μ — are exactly those of Model- A .

9.2 Fundamental PGF equation

Lemma 4 (Fundamental equation for Model- B_2^H). *Assume positive recurrence, so that the stationary distribution exists and $P(x, y) = \sum_{n_1, n_2 \geq 0} \pi(n_1, n_2) x^{n_1} y^{n_2}$ is analytic on the closed unit polydisc. Then $P(x, y)$ and the boundary function $P_y(y) = P(0, y)$ satisfy*

$$\begin{aligned} & [(\lambda_1 + \lambda_2 + \mu)xy - \mu y - \lambda_1 x^2 y - \lambda_2 x y^2 + \gamma_1 y(x - y)] P(x, y) \\ &= [\mu(x - y) + \gamma_1 y(x - y)] P_y(y) - \mu x(1 - y) \pi(0, 0). \end{aligned} \quad (9.6)$$

Proof. Apart from the jockeying terms, equations (9.2)–(9.5) coincide with the balance equations of Model-A. Multiplying each balance equation by $x^{n_1}y^{n_2}$, summing over its domain and multiplying through by xy , the Model-A part reproduces exactly the left- and right-hand sides of the Model-A fundamental equation (5.2),

$$[(\lambda_1 + \lambda_2 + \mu)xy - \mu y - \lambda_1 x^2 y - \lambda_2 x y^2] P(x, y) = \mu(x - y)P_y(y) - \mu x(1 - y)\pi(0, 0),$$

so it remains only to collect the jockeying contribution. Its *out-flow* appears in every state with $n_1 \geq 1$:

$$\gamma_1 \sum_{n_1 \geq 1} \sum_{n_2 \geq 0} \pi(n_1, n_2) x^{n_1} y^{n_2} = \gamma_1 [P(x, y) - P_y(y)],$$

while its *in-flow* enters every state with $n_2 \geq 1$ from $(n_1 + 1, n_2 - 1)$:

$$\gamma_1 \sum_{n_1 \geq 0} \sum_{n_2 \geq 1} \pi(n_1 + 1, n_2 - 1) x^{n_1} y^{n_2} = \gamma_1 \frac{y}{x} \sum_{m_1 \geq 1} \sum_{m_2 \geq 0} \pi(m_1, m_2) x^{m_1} y^{m_2} = \gamma_1 \frac{y}{x} [P(x, y) - P_y(y)],$$

where we substituted $m_1 = n_1 + 1$, $m_2 = n_2 - 1$. The net jockeying contribution (out minus in) is therefore $\gamma_1(1 - \frac{y}{x})[P(x, y) - P_y(y)]$; multiplying by xy gives $\gamma_1 y(x - y)[P(x, y) - P_y(y)]$. Adding this to the Model-A equation and moving the P_y part to the right-hand side yields (9.6). $\square$

Remark. Equation (9.6) contains *no derivative* of P . The flat jockeying rate has contributed the algebraic factor $\gamma_1 y(x - y)$ to the coefficient of $P(x, y)$ where Model-B₂, with its linear rate $\gamma_1 n_1$, contributes the convective term $\gamma_1 xy(x - y) \partial P / \partial x$. Setting $\gamma_1 = 0$ recovers (5.2) verbatim.

9.3 Closed-form solution

The following theorem is the main result of this section; its proof follows the analytical (Kernel Method) treatment of Model-A in §5.1.1.

Theorem 4 (Closed-form PGF for Model-B₂^H). *Consider the two-class non-preemptive priority queue with Poisson arrival rates λ_1, λ_2 , common exponential service rate μ , and head-of-line class-1 jockeying at rate $\gamma_1 > 0$ (with $\gamma_2 = \theta_1 = \theta_2 = 0$). Under the stability condition (9.1), the joint PGF $P(x, y) = \mathbb{E}[x^{N_1} y^{N_2}]$ is the rational function*

$$P(x, y) = \frac{\mu \rho(1 - \rho)(1 - y)}{\lambda_1 (x^*(y) - y)(x^+(y) - x)}, \quad (9.7)$$

where $x^*(y)$ and $x^+(y)$ are the two roots of the kernel quadratic

$$\lambda_1 x^2 - [\lambda_1 + \lambda_2(1 - y) + \mu + \gamma_1]x + (\mu + \gamma_1 y) = 0, \quad (9.8)$$

with $x^*(y)$ the root inside the unit disk. This is the same rational form as Model-A (Theorem 1); the two differ only in the kernel, where μ has been replaced by $\mu + \gamma_1 y$.

Corollary 5. Under (9.1), the empty-state probabilities of Model-B₂^H coincide with those of Model-A,

$$\pi_0 = 1 - \rho, \quad \pi(0, 0) = \rho(1 - \rho).$$

In particular $P(1, 1) = 1 - \pi_0 = \rho$ and the diagonal generating function satisfies $P(z, z) = \pi(0, 0)/(1 - \rho z)$.

Proof of Corollary 5. Set $x = y = z$ in the fundamental equation (9.6). The jockeying factor $\gamma_1 y(x - y)$ vanishes on the diagonal — a direct consequence of the fact that jockeying conserves the total number of customers — and so does the term $\mu(x - y)P_y(y)$, leaving

$$[(\lambda_1 + \lambda_2 + \mu)z^2 - \mu z - (\lambda_1 + \lambda_2)z^3]P(z, z) = -\mu z(1 - z)\pi(0, 0).$$

Dividing by z and factoring the bracket as $-(\lambda z - \mu)(z - 1)$, with $\lambda = \lambda_1 + \lambda_2$,

$$(\lambda z - \mu)(1 - z)P(z, z) = -\mu(1 - z)\pi(0, 0).$$

Cancelling $(1 - z) \neq 0$ for $z \in [0, 1)$ gives the diagonal identity

$$P(z, z) = \frac{\mu \pi(0, 0)}{\mu - \lambda z} = \frac{\pi(0, 0)}{1 - \rho z}. \quad (9.9)$$

Setting $z = 1$ and using the normalisation $P(1, 1) = 1 - \pi_0$ yields $1 - \pi_0 = \pi(0, 0)/(1 - \rho)$. The idle balance (9.5) gives $\lambda \pi_0 = \mu \pi(0, 0)$, i.e. $\pi(0, 0) = \rho \pi_0$; substituting, $(1 - \pi_0)(1 - \rho) = \rho \pi_0$, hence $\pi_0 = 1 - \rho$ and $\pi(0, 0) = \rho(1 - \rho)$. $\square$

9.3.1 Analytical approach for Model- B_2^H

Proof of Theorem 4. We apply the *Kernel Method* to the fundamental equation (9.6). Fix $y \in (0, 1]$ and seek the values of x at which the coefficient of $P(x, y)$ — the *kernel* — vanishes. Since $P(x, y)$ is a PGF it is bounded on the closed unit disk, so at any such root the right-hand side of (9.6) must vanish as well. Dividing the kernel by $y \neq 0$,

$$0 = (\lambda_1 + \lambda_2 + \mu)x - \mu - \lambda_1 x^2 - \lambda_2 xy + \gamma_1(x - y),$$

and rearranging into a quadratic in x gives exactly (9.8),

$$\lambda_1 x^2 - [\lambda_1 + \lambda_2(1 - y) + \mu + \gamma_1]x + (\mu + \gamma_1 y) = 0,$$

whose two roots are

$$x^\pm(y) = \frac{A(y) \pm \sqrt{A(y)^2 - 4\lambda_1(\mu + \gamma_1 y)}}{2\lambda_1}, \quad A(y) := \lambda_1 + \lambda_2(1 - y) + \mu + \gamma_1.$$

This is precisely the Model-A kernel quadratic (5.8) with μ replaced by $\mu + \gamma_1 y$. We label the root carrying the negative sign $x^*(y)$ and the other $x^+(y)$. By Vieta's formulas the product of the roots is $x^*(y)x^+(y) = (\mu + \gamma_1 y)/\lambda_1$.

To locate the roots relative to the unit disk we evaluate the kernel quadratic, written as $f(x) := \lambda_1 x^2 - A(y)x + (\mu + \gamma_1 y)$, at $x = 1$:

$$f(1) = \lambda_1 - [\lambda_1 + \lambda_2(1 - y) + \mu + \gamma_1] + (\mu + \gamma_1 y) = -(1 - y)(\lambda_2 + \gamma_1) < 0 \quad \text{for } y \in [0, 1).$$

Since f is an upward parabola ($\lambda_1 > 0$) with $f(1) < 0$, the point $x = 1$ lies strictly between its two roots, so $x^*(y) < 1 < x^+(y)$ for all $y \in [0, 1)$. Hence $x^*(y)$ is the unique root inside the unit disk, and at $y = 1$ one checks $f(1) = 0$ with $x^*(1) = 1$ and $x^+(1) = (\mu + \gamma_1)/\lambda_1 > 1$ (the latter exceeding 1 because $\lambda_1 < \mu < \mu + \gamma_1$ under stability). For later use, we record $\frac{d}{dy}x^*(1)$. The cleanest route is implicit differentiation of $f(x, y) := \lambda_1 x^2 - A(y)x + (\mu + \gamma_1 y) = 0$ at the point $(x, y) = (1, 1)$, where $f_y = \lambda_2 x + \gamma_1$ and $f_x = 2\lambda_1 x - A(y)$ evaluate to $f_y(1, 1) = \lambda_2 + \gamma_1$ and $f_x(1, 1) = \lambda_1 - \mu - \gamma_1$, so $\frac{d}{dy}x^*(1) = -f_y/f_x$:

$$\frac{d}{dy}x^*(1) = \frac{\lambda_2 + \gamma_1}{\mu + \gamma_1 - \lambda_1}. \quad (9.10)$$

Unlike Model-A's (5.9), the numerator carries an extra γ_1 : the constant term $\mu + \gamma_1 y$ of the kernel (9.8) is itself y -dependent here, whereas in Model-A (and in Model- C_2^H below) the constant term is independent of y .

Setting $x = x^*(y)$ in (9.6) and equating the right-hand side to zero,

$$[\mu(x^*(y) - y) + \gamma_1 y(x^*(y) - y)]P_y(y) - \mu x^*(y)(1 - y)\pi(0, 0) = 0,$$

that is, $(x^*(y) - y)(\mu + \gamma_1 y)P_y(y) = \mu x^*(y)(1 - y)\pi(0, 0)$, whence the boundary function is

$$P_y(y) = \frac{\mu x^*(y)(1 - y)\pi(0, 0)}{(\mu + \gamma_1 y)(x^*(y) - y)}. \quad (9.11)$$

As a consistency check, set $y = 0$ in (9.11): the factor $\mu + \gamma_1 y$ becomes μ and $1 - y$ becomes 1, so $P_y(0) = \mu x^*(0)\pi(0, 0)/(\mu x^*(0)) = \pi(0, 0)$ (using $x^*(0) > 0$, the smaller of the two positive roots), recovering the boundary value $P_y(0) = \pi(0, 0)$.

Finally, substitute (9.11) and $\pi(0, 0) = \rho(1 - \rho)$ (Corollary 5) back into (9.6). Writing the kernel as $-\lambda_1 y(x - x^*(y))(x - x^+(y))$ and simplifying the right-hand side,

$$\begin{aligned} [\mu(x - y) + \gamma_1 y(x - y)]P_y(y) - \mu x(1 - y)\pi(0, 0) &= \mu(1 - y)\pi(0, 0) \left[\frac{(x - y)x^*(y)}{x^*(y) - y} - x \right] \\ &= \mu(1 - y)\pi(0, 0) \frac{y(x - x^*(y))}{x^*(y) - y}, \end{aligned}$$

where the bracket collapses because $(x - y)x^* - x(x^* - y) = y(x - x^*)$. Dividing by the factored kernel, the common factor $(x - x^*(y))$ cancels and we obtain

$$P(x, y) = \frac{\mu(1 - y)\pi(0, 0)y(x - x^*(y))/(x^*(y) - y)}{-\lambda_1 y(x - x^*(y))(x - x^+(y))} = \frac{\mu\rho(1 - \rho)(1 - y)}{\lambda_1(x^*(y) - y)(x^+(y) - x)},$$

which is (9.7) and concludes the proof. $\square$

9.3.2 Probabilistic approach for Model- B_2^H

For Model- A , the boundary function $P_y(y)$ admits an independent derivation by the *Pollaczek-Khinchine* (PK) formula (§5.1.2): conditional on $N_1 = 0$, class 2 experiences an $M/G/1$ queue whose service times are i.i.d. copies of a class-1 busy period, and the PK formula then delivers $P_y(y)$. *This route does not transfer to Model- B_2^H* . Head-of-line jockeying injects class-1 customers into Queue 2 at the instants when $N_1 \geq 1$; the resulting arrival stream into Queue 2 is the superposition of the exogenous Poisson stream λ_2 and a state-dependent jockeying stream. This merged stream is *neither* Poisson *nor* independent of the system state — it is correlated with N_1 . Since the PK formula requires PASTA, which in turn demands both a Poisson arrival process *and* the lack-of-anticipation property, and the merged input violates both, class 2 does not see a renewal $M/G/1$ input and no PK representation of $P_y(y)$ is available. The analytical derivation of §9.3.1 is consequently the only route, and we retain this subsection solely to record why the probabilistic one is unavailable — in the same spirit as the discussion of Models B and B_2 .

9.4 Limits and sanity checks

As $\gamma_1 \rightarrow 0^+$ the kernel quadratic (9.8) reduces to Model- A 's quadratic (5.8), so $x^*(y) \rightarrow x_A^*(y)$, and formula (9.7) reduces to the factored form of Theorem 1; the empty-state probabilities already equal those of Model- A by Corollary 5. Setting $x = y = 1$ in (9.7) through the diagonal identity (9.9) gives $P(1, 1) = \rho = 1 - \pi_0$, confirming normalisation, and the diagonal $P(z, z) = \pi(0, 0)/(1 - \rho z)$ is exactly that of the combined $M/M/1(\lambda, \mu)$ queue — as it must be, since jockeying conserves the total customer count. Beyond this $\gamma_1 \rightarrow 0^+$ limit, B_2^H itself occupies the $k = 1$ rung of the head-of-line ladder discussed in Section 13.5, which interpolates towards the length-proportional Model- B_2 of Section 7 as $k \rightarrow \infty$.

9.5 Mean queue lengths and mean waiting times

9.5.1 Class-1 marginal PGF and mean queue length

Setting $y = 1$ in the fundamental equation (9.6), the term $\mu x(1 - y)\pi(0, 0)$ vanishes and the coefficient of $P(x, 1)$ factors as $-(x - 1)(\lambda_1 x - (\mu + \gamma_1))$, while the right-hand side becomes $(x - 1)(\mu + \gamma_1)P_y(1)$. Cancelling $(x - 1)$ for $x \neq 1$,

$$P(x, 1) = \frac{(\mu + \gamma_1)P_y(1)}{\mu + \gamma_1 - \lambda_1 x} = \frac{P_y(1)}{1 - \rho_B x}, \quad \rho_B := \frac{\lambda_1}{\mu + \gamma_1}. \quad (9.12)$$

This is a geometric class-1 marginal: the priority queue behaves as an $M/M/1$ queue whose effective *clearance rate* is $\mu + \gamma_1$, since a waiting class-1 customer at the head of Queue 1 departs that queue either by entering service (rate μ) or by jockeying (rate γ_1). The constant $P_y(1)$ is pinned by the normalisation $P(1, 1) = 1 - \pi_0 = \rho$, giving $P_y(1) = \rho(1 - \rho_B)$, so that $P(x, 1) = \rho(1 - \rho_B)/(1 - \rho_B x)$. Differentiating by the quotient rule and evaluating at $x = 1$,

$$\mathbb{E}[N_1] = \left. \frac{d}{dx} P(x, 1) \right|_{x=1} = \frac{\rho(1 - \rho_B)\rho_B}{(1 - \rho_B)^2} = \frac{\rho\rho_B}{1 - \rho_B}. \quad (9.13)$$

This is Model- A 's formula (5.22) with the priority load ρ_1 replaced by the *jockeying-reduced* load $\rho_B = \lambda_1/(\mu + \gamma_1) \leq \rho_1$.

9.5.2 Total queue length from the diagonal

Because jockeying conserves the total number of customers, the combined process $N_1 + N_2$ (together with the in-service customer) is the queue-length process of an ordinary $M/M/1(\lambda, \mu)$ queue, and the diagonal generating function (9.9) is $P(z, z) = \rho(1 - \rho)/(1 - \rho z)$. Since $\left. \frac{d}{dz} P(z, z) \right|_{z=1} = \mathbb{E}[N_1] + \mathbb{E}[N_2]$, differentiating gives

$$\mathbb{E}[N_1] + \mathbb{E}[N_2] = \left. \frac{d}{dz} \frac{\rho(1 - \rho)}{1 - \rho z} \right|_{z=1} = \frac{\rho^2(1 - \rho)}{(1 - \rho)^2} = \frac{\rho^2}{1 - \rho}, \quad (9.14)$$

identical to the combined $M/M/1$ result of Model- A (5.24).

9.5.3 Class-2 mean queue length by subtraction

Subtracting (9.13) from (9.14),

$$\begin{aligned}\mathbb{E}[N_2] &= \frac{\rho^2}{1-\rho} - \frac{\rho\rho_B}{1-\rho_B} = \frac{\rho[\rho(1-\rho_B) - \rho_B(1-\rho)]}{(1-\rho)(1-\rho_B)} \\ &= \frac{\rho(\rho - \rho_B)}{(1-\rho)(1-\rho_B)}.\end{aligned}\tag{9.15}$$

The excess load admits a transparent decomposition,

$$\rho - \rho_B = \frac{\lambda_1 + \lambda_2}{\mu} - \frac{\lambda_1}{\mu + \gamma_1} = \frac{(\lambda_1 + \lambda_2)(\mu + \gamma_1) - \lambda_1\mu}{\mu(\mu + \gamma_1)} = \frac{\lambda_2\mu + \lambda_1\gamma_1 + \lambda_2\gamma_1}{\mu(\mu + \gamma_1)} = \rho_2 + \frac{\lambda_1\gamma_1}{\mu(\mu + \gamma_1)},$$

so $\rho - \rho_B$ exceeds the bare class-2 load ρ_2 precisely by the load that jockeying transfers from Queue 1 into Queue 2. At $\gamma_1 = 0$ it collapses to ρ_2 and (9.15) recovers Model-A's $\mathbb{E}[N_2] = \rho\rho_2/[(1-\rho_1)(1-\rho)]$.

9.5.4 Mean waiting times via Little's Law

We apply **Little's Law** to each queue separately, recalling that N_i counts only the customers *waiting*, so $\mathbb{E}[N_i] = \Lambda_i \mathbb{E}[W_i]$ with Λ_i the long-run rate at which customers *enter* Queue i and $\mathbb{E}[W_i]$ their mean residence time in that queue (from entry until they leave it, by service entry or by jockeying).

Every class-1 customer enters Queue 1 exactly once, by exogenous arrival, so $\Lambda_1 = \lambda_1$ and

$$\mathbb{E}[W_1] = \frac{\mathbb{E}[N_1]}{\lambda_1} = \frac{\rho\rho_B}{\lambda_1(1-\rho_B)} = \frac{\rho}{\mu + \gamma_1 - \lambda_1}.$$

Compared with Model-A's $\mathbb{E}[W_1] = \rho/(\mu - \lambda_1)$ (5.25), the denominator is enlarged by γ_1 : head-of-line jockeying drains the priority queue faster and shortens class-1 waits.

Queue 2 receives a *mixed* input: the exogenous Poisson stream at rate λ_2 together with the jockeying stream. The latter fires at rate γ_1 whenever $N_1 \geq 1$, so its long-run rate is $\gamma_1 \mathbb{P}(N_1 \geq 1)$. The probability that a class-1 customer is waiting follows from the geometric marginal (9.12) by subtracting the $N_1 = 0$ mass from the busy mass,

$$\mathbb{P}(N_1 \geq 1) = P(1, 1) - P(0, 1) = \rho - \rho(1 - \rho_B) = \rho\rho_B,$$

since $P(1, 1) = 1 - \pi_0 = \rho$ and $P(0, 1) = P_y(1) = \rho(1 - \rho_B)$. The total entrance rate into Queue 2 is therefore $\Lambda_2 = \lambda_2 + \gamma_1\rho\rho_B$, and Little's Law gives the mean residence time per Queue-2 entrant,

$$\mathbb{E}[W_2] = \frac{\mathbb{E}[N_2]}{\Lambda_2} = \frac{\rho(\rho - \rho_B)}{(1-\rho)(1-\rho_B)(\lambda_2 + \gamma_1\rho\rho_B)}.\tag{9.16}$$

As $\gamma_1 \rightarrow 0^+$ the jockeying stream disappears, $\Lambda_2 \rightarrow \lambda_2$ and $\rho_B \rightarrow \rho_1$, and (9.16) recovers Model-A's $\mathbb{E}[W_2] = \rho/[\mu(1 - \rho_1)(1 - \rho)]$ (5.26).

10 Analysis of Model- C_2^H ($\theta_1 > 0$ head-of-line, $\gamma_1 = \gamma_2 = \theta_2 = 0$)

Model- C_2^H is the abandonment counterpart of Model- B_2^H , and a simplified variant of the class-1 abandonment model Model- C_2 . In Model- C_2 *every* waiting class-1 customer is impatient and reneges at the linear rate $\theta_1 n_1$, turning the class-1 subsystem into an $M/M/1+M$ (Erlang-A) queue and the fundamental equation into a first-order ODE. Here only the *customer at the head of Queue 1* may abandon, at the fixed rate θ_1 whenever $n_1 \geq 1$; the aggregate abandonment rate is $\theta_1 \mathbf{1}_{\{n_1 \geq 1\}}$ rather than $\theta_1 n_1$. As with Model- B_2^H , the flat rate contributes an *algebraic* term to the fundamental equation rather than a derivative, and the Kernel Method applies directly. Unlike jockeying, abandonment is a *true departure*: it removes a customer from the system and so breaks the conservation of the total customer count that Model- B_2^H enjoyed. The empty-state probabilities and the class-2 mean queue length consequently require a separate argument.

Stability. Because only the head of Queue 1 abandons, the aggregate abandonment rate is bounded by θ_1 regardless of queue length, so — in contrast to the Erlang-A class-1 subsystem of

Model- C_2 , which is positive recurrent for all loads — a genuine stability condition is required. As shown in Corollary 6 below, the chain is positive recurrent if and only if $\pi_0 > 0$, that is

$$(\mu + \theta_1)(\mu - \lambda) + \lambda_1 \theta_1 > 0, \quad \lambda = \lambda_1 + \lambda_2. \quad (10.1)$$

This is *weaker* than $\rho < 1$: if $\rho < 1$ then $\mu - \lambda > 0$ and the condition holds trivially, but it can also hold with $\rho \geq 1$ provided θ_1 is large enough — head abandonment accelerates the clearance of the priority queue and can stabilise the system at offered loads exceeding the service capacity. Moreover (10.1) implies $\rho_C := \lambda_1/(\mu + \theta_1) < 1$, so the priority subqueue is itself stable whenever the condition holds.

10.1 Balance equations

The abandonment transition $(n_1, n_2) \rightarrow (n_1 - 1, n_2)$ occurs at the fixed rate θ_1 whenever $n_1 \geq 1$. Since both a service completion and a head abandonment from $(n_1 + 1, n_2)$ land in (n_1, n_2) , the two mechanisms combine into the coefficient $\mu + \theta_1$ on the in-flows. On S ,

$$\begin{aligned} & [\lambda_1 + \lambda_2 + \mu + \theta_1] \pi(n_1, n_2) \\ &= (\mu + \theta_1) \pi(n_1 + 1, n_2) + \lambda_1 \pi(n_1 - 1, n_2) + \lambda_2 \pi(n_1, n_2 - 1) \quad n_1 \geq 1, n_2 \geq 1, \quad (10.2) \\ & [\lambda_1 + \lambda_2 + \mu + \theta_1] \pi(n_1, 0) = (\mu + \theta_1) \pi(n_1 + 1, 0) + \lambda_1 \pi(n_1 - 1, 0) \quad n_1 \geq 1, n_2 = 0, \\ & [\lambda_1 + \lambda_2 + \mu] \pi(0, n_2) \\ &= (\mu + \theta_1) \pi(1, n_2) + \lambda_2 \pi(0, n_2 - 1) + \mu \pi(0, n_2 + 1) \quad n_1 = 0, n_2 \geq 1, \\ & [\lambda_1 + \lambda_2 + \mu] \pi(0, 0) = (\lambda_1 + \lambda_2) \pi_0 + (\mu + \theta_1) \pi(1, 0) + \mu \pi(0, 1), \\ & (\lambda_1 + \lambda_2) \pi_0 = \mu \pi(0, 0). \end{aligned} \quad (10.3)$$

The contrast with Model- C_2 is the same as before: there the in-flow coefficients $\mu + \theta_1(n_1 + 1)$ and the out-flow coefficient $\theta_1 n_1$ grow with the queue length, whereas here they are the constants $\mu + \theta_1$ and θ_1 throughout.

10.2 Fundamental PGF equation

Lemma 5 (Fundamental equation for Model- C_2^H). *Assume positive recurrence. Then the joint PGF $P(x, y)$ and the boundary function $P_y(y) = P(0, y)$ satisfy*

$$\begin{aligned} & [(\lambda_1 + \lambda_2 + \mu)xy - \mu y - \lambda_1 x^2 y - \lambda_2 x y^2 + \theta_1 y(x - 1)] P(x, y) \\ &= [\mu(x - y) + \theta_1 y(x - 1)] P_y(y) - \mu x(1 - y) \pi(0, 0). \end{aligned} \quad (10.4)$$

Proof. As in Lemma 4, the non-abandonment terms of (10.2)–(10.3) reproduce the Model-A fundamental equation (5.2) after multiplication by $x^{n_1} y^{n_2}$, summation, and multiplication by xy . The abandonment *out-flow* appears in every state with $n_1 \geq 1$,

$$\theta_1 \sum_{n_1 \geq 1} \sum_{n_2 \geq 0} \pi(n_1, n_2) x^{n_1} y^{n_2} = \theta_1 [P(x, y) - P_y(y)],$$

and its *in-flow* enters every state (n_1, n_2) from $(n_1 + 1, n_2)$,

$$\theta_1 \sum_{n_1 \geq 0} \sum_{n_2 \geq 0} \pi(n_1 + 1, n_2) x^{n_1} y^{n_2} = \theta_1 x^{-1} \sum_{m_1 \geq 1} \sum_{n_2 \geq 0} \pi(m_1, n_2) x^{m_1} y^{n_2} = \theta_1 x^{-1} [P(x, y) - P_y(y)],$$

with $m_1 = n_1 + 1$. The net contribution is $\theta_1(1 - x^{-1})[P(x, y) - P_y(y)]$; multiplying by xy gives $\theta_1 y(x - 1)[P(x, y) - P_y(y)]$. Adding to the Model-A equation and rearranging yields (10.4). $\square$

Remark. Again no derivative of P appears: the algebraic factor $\theta_1 y(x - 1)$ replaces the convective term $\theta_1 x y(x - 1) \partial P / \partial x$ of Model- C_2 . Setting $\theta_1 = 0$ recovers (5.2). Unlike the jockeying factor of Model- B_2^H , the abandonment factor $\theta_1 y(x - 1)$ does *not* vanish on the diagonal $x = y$, which is the analytic signature of the broken conservation. This algebraic collapse—no derivative term, and hence no singularity to resolve—is the head-of-line regime anticipated in the remark following Lemma 1.

10.3 Closed-form solution

Theorem 5 (Closed-form PGF for Model- C_2^H). *Consider the two-class non-preemptive priority queue with Poisson arrival rates λ_1, λ_2 , common exponential service rate μ , and head-of-line class-1 abandonment at rate $\theta_1 > 0$ (with $\gamma_1 = \gamma_2 = \theta_2 = 0$). Under the stability condition (10.1), the joint PGF $P(x, y) = \mathbb{E}[x^{N_1} y^{N_2}]$ is the rational function*

$$P(x, y) = \frac{\mu(\mu + \theta_1)(1 - y)\pi(0, 0)}{\lambda_1(x^+(y) - x)[(\mu + \theta_1 y)x^*(y) - y(\mu + \theta_1)]}, \quad (10.5)$$

where $x^*(y)$ and $x^+(y)$ are the two roots of the kernel quadratic

$$\lambda_1 x^2 - [\lambda_1 + \lambda_2(1 - y) + \mu + \theta_1]x + (\mu + \theta_1) = 0, \quad (10.6)$$

with $x^*(y)$ the root inside the unit disk, and $\pi(0, 0)$ is given by Corollary 6. As $\theta_1 \rightarrow 0^+$ the formula recovers Model-A (Theorem 1) exactly.

Corollary 6. *Under (10.1), the empty-state probabilities of Model- C_2^H are*

$$\pi_0 = \frac{(\mu + \theta_1)(\mu - \lambda) + \lambda_1 \theta_1}{\mu(\mu + \theta_1) + \lambda_1 \theta_1}, \quad \pi(0, 0) = \frac{\lambda[(\mu + \theta_1)(\mu - \lambda) + \lambda_1 \theta_1]}{\mu[\mu(\mu + \theta_1) + \lambda_1 \theta_1]}, \quad (10.7)$$

with $\lambda = \lambda_1 + \lambda_2$. At $\theta_1 \rightarrow 0^+$ these reduce to $\pi_0 \rightarrow 1 - \rho$ and $\pi(0, 0) \rightarrow \rho(1 - \rho)$, recovering Model-A.

Proof of Corollary 6. Two evaluations of the fundamental equation (10.4) suffice. First set $y = 1$: the term $\mu x(1 - y)\pi(0, 0)$ vanishes, the coefficient of $P(x, 1)$ factors as $-(x - 1)(\lambda_1 x - (\mu + \theta_1))$, and the right-hand side becomes $(x - 1)(\mu + \theta_1)P_y(1)$. Cancelling $(x - 1)$,

$$P(x, 1) = \frac{(\mu + \theta_1)P_y(1)}{\mu + \theta_1 - \lambda_1 x} = \frac{P_y(1)}{1 - \rho_C x}, \quad \rho_C := \frac{\lambda_1}{\mu + \theta_1}. \quad (10.8)$$

This geometric marginal shows that the priority queue clears at the effective rate $\mu + \theta_1$ (service completions plus head abandonments). Evaluating at $x = 1$ and using $P(1, 1) = 1 - \pi_0$ gives

$$1 - \pi_0 = \frac{P_y(1)}{1 - \rho_C} \implies P_y(1) = (1 - \pi_0)(1 - \rho_C). \quad (10.9)$$

Second, set $x = 1$ in (10.4). The coefficient of $P(1, y)$ collapses to $\lambda_2 y(1 - y)$ and the right-hand side to $\mu(1 - y)[P_y(y) - \pi(0, 0)]$, yielding the class-2 marginal relation

$$\lambda_2 y P(1, y) = \mu[P_y(y) - \pi(0, 0)]. \quad (10.10)$$

Letting $y \rightarrow 1$ and using (10.9) and the idle balance (10.3), $\pi(0, 0) = \lambda\pi_0/\mu$,

$$\lambda_2(1 - \pi_0) = \mu[P_y(1) - \pi(0, 0)] = \mu(1 - \pi_0)(1 - \rho_C) - \lambda\pi_0.$$

This is a linear equation in π_0 . Expanding and collecting the π_0 terms on the left,

$$[\mu(1 - \rho_C) + \lambda - \lambda_2]\pi_0 = \mu(1 - \rho_C) - \lambda_2, \quad \text{i.e.} \quad \pi_0 = \frac{\mu(1 - \rho_C) - \lambda_2}{\mu(1 - \rho_C) + \lambda_1},$$

where we used $\lambda - \lambda_2 = \lambda_1$. Substituting $\mu(1 - \rho_C) = \mu(\mu + \theta_1 - \lambda_1)/(\mu + \theta_1)$ and multiplying numerator and denominator by $\mu + \theta_1$, the numerator becomes $\mu(\mu + \theta_1 - \lambda_1) - \lambda_2(\mu + \theta_1) = (\mu + \theta_1)(\mu - \lambda) + \lambda_1 \theta_1$ and the denominator $\mu(\mu + \theta_1 - \lambda_1) + \lambda_1(\mu + \theta_1) = \mu(\mu + \theta_1) + \lambda_1 \theta_1$, yielding the first identity in (10.7); the second follows from $\pi(0, 0) = \lambda\pi_0/\mu$. The denominator $\mu(\mu + \theta_1) + \lambda_1 \theta_1$ is strictly positive, so $\pi_0 > 0$ is equivalent to the stability condition (10.1). $\square$

10.3.1 Analytical approach for Model- C_2^H

Proof of Theorem 5. We again apply the Kernel Method to (10.4). Dividing the coefficient of $P(x, y)$ by $y \neq 0$ and rearranging gives the kernel quadratic (10.6),

$$\lambda_1 x^2 - [\lambda_1 + \lambda_2(1 - y) + \mu + \theta_1]x + (\mu + \theta_1) = 0,$$

which is the Model- A kernel quadratic with μ replaced by $\mu + \theta_1$ *both* in the linear coefficient and in the constant term. Its roots are

$$x^\pm(y) = \frac{A(y) \pm \sqrt{A(y)^2 - 4\lambda_1(\mu + \theta_1)}}{2\lambda_1}, \quad A(y) := \lambda_1 + \lambda_2(1 - y) + \mu + \theta_1, \quad (10.11)$$

with product $x^*(y)x^+(y) = (\mu + \theta_1)/\lambda_1$ by Vieta's formulas. Evaluating the quadratic $f(x) := \lambda_1 x^2 - A(y)x + (\mu + \theta_1)$ at $x = 1$,

$$f(1) = \lambda_1 - A(y) + (\mu + \theta_1) = -\lambda_2(1 - y) < 0 \quad \text{for } y \in [0, 1],$$

so $x = 1$ lies between the two roots and $x^*(y) < 1 < x^+(y)$ for $y \in [0, 1]$; at $y = 1$, $x^*(1) = 1$ and $x^+(1) = (\mu + \theta_1)/\lambda_1 > 1$ (the latter by $\rho_C < 1$). Hence $x^*(y)$ is the unique root inside the unit disk. Differentiating (10.11) at $y = 1$,

$$\frac{d}{dy}x^*(1) = \frac{\lambda_2}{\mu + \theta_1 - \lambda_1}, \quad (10.12)$$

the analogue of (5.9) and (9.10).

Setting $x = x^*(y)$ in (10.4) forces the right-hand side to vanish,

$$[\mu(x^*(y) - y) + \theta_1 y(x^*(y) - 1)]P_y(y) = \mu x^*(y)(1 - y)\pi(0, 0),$$

and, regrouping the bracket as $\mu(x^* - y) + \theta_1 y(x^* - 1) = (\mu + \theta_1 y)x^* - y(\mu + \theta_1)$, the boundary function is

$$P_y(y) = \frac{\mu x^*(y)(1 - y)\pi(0, 0)}{(\mu + \theta_1 y)x^*(y) - y(\mu + \theta_1)}. \quad (10.13)$$

At $y = 0$ the denominator is $\Delta(0) = \mu x^*(0)$, so $P_y(0) = \mu x^*(0)\pi(0, 0)/(\mu x^*(0)) = \pi(0, 0)$ (using $x^*(0) > 0$, the smaller positive root), as required.

Substituting (10.13) into (10.4), writing the kernel as $-\lambda_1 y(x - x^*(y))(x - x^+(y))$, and abbreviating the denominator of (10.13) by $\Delta(y) := (\mu + \theta_1 y)x^*(y) - y(\mu + \theta_1)$, the right-hand side becomes

$$\begin{aligned} & [\mu(x - y) + \theta_1 y(x - 1)]P_y(y) \\ & - \mu x(1 - y)\pi(0, 0) = \frac{\mu(1 - y)\pi(0, 0)}{\Delta(y)} \left\{ [\mu(x - y) + \theta_1 y(x - 1)]x^*(y) - x\Delta(y) \right\} \\ & = \frac{\mu(1 - y)\pi(0, 0)}{\Delta(y)} (\mu + \theta_1)y(x - x^*(y)), \end{aligned}$$

where, after substituting $\Delta(y) = (\mu + \theta_1 y)x^*(y) - y(\mu + \theta_1)$, the curly bracket is *linear* in $x^*(y)$ and collapses term-by-term to $(\mu + \theta_1)y(x - x^*(y))$ (no use of the kernel relation is required). Dividing by the factored kernel, the factor $(x - x^*(y))$ cancels and

$$P(x, y) = \frac{\mu(\mu + \theta_1)(1 - y)\pi(0, 0)y(x - x^*(y))/\Delta(y)}{-\lambda_1 y(x - x^*(y))(x - x^+(y))} = \frac{\mu(\mu + \theta_1)(1 - y)\pi(0, 0)}{\lambda_1(x^+(y) - x)\Delta(y)},$$

which is (10.5) and concludes the proof. $\square$

10.3.2 Probabilistic approach for Model- C_2^H

As for Model- B_2^H , the Pollaczek-Khinchine route used for Models A and C_2 does not transfer. In Model- C_2 the class-1 subsystem seen by class 2 is a clean $M/M/1+M$ (Erlang-A) queue whose busy period B_C is an i.i.d. renewal quantity, and class 2 sees an $M/G/1$ queue with service B_C . Under head-only abandonment this clean renewal structure is lost: abandonment acts at the fixed rate θ_1 only while a class-1 customer waits ($N_1 \geq 1$) but is inactive when the single class-1 present is in service, so the departure rate of the class-1 subsystem is state-dependent at its boundary and the class-1 busy period is not the busy period of any single $M/M/1$ or $M/M/1+M$ queue. Consequently class 2 does not see a renewal $M/G/1$ input and no Pollaczek-Khinchine representation of $P_y(y)$ is available. We retain this subsection, in parallel with the other models, only to record why the analytical route of §10.3.1 is the operative one.

10.4 Limits and sanity checks

As $\theta_1 \rightarrow 0^+$ the kernel quadratic (10.6) reduces to Model- A 's, so $x^*(y) \rightarrow x_A^*(y)$ and $\Delta(y) \rightarrow \mu(x_A^*(y) - y)$; together with $\pi(0, 0) \rightarrow \rho(1 - \rho)$ from Corollary 6, formula (10.5) reduces to the factored Model- A PGF of Theorem 1. Setting $x = 1$ in (10.5) and $y \rightarrow 1$ reproduces $P(1, 1) = 1 - \rho_0$ via (10.9), confirming normalisation.

10.5 Mean queue lengths and mean waiting times

10.5.1 Class-1 marginal PGF and mean queue length

The class-1 marginal is the geometric law (10.8); pinning the constant by $P(1,1) = 1 - \pi_0$ as in (10.9) gives $P(x,1) = (1 - \pi_0)(1 - \rho_C)/(1 - \rho_C x)$. Differentiating and evaluating at $x = 1$,

$$\mathbb{E}[N_1] = \left. \frac{d}{dx} P(x,1) \right|_{x=1} = \frac{(1 - \pi_0) \rho_C}{1 - \rho_C}, \quad \rho_C = \frac{\lambda_1}{\mu + \theta_1}, \quad (10.14)$$

with $1 - \pi_0 = \lambda(\mu + \theta_1)/[\mu(\mu + \theta_1) + \lambda_1 \theta_1]$ from Corollary 6. As in Model- B_2^H , the priority queue behaves as an $M/M/1$ queue with effective clearance rate $\mu + \theta_1$, the only difference from Model- B_2^H 's (9.13) being the prefactor $1 - \pi_0$, which here is no longer equal to ρ because abandonment alters the idle probability.

10.5.2 Class-2 mean queue length

Abandonment is a true departure and so breaks the *conservation of the total customer count* that underlies the diagonal argument: in Model- B_2^H the jockeying factor $\gamma_1 y(x - y)$ vanishes on the diagonal $x = y$, whereas here the abandonment factor $\theta_1 y(x - 1)$ does not. Concretely, setting $x = y = z$ in (10.4) leaves the term $\theta_1 z(z - 1)P_y(z)$ on the right-hand side, so the diagonal equation no longer closes and the route that furnished $\mathbb{E}[N_1] + \mathbb{E}[N_2]$ for Model- B_2^H is unavailable. We instead differentiate the class-2 marginal. From the marginal relation (10.10), $P(1,y) = \mu[P_y(y) - \pi(0,0)]/(\lambda_2 y)$, and since $\mathbb{E}[N_2] = \left. \frac{d}{dy} P(1,y) \right|_{y=1}$,

$$\mathbb{E}[N_2] = \frac{\mu}{\lambda_2} P'_y(1) - (1 - \pi_0), \quad (10.15)$$

where we used $\left. \frac{\mu}{\lambda_2} [P_y(1) - \pi(0,0)] \right|_{y=1} = P(1,1) = 1 - \pi_0$. It remains to compute $P'_y(1)$ from (10.13). Both the numerator $\mu x^*(y)(1 - y)\pi(0,0)$ and the denominator $\Delta(y) = (\mu + \theta_1 y)x^*(y) - y(\mu + \theta_1)$ vanish at $y = 1$ (since $x^*(1) = 1$ gives $\Delta(1) = (\mu + \theta_1) - (\mu + \theta_1) = 0$), so $P_y(y)$ is a 0/0 form there. Writing $N(y)$ and $\Delta(y)$ for numerator and denominator, a second-order Taylor expansion about $y = 1$ gives the value and the slope

$$P_y(1) = \frac{N'(1)}{\Delta'(1)} = (1 - \pi_0)(1 - \rho_C), \quad P'_y(1) = \frac{N''(1)\Delta'(1) - N'(1)\Delta''(1)}{2\Delta'(1)^2},$$

which is why both the first and the second root sensitivities of x^* at $y = 1$ are needed, namely (10.12) together with

$$\left. \frac{d^2}{dy^2} x^*(y) \right|_{y=1} = \frac{2(\mu + \theta_1) \lambda_2^2}{(\mu + \theta_1 - \lambda_1)^3}$$

(the latter obtained by differentiating $x^*(y)$ in (10.11) twice). Substituting these into $P'_y(1)$, then into (10.15), and simplifying yields the closed form

$$\mathbb{E}[N_2] = \frac{\lambda_2 \lambda (\mu + \theta_1) [(\mu + \theta_1)^2 - \lambda_1 \theta_1]}{(\mu + \theta_1 - \lambda_1) [\mu(\mu + \theta_1) + \lambda_1 \theta_1] [(\mu + \theta_1)(\mu - \lambda) + \lambda_1 \theta_1]}, \quad (10.16)$$

with $\lambda = \lambda_1 + \lambda_2$. The last bracket in the denominator is exactly the stability numerator of π_0 from (10.7), so $\mathbb{E}[N_2]$ is finite precisely under the stability condition (10.1) and diverges as $\pi_0 \rightarrow 0^+$. As $\theta_1 \rightarrow 0^+$ the right-hand side reduces to $\lambda_2 \lambda / [(\mu - \lambda_1)(\mu - \lambda)] = \rho \rho_2 / [(1 - \rho_1)(1 - \rho)]$, recovering Model- A 's (5.24).

10.5.3 Mean waiting times via Little's Law

We apply **Little's Law** to each queue. Every class-1 customer enters Queue 1 once by exogenous arrival and later leaves it either by entering service or by abandoning from the head; in both cases it leaves Queue 1, so the entrance rate is $\Lambda_1 = \lambda_1$ and

$$\mathbb{E}[W_1] = \frac{\mathbb{E}[N_1]}{\lambda_1} = \frac{(1 - \pi_0) \rho_C}{\lambda_1(1 - \rho_C)}.$$

Here $\mathbb{E}[W_1]$ is the mean time a class-1 customer spends *waiting* in Queue 1, averaged over all class-1 arrivals, whether they are eventually served or abandon. Queue 2, by contrast, receives no

jockeying input and its customers never abandon ($\theta_2 = 0$), so every class-2 arrival is eventually served and the entrance rate is simply $\Lambda_2 = \lambda_2$:

$$\mathbb{E}[W_2] = \frac{\mathbb{E}[N_2]}{\lambda_2} = \frac{\lambda(\mu + \theta_1)[(\mu + \theta_1)^2 - \lambda_1\theta_1]}{(\mu + \theta_1 - \lambda_1)[\mu(\mu + \theta_1) + \lambda_1\theta_1][(\mu + \theta_1)(\mu - \lambda) + \lambda_1\theta_1]}.$$

Both reduce to the corresponding Model-*A* expressions as $\theta_1 \rightarrow 0^+$.

10.6 Comparison of the head-of-line models

Table 4 summarises the structural differences between the two head-of-line variants and their full-rate counterparts. The common thread is that replacing a length-proportional mechanism rate ($\gamma_1 n_1$ or $\theta_1 n_1$) by a head-of-line rate ($\gamma_1 \mathbf{1}_{\{n_1 \geq 1\}}$ or $\theta_1 \mathbf{1}_{\{n_1 \geq 1\}}$) converts the governing ODE into an algebraic equation, so that $P(x, y)$ becomes rational. For Model- B_2^H the result is structurally identical to Model-*A*, the sole change being $\mu \rightarrow \mu + \gamma_1 y$ in the kernel; for Model- C_2^H the broken conservation introduces the extra factor $(\mu + \theta_1)$ and the modified boundary denominator $\Delta(y)$, and shifts the stability boundary to (10.1).

Table 4: Structural comparison across the priority-queue models.

Property	A	B ₂	B ₂ ^H	C ₂	C ₂ ^H
Mechanism rate	—	$\gamma_1 n_1$	$\gamma_1 \mathbf{1}_{\{n_1 \geq 1\}}$	$\theta_1 n_1$	$\theta_1 \mathbf{1}_{\{n_1 \geq 1\}}$
PGF equation type	algebraic	ODE	algebraic	ODE	algebraic
$P(x, y)$ form	rational	integral+ $_1F_1$	rational	integral+ $_1F_1$	rational
Same form as Model- <i>A</i> ?	✓	×	✓	×	(similar)
$\pi(0, 0)$	$\rho(1 - \rho)$	$\rho(1 - \rho)$	$\rho(1 - \rho)$	(8.1)	(10.7)
π_0	$1 - \rho$	$1 - \rho$	$1 - \rho$	(8.2)	(10.7)
Conserves total N ?	✓	✓	✓	×	×
Stability condition	$\rho < 1$	$\rho < 1$	$\rho < 1$	$\lambda_2 \mathbb{E}[B_C] < 1$	(10.1)

11 Comparison of the Solved Models

This section compares the five solved models: the baseline (*A*), the jockeying extension (B_2), the abandonment extension (C_2), and their head-of-line variants (B_2^H , C_2^H). The first three are the tractable specialisations of Model-*X*, each obtained by zeroing all but one of the extra mechanisms; the head-of-line variants are the standalone simplifications of §§9–10. Each produces a closed-form or integral-form expression for $P(x, y)$. Table 5 collects the key structural facts for the three full-rate models (the head-of-line variants are compared in Table 4 and rejoin the numerical comparison below); the remainder of this section traces the consequences.

Table 5: Structural comparison of Models A, B₂, and C₂.

	Model-A	Model-B ₂	Model-C ₂
Extra mechanism	none	jockeying 1 → 2 ($\gamma_1 > 0$)	class-1 abandon- ment ($\theta_1 > 0$)
Stability condition	$\rho < 1$	$\rho < 1$	$\rho_2 \cdot {}_1F_1 < 1$
Fundamental equation	algebraic PGF equation	1st-order PDE in x, y	1st-order ODE in x
Pinning of $P_y(y)$	zero of kernel: $x = x^*(y)$	analyticity at $x = y$	boundedness at $x = 1$
Probabilistic route	yes (P-K with B)	no (jockeying breaks Poisson)	yes (P-K with B_C)
$P_y(y)$	rational in $x^*(y)$	Kummer ratio (7.20)	P-K / Kummer (8.8)
$P(x, y)$	closed form (5.2)	integral form (7.3)	integral form (8.16)
$\pi(0, 0)$	$\rho(1 - \rho)$	$\rho(1 - \rho)$	$(\rho_1 + \rho_2)\pi_0$ via $\mathbb{E}[B_C]$
π_0	$1 - \rho$	$1 - \rho$	$> 1 - \rho$ (abandonment increases idle time)
$P(z, z)$	$\rho(1 - \rho)/(1 - \rho z)$	$\rho(1 - \rho)/(1 - \rho z)$	no closed form
$\mathbb{E}[N_1] + \mathbb{E}[N_2]$	$\rho^2/(1 - \rho)$	$\rho^2/(1 - \rho)$	$< \rho^2/(1 - \rho)$
$\mathbb{E}[N_1]$ closed form	$\rho\rho_1/(1 - \rho_1)$	no	no

Model-A: the rational baseline

Model-A imposes no extra departure mechanism. Class-1 operates as an independent $M/M/1(\lambda_1, \mu)$ queue; class-2 sees an $M/G/1$ with service time equal to a class-1 busy period (by PASTA and the priority decomposition of §3.2). The Pollaczek–Khinchine formula (3.8) delivers $P_y(y)$ as a rational function of $x^*(y)$, and $P(x, y)$ is a ratio of polynomials in x , y , and $x^*(y)$. All performance measures — π_0 , $\pi(0, 0)$, $\mathbb{E}[N_i]$, $\mathbb{E}[W_i]$ — are in closed form. Model-A also furnishes the structural identities that persist across the hierarchy: $P(z, z) = \rho(1 - \rho)/(1 - \rho z)$ (total queue is $M/M/1$), $\pi_0 = 1 - \rho$, and $\pi(0, 0) = \rho(1 - \rho)$.

Effect of jockeying: Model-A → Model-B₂

Adding one-way jockeying ($\gamma_1 > 0$, $\gamma_2 = 0$) couples the two queues: a class-1 customer waiting in Queue 1 may defect to Queue 2 at rate γ_1 . The consequences are asymmetric.

What is preserved. Jockeying transfers customers but does not remove them, so the total-customer process remains a standard $M/M/1(\lambda, \mu)$ birth–death chain. Consequently $P(z, z)$, π_0 , $\pi(0, 0)$, and $\mathbb{E}[N]$ are identical to Model-A for all $\gamma_1 > 0$.

What is destroyed. Jockeying breaks the per-class Poisson/renewal structure: a class-2 customer can no longer identify its effective service time with a class-1 busy period, because class-1 customers now arrive at Queue 2 from inside the system as well as from outside. The P-K route is therefore unavailable.

What is new. The fundamental PGF equation gains the convection term $\gamma_1 xy(x - y)\partial_x P$, which reduces it to a first-order ODE in x for each fixed y (since $\gamma_2 = 0$ kills the ∂_y term). The solution closes via an integrating factor with a branch singularity at the *interior* point $x = y$; the branch coefficient must vanish for P to be holomorphic, and this analyticity condition pins $P_y(y)$ as a Kummer ratio. The individual means $\mathbb{E}[N_1]$ and $\mathbb{E}[N_2]$ change with γ_1 (class-1 becomes shorter, class-2 longer), with no closed form for the split.

Effect of abandonment: Model-A → Model-C₂

Adding class-1 abandonment ($\theta_1 > 0$, $\theta_2 = 0$, $\gamma_i = 0$) allows each waiting class-1 customer to leave the system at rate θ_1 .

What is preserved. Class-2 customers still see Poisson arrivals (abandonment only affects class-1; the class-2 arrival process is undisturbed). PASTA therefore still applies, and the $M/G/1$

subsystem for class-2 is intact — now with effective service time B_C (the busy period of the $M/M/1 + M$ queue of §3.3) replacing B . The P-K formula (3.8) applies with B_C and delivers $P_y(y)$.

What is destroyed. Abandonment is a true departure: it reduces the total number of customers and breaks the conservation property. Consequently $P(z, z)$ no longer has the simple $M/M/1$ form (5.23), $\pi_0 > 1 - \rho$, $\pi(0, 0)$ requires $\mathbb{E}[B_C]$, and $\mathbb{E}[N] < \rho^2/(1 - \rho)$. The stability condition is also generalised: $\rho < 1$ is no longer necessary; the relevant condition is $\lambda_2 \mathbb{E}[B_C] < 1$, which is strictly weaker for all $\theta_1 > 0$.

What is new. The fundamental equation gains the convection term $\theta_1 xy(x - 1)\partial_x P$, which already reduces it to a first-order ODE in x for each fixed y (with $\theta_2 = 0$ killing the ∂_y term). The integrating factor now has its singularity at the *boundary* $x = 1$ with a positive exponent, so μ_I vanishes there and mere boundedness of P forces $\int_0^1 q \mu_I dt = 0$. This determines $P_y(y)$ without any reference to holomorphy. Both $\mathbb{E}[N_1]$ and $\mathbb{E}[N_2]$ decrease relative to Model-A as θ_1 increases: class-2 benefits indirectly because class-1 abandonments free the server sooner.

Structural comparison of Models B_2 and C_2

Both models reduce the fundamental equation to a first-order linear ODE in x and both close in confluent hypergeometric functions; yet the mechanism that pins $P_y(y)$ is structurally different. The contrast is as follows.

	Model- C_2 ($\theta_1 > 0$)	Model- B_2 ($\gamma_1 > 0$)
Convection coefficient	$\theta_1 x(x - 1)$	$\gamma_1 x(x - y)$
Operative singular point	boundary $x = 1$	interior $x = y$
Exponent at singularity	$\beta(y) = \lambda_2(1 - y)/\theta_1 > 0$ ($\mu_I \rightarrow 0$)	$\beta(y) = (1 - y)(\lambda y - \mu)/(\gamma_1 y) < 0$ ($\mu_I \rightarrow \infty$)
Pinning principle	<i>boundedness</i> : $\int_0^1 q \mu_I dt = 0$	<i>analyticity</i> : branch coefficient = 0
$\pi(0, 0)$	$M/G/1$ +renewal (conservation broken)	$\rho(1 - \rho)$ exactly (conservation holds)
Probabilistic backbone	class-2 is $M/G/1$; P-K with B_C	class-1 is $M/M/1+M$ (reneging = γ_1); no P-K

In words: in Model- C_2 the abandonment convection $\theta_1 x(x - 1)$ places the singularity at the boundary $x = 1$ with a positive integrating-factor exponent, so μ_I vanishes there and mere boundedness forces the vanishing integral that determines $P_y(y)$. The Poisson character of the class-2 stream survives abandonment, so P-K delivers $P_y(y)$ probabilistically and an $M/G/1$ renewal argument delivers $\pi(0, 0)$. In Model- B_2 , the jockeying convection $\gamma_1 x(x - y)$ places the singularity at the interior point $x = y$ with a negative exponent, so μ_I blows up, the homogeneous mode vanishes, and only the stronger requirement of holomorphy — the absence of the non-integer branch $(x - y)^{|\beta|}$ — pins $P_y(y)$. Jockeying destroys the per-class Poisson/renewal structure, so no P-K identity is available; what remains probabilistic is the total-customer $M/M/1$ behaviour and the $M/M/1+M$ reading of the class-1 marginal, which is the origin of the ${}_1F_1$ factors.

Common Erlang-A structure

A common object appears in both B_2 and C_2 :

$${}_1F_1\left(1; \frac{\mu}{r} + 1; \frac{\lambda_1}{r}\right), \quad (11.1)$$

with $r = \gamma_1$ (Model- B_2) or $r = \theta_1$ (Model- C_2). This is the Erlang-A utilisation factor: in C_2 it equals $\mu \mathbb{E}[B_C]$ (the mean class-1 busy period scaled by μ), while in B_2 it controls the Kummer ratio in $P_y(y)$ through the exponent $\alpha(y) = \mu/(\gamma_1 y)$. In both cases the function evaluates at a negative argument $(-\lambda_1/r)$ in C_2 and $(-\lambda_1 y/\gamma_1)$ in B_2 , and the ratio ${}_1F_1(\alpha + 1; \cdot; \cdot)/{}_1F_1(\alpha; \cdot; \cdot)$ acts as a boundary correction to the Model-A rational formula.

Model- B_2 degenerates to Model-A only in the singular limit $\gamma_1 \rightarrow 0^+$, in which $\alpha, \beta \rightarrow \infty$ and the Kummer ratio collapses to the rational expression $x^*(y)(1 - y)/(x^*(y) - y)$ of Model-A; this limit is not uniform, consistent with γ_1 multiplying the highest-order term of the ODE. Model- C_2

degenerates to Model-A uniformly as $\theta_1 \rightarrow 0^+$: $\mathbb{E}[B_C] \rightarrow (\mu - \lambda_1)^{-1}$, and all C_2 formulas recover their Model-A counterparts.

11.1 Traffic-intensity sweep

The structural comparison above is now made quantitative. Holding the service rate at $\mu = 1$ and the arrival split fixed at $\lambda_1:\lambda_2 = 3:4$, the offered load $\rho = \lambda_1 + \lambda_2$ is swept across $[0.10, 0.95]$ and the stationary performance of the three models is solved exactly on the truncated CTMC. Figure 8 collects the mean queue length $\mathbb{E}[N]$, the priority-class queue $\mathbb{E}[N_1]$, the priority waiting time $\mathbb{E}[W_1] = \mathbb{E}[N_1]/\lambda_1$ (Little's Law), and the idle probability π_0 ; Figure 9 isolates $\pi(0, 0)$ against the Model-A parabola $\rho(1 - \rho)$.




Figure 8: Performance versus traffic intensity ρ for the five solved models — the baseline A , the jockeying pair B_2 , B_2^H ($\gamma_1 = 0.5$), and the abandonment pair C_2 , C_2^H ($\theta_1 = 0.5$) — at $\mu = 1$ and $\lambda_1:\lambda_2 = 3:4$; the vertical line marks the canonical $\rho = 0.70$. *Top-left*: total queue $\mathbb{E}[N]$ — A , B_2 and B_2^H coincide because jockeying conserves N (Corollaries 1, 5), while C_2 and C_2^H lie strictly below because abandonment removes customers (C_2^H by less). *Top-right*: priority queue $\mathbb{E}[N_1]$ — the head-of-line variants drain it less than their full-rate cousins. *Bottom-left*: priority waiting time $\mathbb{E}[W_1]$. *Bottom-right*: idle probability π_0 , with the $1 - \rho$ law (Models A , B_2 , B_2^H) shown dotted; C_2 and C_2^H sit above it for every ρ , as predicted by Corollaries 4, 6.

Two mechanisms produce two qualitatively different responses. Jockeying is internal: the $1 \rightarrow 2$ transition relabels a waiting customer without changing the total count, so the B_2 and A curves for $\mathbb{E}[N]$ and π_0 are indistinguishable across the entire sweep — the server sees the same workload. Abandonment is a true departure at rate $\theta_1 n_1$, so C_2 carries a strictly lighter queue and a strictly larger idle fraction: the priority backlog is bled off before it can be served. At $\rho = 0.70$ (Table 6) this gives $\pi_0 = 0.3696$ for C_2 against 0.3000 for A and B_2 , and $\mathbb{E}[N] = 0.8358$ against 1.6333.

The divergence widens with load. At $\rho = 0.90$ the baseline class-2 queue reaches $\mathbb{E}[N_2] = 7.5346$, whereas C_2 holds it to 1.9245 — a factor of nearly four — because the self-regulating valve $\theta_1 n_1$ acts most strongly precisely when the queue is long. The priority queue $\mathbb{E}[N_1]$ is bounded for C_2 for the same reason and never exceeds 0.2292 over the sweep, while the A value grows monotonically toward the $\rho \rightarrow 1$ singularity. We note for completeness that near criticality the dense CTMC truncation leaves a boundary tail mass of order 6×10^{-4} for A and B_2 at $\rho = 0.95$ (against 1.5×10^{-7} for C_2 , whose tail is light); the $\rho \geq 0.92$ points for the non-abandoning models therefore carry a small downward truncation bias and should be read as lower bounds.




Figure 9: $\pi(0,0)$ versus ρ . The three customer-conserving models A , B_2 and B_2^H lie exactly on the parabola $\rho(1 - \rho)$ (Corollaries 1, 3, 5: jockeying preserves N , hence the busy-but-empty probability), whereas the abandonment models C_2 and C_2^H rise strictly above it — abandonment inflates $\pi(0,0)$ through the enlarged idle probability $\pi(0,0) = \rho \pi_0$ (Corollaries 4, 6), the full rate (C_2) lifting it more than the head-of-line rate (C_2^H).

The invariance of $\pi(0,0) = \rho(1 - \rho)$ along the A and B_2 curves of Figure 9 is not a numerical coincidence but a direct reading of Corollary 3: the event “server busy, both queues empty” depends only on the total occupancy N , and jockeying leaves the law of N unchanged. For C_2 the identity $\pi(0,0) = \rho \pi_0$ still holds, but $\pi_0 > 1 - \rho$, so the curve is lifted — at $\rho = 0.70$, $\pi(0,0) = 0.2587$ against the baseline 0.2100 (Table 6). The lift is also asymmetric: whereas the conserving parabola $\rho(1 - \rho)$ is symmetric about $\rho = \frac{1}{2}$, the C_2 curve in Figure 9 attains its maximum to the *right* of that point, the signature of abandonment renormalising the effective load — the busy-but-empty probability keeps climbing past the symmetric-load peak because the self-regulating valve blunts congestion fastest exactly where the baseline is heaviest, before the shrinking busy fraction finally pulls it down.

The two head-of-line variants B_2^H and C_2^H added to Figures 8–9 confirm that the qualitative split survives the head-of-line simplification but in a muted form. Model- B_2^H conserves N exactly like B_2 (Corollary 5), so its $\mathbb{E}[N]$ and π_0 curves are indistinguishable from the A/B_2 pair and lie on the same $\rho(1 - \rho)$ parabola in Figure 9; it departs only in the split of $\mathbb{E}[N_1]$ (at $\rho = 0.70$, $\mathbb{E}[N_1] = 0.1750$ against 0.1545 for B_2 , Table 6). Model- C_2^H tracks C_2 but bleeds the priority queue more gently — at $\rho = 0.70$ it gives $\mathbb{E}[N] = 0.9015$ and $\pi_0 = 0.3636$, both lying between the baseline and the full-rate C_2 values (0.8358 and 0.3696). The structural reason for this systematic “weaker-mechanism” behaviour is isolated in §11.5.

Table 6 reads off the full set of stationary descriptors at three loads; the class-1 loss accounting — throughput, abandonment flow, and the loss fraction $P_{\text{loss}}^{(1)}$ — is carried in the compact companion Table 7 so the main table stays within the margins.

Two columns deserve a reviewer’s caution. First, the priority waiting time $\mathbb{E}[W_1]$ is smallest for C_2 at every load (Table 8: 0.5941 at $\rho = 0.90$ against 0.6808 for B_2 and 1.4651 for A), but $\mathbb{E}[W_i] = \mathbb{E}[N_i]/\lambda_i$ is the mean time a customer is *counted in queue* i , and that residence is shortened by every departure mode — service, *and* attrition. In C_2 a class-1 customer leaves the count by being served *or* by abandoning; in B_2 by being served *or* by jockeying into class 2. A smaller $\mathbb{E}[W_1]$ therefore does not by itself certify faster service, and cross-model $\mathbb{E}[W_1]$ comparisons conflate latency reduction with customer loss.

The magnitude of that confound is exactly the class-1 loss fraction $L_1 = \theta_1 \mathbb{E}[N_1]/\lambda_1$ derived in §11.4 (eq. (11.3)): in C_2 at $\rho = 0.70$ a fraction $L_1 = 0.232$ of class-1 arrivals leaves by abandonment rather than by service, so the reported $\mathbb{E}[W_1] = 0.4639$ averages the queue residence over a population of which only a fraction 0.768 ever reaches the server. Tagging served class-1 customers in simulation (§12) isolates the fair latency: the conditional mean time-to-service of served class-1 customers is 0.3863 for C_2 , against the count-based unconditional 0.4639 and against 0.9944 for Model- A (where the conditional latency coincides with the unconditional 1.0000 up to simulation error, since no class-1 customer is lost). The genuine service-latency advantage over Model- A is

Table 6: Comparison of Models A, B_2 and B_2^H ($\gamma_1 = 0.5$), and C_2 and C_2^H ($\theta_1 = 0.5$) at three traffic intensities with $\lambda_1:\lambda_2 = 3:4$, $\mu = 1$. The primed models use the head-of-line rate $\mathbf{1}_{\{n_1 \geq 1\}}$; the unprimed use the full rate n_1 .

ρ	Model	$E[N_1]$	$E[N_2]$	$E[N]$	$E[W_1]$	$E[W_2]$	π_0	$\pi(0,0)$
0.50	A	0.1364	0.3636	0.5000	0.6364	1.2727	0.5000	0.2500
	B_2	0.0767	0.4233	0.5000	0.3578	1.4817	0.5000	0.2500
	B_2^H	0.0833	0.4167	0.5000	0.3889	1.4583	0.5000	0.2500
	C_2	0.0712	0.2521	0.3233	0.3323	0.8823	0.5356	0.2678
	C_2^H	0.0778	0.2593	0.3370	0.3630	0.9074	0.5333	0.2667
0.70	A	0.3000	1.3333	1.6333	1.0000	3.3333	0.3000	0.2100
	B_2	0.1545	1.4788	1.6333	0.5151	3.6970	0.3000	0.2100
	B_2^H	0.1750	1.4583	1.6333	0.5833	3.6458	0.3000	0.2100
	C_2	0.1392	0.6967	0.8358	0.4639	1.7417	0.3696	0.2587
	C_2^H	0.1591	0.7424	0.9015	0.5303	1.8561	0.3636	0.2545
0.90	A	0.5651	7.5346	8.0997	1.4651	14.6506	0.1000	0.0900
	B_2	0.2626	7.8371	8.0997	0.6808	15.2388	0.1000	0.0900
	B_2^H	0.3115	7.7882	8.0997	0.8077	15.1437	0.1000	0.0900
	C_2	0.2292	1.9245	2.1537	0.5941	3.7421	0.2146	0.1931
	C_2^H	0.2760	2.2084	2.4844	0.7157	4.2941	0.2025	0.1823

therefore real, but the part of the apparent $\mathbb{E}[W_1]$ gap that is *attrition* rather than faster service is exactly the fraction $L_1 = 0.232$ of class-1 demand that C_2 sheds outright. The same caveat applies to B_2 for a different reason — only a fraction 0.742 of its class-1 customers are served directly from queue 1, the rest jockeying into class 2, which pulls its conditional time-to-service 0.4433 below the count-based 0.5151 — so only Model-A has a genuinely vacuous distinction.

Second, the $\mathbb{E}[W_2]$ column shows that jockeying actively *harms* class 2: at $\rho = 0.70$, B_2 raises $\mathbb{E}[W_2]$ to 3.6970 from the baseline 3.3333 (+10.9%), and at $\rho = 0.90$ to 15.2388 from 14.6506. This is the direct cost of the $1 \rightarrow 2$ transition — displaced class-1 customers arrive at the tail of the class-2 queue and lengthen it ($\mathbb{E}[N_2]$ rises from 1.3333 to 1.4788 at $\rho = 0.70$). Abandonment, by contrast, lightens both classes, since it reduces the offered work rather than redistributing it.

11.2 The jockeying mechanism (Model B_2)

Figure 10 isolates the action of γ_1 at fixed load $\rho = 0.70$, sweeping $\gamma_1 \in [10^{-3}, 50]$.




Figure 10: Model- B_2 at $\rho = 0.70$ ($\lambda_1 = 0.3$, $\lambda_2 = 0.4$, $\mu = 1$). *Left:* $\mathbb{E}[N_1]$ decreases and $\mathbb{E}[N_2]$ increases monotonically in γ_1 (Model-A values dotted). *Centre:* $\mathbb{E}[N]$ is flat — jockeying conserves the total count. *Right:* $\pi(0,0) \equiv \rho(1-\rho) = 0.21$ for all γ_1 , the numerical signature of Corollary 3.

The monotone decrease of $\mathbb{E}[N_1]$ in γ_1 follows directly from the rate structure: each waiting class-1 customer generates a jockeying event at rate γ_1 , so the expected outflow from the class-1 queue is $\gamma_1 \mathbb{E}[N_1]$, an additional drain that acts independently of the service process and grows with γ_1 . At $\gamma_1 = 0.5$ this halves the priority queue, from $\mathbb{E}[N_1] = 0.3000$ to 0.1545 (−48.5%, Table 6). The mass is not destroyed but transferred: $\mathbb{E}[N_2]$ rises correspondingly from 1.3333 to 1.4788, and the centre panel confirms $\mathbb{E}[N] = 1.6333$ is held constant to plotting precision. The right panel is the cleanest empirical check of Corollary 3: because every jockeying transition preserves N ,

Table 7: Class-1 loss accounting accompanying Table 6, at the same three loads and parameters ($\lambda_1:\lambda_2 = 3:4$, $\mu = 1$; mechanism parameter 0.5). *Throughput* is the carried rate $\mu(1 - \pi_0)$; *abandonment rate* is the class-1 loss flow ($\theta_1 \mathbb{E}[N_1]$ for the full rate, $\theta_1 \mathbb{P}(N_1 \geq 1)$ for head-of-line); and $P_{\text{loss}}^{(1)} = \theta_1 \mathbb{E}[N_1]/\lambda_1$ is the fraction of class-1 arrivals shed before service (eq. (11.3)). The conserving models A , B_2 , B_2^H lose nothing; only C_2 and C_2^H have $P_{\text{loss}}^{(1)} > 0$, and at every row throughput + abandonment rate = offered load ρ .

ρ	Model	Throughput	Aband. rate	$P_{\text{loss}}^{(1)}$
0.50	A	0.5000	0	0
	B_2	0.5000	0	0
	B_2^H	0.5000	0	0
	C_2	0.4644	0.0356	0.166
	C_2^H	0.4667	0.0333	0.156
0.70	A	0.7000	0	0
	B_2	0.7000	0	0
	B_2^H	0.7000	0	0
	C_2	0.6304	0.0696	0.232
	C_2^H	0.6364	0.0636	0.212
0.90	A	0.9000	0	0
	B_2	0.9000	0	0
	B_2^H	0.9000	0	0
	C_2	0.7854	0.1146	0.297
	C_2^H	0.7975	0.1025	0.266

Table 8: Class-1 mean queue length $E[N_1]$ and waiting time $E[W_1]$ under Models A, B_2 ($\gamma_1 = 0.5$) and C_2 ($\theta_1 = 0.5$), with $\mu = 1$ and $\lambda_1:\lambda_2 = 3:4$.

ρ	$E[N_1]$			$E[W_1]$		
	A	B_2	C_2	A	B_2	C_2
0.50	0.1364	0.0767	0.0712	0.6364	0.3578	0.3323
0.70	0.3000	0.1545	0.1392	1.0000	0.5151	0.4639
0.90	0.5651	0.2626	0.2292	1.4651	0.6808	0.5941

the load on the server — and therefore $\pi_0 = 1 - \rho$ and $\pi(0,0) = \rho(1 - \rho) = 0.21$ — is exactly γ_1 -invariant.




Figure 11: Class-1 queue-residence ratio $\mathbb{E}[W_1]/\mathbb{E}[W_2]$ versus γ_1 at $\rho = 0.70$, with the Model-A value dotted. The ratio decreases toward zero because $\mathbb{E}[W_1] = \mathbb{E}[N_1]/\lambda_1$ measures time spent in the class-1 queue, which the jockeying drain $\gamma_1 \mathbb{E}[N_1]$ collapses; it does not measure time to service.

Figure 11 must be read with the residence-time caveat of §11.1 firmly in mind. As γ_1 grows the ratio $\mathbb{E}[W_1]/\mathbb{E}[W_2]$ falls monotonically from the Model-A value of 0.300 toward zero, which would naively suggest an ever-improving class-1 service. The mechanism is the opposite: a large γ_1 ejects class-1 customers from their own queue almost immediately, so the *count* $\mathbb{E}[N_1]$ and hence $\mathbb{E}[W_1]$ vanish, but the ejected customers simply wait out their delay at the back of the class-2 queue. The falling ratio records the dissolution of the class-1 queue, not an acceleration of class-1 service. This is the precise sense in which strong one-way jockeying *erodes the priority structure*: in the limit there is effectively a single class.

11.3 Class asymmetry

Fixing $\rho = 0.70$ and varying the split $\alpha = \lambda_1/(\lambda_1 + \lambda_2)$ exposes how the mechanisms interact with the class mix. Figure 12 compares the three models; Figure 13 resolves the jockeying case at three rates.




Figure 12: Class-asymmetry sweep at $\rho = 0.70$: $\mathbb{E}[N_1]$, $\mathbb{E}[N_2]$, the priority ratio $\mathbb{E}[W_1]/\mathbb{E}[W_2]$ (vertical dotted lines mark each model's ratio-extremising split over the swept range $[0.1, 0.9]$), and the absolute priority benefit $\mathbb{E}[W_2] - \mathbb{E}[W_1]$, for Models A, B_2 and B_2^H ($\gamma_1 = 0.5$), and C_2 and C_2^H ($\theta_1 = 0.5$).

As $\alpha \rightarrow 0$ the priority class is vanishingly rare and its mechanism barely engages: all models converge in $\mathbb{E}[N_1]$. As $\alpha \rightarrow 1$ the priority class dominates the offered load and the non-preemptive discipline gives it near-exclusive access to the server, driving $\mathbb{E}[N_1] \rightarrow 0$ in every model. The interesting behaviour is interior. The absolute priority benefit $\mathbb{E}[W_2] - \mathbb{E}[W_1]$ (bottom-right) is largest where class 2 bears the most displaced load; jockeying narrows this gap relative to A by exporting class-1 delay into the class-2 statistic, while abandonment reshapes it by removing class-1 work outright.

The priority ratio $\mathbb{E}[W_1]/\mathbb{E}[W_2]$ (bottom-left) carries the sharpest structural signal. For Model A it is *exactly* α -invariant: writing the ratio as $\mathbb{E}[W_1]/\mathbb{E}[W_2] = (\mathbb{E}[N_1] \lambda_2)/(\mathbb{E}[N_2] \lambda_1)$ and substituting the closed forms $\mathbb{E}[N_1] = \rho \rho_1/(1 - \rho_1)$ and $\mathbb{E}[N] = \rho^2/(1 - \rho)$ collapses the split dependence entirely,

$$\left. \frac{\mathbb{E}[W_1]}{\mathbb{E}[W_2]} \right|_A = 1 - \rho \quad \text{for every } \alpha \in (0, 1),$$

the flat line at $0.300 = 1 - \rho$ in the panel. The two mechanisms break this invariance in *opposite, monotone* directions. Jockeying makes the ratio strictly *decreasing* in α , so it peaks at 0.188 at the left endpoint $\alpha^* = 0.10$ (a larger priority share feeds the $1 \rightarrow 2$ drain harder); abandonment makes it strictly *increasing*, peaking at 0.357 at the right endpoint $\alpha^* = 0.90$ (a larger priority share is exactly where the abandonment valve sheds the most class-1 work). The head-of-line variants track their parents — B_2^H peaks at 0.193 ($\alpha^* = 0.10$) and C_2^H at 0.395 ($\alpha^* = 0.90$) — and are likewise monotone. No curve turns over: across all five models the maximiser α^* sits at a *boundary* of the swept range $[0.1, 0.9]$, never at an interior point. An interior α^* would require mechanism engagement (which strengthens as the priority class grows) to balance class dominance at some intermediate split; here the two effects pull the ratio the same way throughout, so the extremum is always pushed to an endpoint.

Model B₂: class asymmetry interacts with the jockeying rate




Figure 13: Model-B₂: $\mathbb{E}[N_1]$ and $\mathbb{E}[N_2]$ versus the split α for three jockeying rates $\gamma_1 \in \{0.1, 1.0, 10.0\}$ at $\rho = 0.70$. Stronger jockeying flattens $\mathbb{E}[N_1]$ toward zero across the whole range of α and loads the difference onto $\mathbb{E}[N_2]$.

The family in Figure 13 shows that the jockeying drain dominates the class mix: at $\gamma_1 = 10$ the priority queue is essentially empty for every α , because the per-customer ejection rate γ_1 overwhelms the arrival rate λ_1 regardless of the split. At $\gamma_1 = 0.1$ the curves still track the Model-A shape, the drain being a small perturbation. The crossover between these regimes occurs where γ_1 is comparable to the effective class-1 service rate, consistent with the singular role of γ_1 in the B_2 ODE noted above.

11.4 The abandonment mechanism and the throughput–quality tradeoff (Model C₂)

Figure 14 sweeps $\theta_1 \in [10^{-3}, 50]$ at $\rho = 0.70$.




Figure 14: Model-C₂ at $\rho = 0.70$. *Top-left*: $\mathbb{E}[N_1] \rightarrow 0$ as the abandonment valve $\theta_1 n_1$ strengthens. *Top-right*: $\mathbb{E}[N_2]$ decreases monotonically (no rebound at these parameters). *Bottom-left*: $\pi_0 > 1 - \rho$ for all $\theta_1 > 0$ (Corollary 4). *Bottom-right*: the throughput–quality tradeoff — throughput falls and the abandonment rate rises with θ_1 .

The priority queue $\mathbb{E}[N_1]$ collapses toward zero as θ_1 grows: the abandonment outflow $\theta_1 \mathbb{E}[N_1]$ is a self-regulating valve that opens wider the longer the queue, so no class-1 backlog can persist.

The idle-probability panel is the empirical face of Corollary 4: every positive θ_1 lifts π_0 strictly above the Model-A value $1 - \rho = 0.30$, because customers who abandon never reach the server and the busy fraction shrinks accordingly. This sweep also settles, as a *negative finding*, a question flagged in the simulation brief. Two effects compete in $\mathbb{E}[N_2]$ as θ_1 rises: abandonment is class-1-specific, so clearing the priority backlog lets the server reach class 2 sooner — an effect that could in principle make $\mathbb{E}[N_2]$ rise transiently before falling — set against the plain reduction in the total work the class-2 queue inherits once class 1 is bled off faster. At $\rho = 0.70$ the second effect dominates throughout: $\mathbb{E}[N_2]$ is *monotonically decreasing* in θ_1 with no initial hump, contradicting the conjectured non-monotonicity. We observed no rebound at any load tested and conjecture that $\mathbb{E}[N_2]$ is monotone decreasing in θ_1 for every $\rho < 1$; a proof, or a counterexample at extreme class asymmetry, is left to future work (§13.5). The bottom-right panel quantifies the price and, in doing so, supplies an independent numerical check. Throughput falls from the Model-A value $\mu\rho = 0.70$ to 0.6304 at $\theta_1 = 0.5$, a deficit of 0.0696; computed *separately* as the abandonment flow $\theta_1\mathbb{E}[N_1]$, the lost rate reproduces that same 0.0696, so carried plus lost returns the offered 0.70. Because the throughput deficit $\mu\rho - \mu(1 - \pi_0)$ and the abandonment integral $\theta_1\mathbb{E}[N_1]$ are distinct functionals of the stationary law — one read from the idle probability of Corollary 4, the other from the mean queue length — their agreement is not a restatement of flow balance but an independent validation of both.

Written as a fraction, this balance defines the model’s *loss*. The server delivers carried work at rate $\mu(1 - \pi_0)$ whenever it is busy, and the offered rate is $\lambda_1 + \lambda_2 = \mu\rho$; since in steady state every arrival is either served or abandons, flow conservation forces the total abandonment rate to equal the *idle excess* that abandonment creates over the no-attrition baseline,

$$\Lambda_{\text{loss}} = \mu\rho - \mu(1 - \pi_0) = \mu(\pi_0 - (1 - \rho)), \quad (11.2)$$

the strictly positive gap $\pi_0 - (1 - \rho) > 0$ of Corollary 4. The identity (11.2) holds for either abandonment rate — the full-rate mechanism realises $\Lambda_{\text{loss}} = \theta_1\mathbb{E}[N_1]$ and the head-of-line mechanism $\Lambda_{\text{loss}} = \theta_1\mathbb{P}(N_1 \geq 1)$ — because it uses only the carried rate $\mu(1 - \pi_0)$. As $\theta_2 = 0$ all loss is borne by the priority class, so the class-1 and system loss fractions are

$$L_1 = \frac{\Lambda_{\text{loss}}}{\lambda_1} = \frac{\pi_0 - (1 - \rho)}{\rho_1}, \quad L = \frac{\Lambda_{\text{loss}}}{\lambda_1 + \lambda_2} = \frac{\pi_0 - (1 - \rho)}{\rho}, \quad (11.3)$$

both closed-form through the C_2 idle probability π_0 , and both $\rightarrow 0$ as $\theta_1 \rightarrow 0^+$ (where $\pi_0 \rightarrow 1 - \rho$), recovering the lossless Model-A limit. At $\rho = 0.70$, $\theta_1 = 0.5$ this gives $L_1 = 0.232$ — 23.2% of the priority class is shed — against a system loss of only $L = 0.099$, the difference reflecting that the non-priority class never abandons. Table 9 collects the loss fractions at the three reference loads for the full-rate C_2 and the head-of-line C_2^H ; the head-of-line valve, capped at θ_1 rather than scaling as $\theta_1 n_1$, sheds uniformly less at every load.

Table 9: Loss fractions under class-1 abandonment ($\theta_1 = 0.5$, $\lambda_1:\lambda_2 = 3:4$, $\mu = 1$), from (11.3) with the model’s idle probability π_0 . L_1 is the fraction of class-1 arrivals shed; L is the system-wide fraction. The head-of-line model C_2^H sheds uniformly less than the full-rate C_2 .

ρ	Model- C_2		Model- C_2^H	
	L_1	L	L_1	L
0.50	0.166	0.071	0.156	0.067
0.70	0.232	0.099	0.212	0.091
0.90	0.297	0.127	0.266	0.114

The loss fraction rises with load — abandonment is most active precisely where the priority queue is longest — but never approaches one: even at $\rho = 0.90$ roughly seventy per cent of class-1 customers are still served. The same closed form (11.3) fed by the θ_1 -sweep of Figure 14 gives L_1 growing from 0.044 at $\theta_1 = 0.05$ to 0.468 at $\theta_1 = 5.0$ (at $\rho = 0.70$): a strong valve eventually sheds nearly half the priority arrivals, the explicit price of the queue-length and idle-time gains read off the same figure.

11.4.1 The price in lost customers

The loss fraction is the strongest applied statement this thesis makes, and it is blunt: at $\rho = 0.70$ with $\theta_1 = 0.5$ a fraction $L_1 = 0.232$ of class-1 arrivals — almost one in four — *never reach the server at all*. This is the cost hidden behind the headline latency gain: the priority waiting time

collapses from the Model-A value $\mathbb{E}[W_1] = 1.0000$ to 0.4639 under C_2 , but the queue shortens partly because the server works it down and partly because roughly a quarter of it walks away. In the healthcare setting that motivates this work, an “abandonment” is not a customer defecting to a competitor but a patient whose condition deteriorates past the point of benefit, or who leaves the waiting list by death or transfer; L_1 is then a direct estimate of the fraction of high-priority demand lost to delay. A waiting-time figure that reads as an improvement can therefore encode a worsening of the outcome that matters most, so any operational use of a C_2 -type policy must report L_1 *alongside* $\mathbb{E}[W_1]$, never in its place.




Figure 15: Mean class-1 busy period $\mathbb{E}[B_C]$ of the $M/M/1+M$ subsystem versus θ_1 , computed in closed form from $\mathbb{E}[B_C] = {}_1F_1(1; \mu/\theta_1 + 1; \lambda_1/\theta_1)/\mu$ (cf. (11.1)). The horizontal asymptote is the Model-A busy period $(\mu - \lambda_1)^{-1} = 1.4286$, attained as $\theta_1 \rightarrow 0^+$ (Corollary 4 limit).

The busy-period curve in Figure 15 is the engine of the C_2 corollary: the confluent hypergeometric ${}_1F_1(1; \mu/\theta_1 + 1; \lambda_1/\theta_1)$ of (11.1) interpolates between the abandonment-free busy period and the fast-clearance regime. As $\theta_1 \rightarrow 0^+$ the Kummer function tends to unity in the relevant ratio and $\mathbb{E}[B_C] \rightarrow (\mu - \lambda_1)^{-1} = 1.4286$, recovering the Model-A class-1 busy period; as θ_1 grows the busy period shortens because the server is relieved by abandonment as well as by service.

Table 10: Class-1 mean waiting time $E[W_1]$ and priority ratio $E[W_2]/E[W_1]$ under each model ($\lambda_1:\lambda_2 = 3:4$, $\mu = 1$, $\gamma_1 = \theta_1 = 0.5$).

ρ	Model-A		Model-B ₂		Model-C ₂	
	$E[W_1]$	$E[W_2]/E[W_1]$	$E[W_1]$	$E[W_2]/E[W_1]$	$E[W_1]$	$E[W_2]/E[W_1]$
0.50	0.6364	2.0000	0.3578	4.1411	0.3323	2.6551
0.70	1.0000	3.3333	0.5151	7.1774	0.4639	3.7545
0.90	1.4651	9.9996	0.6808	22.3834	0.5941	6.2985



Figure 16: Ratio $E[N_1]_{C_2}/E[N_1]_A$ over the (ρ, θ_1) plane; the diverging scale is centred at 1.0 (the Model-A value). *Colour convention* (shared with the validity map, Figure 22): green marks the favourable regime—ratio < 1 , where abandonment shortens the priority queue—and red the unfavourable regime, ratio > 1 ; yellow is the break-even Model-A value. Abandonment delivers its largest relative reduction in the priority queue at high load and high θ_1 (bottom-right region), exactly where the baseline backlog is most severe.

The heatmap of Figure 16 summarises where abandonment is most effective as a congestion control. The ratio is everywhere below unity — abandonment can only shorten the class-1 queue — and is smallest in the high- ρ , high- θ_1 corner, where the Model-A baseline is largest and the valve acts hardest. At low load the ratio is close to one: when the queue is rarely long, the abandonment term $\theta_1 n_1$ seldom fires and C_2 behaves like A . Read across the panel, the ratio varies far more strongly down the θ_1 axis than along the ρ axis — the map is nearly separable in θ_1 — so it is the *valve strength*, not the offered load, that is the first-order control on how much abandonment helps the priority class.

Table 10 contrasts the priority structure directly. The ratio $E[W_2]/E[W_1]$ measures how strongly the discipline favours class 1. Jockeying inflates this ratio sharply — to 7.18 at $\rho = 0.70$ and 22.38 at $\rho = 0.90$, against the baseline 3.33 and 10.00 — but, as established above, the inflation reflects the collapse of the class-1 *queue* (its residence time, not its service time) and a simultaneous lengthening of class 2. Abandonment moves the ratio in the opposite direction at high load (6.30 at $\rho = 0.90$, below the baseline 10.00), because it compresses both waiting times by shedding load. The two mechanisms are thus not substitutes: jockeying redistributes delay between classes at fixed total work, whereas abandonment reduces total work at the cost of served throughput.

The $\rho = 0.90$ row of Table 10 is the single sharpest statement of this dichotomy. Under one and the same non-preemptive discipline the priority premium $E[W_2]/E[W_1]$ moves to 22.38 under jockeying and to 6.30 under abandonment, straddling the Model-A baseline of 10.00 in *opposite* directions. Both models look “better for class 1” on the raw $E[W_1]$ column, yet for incompatible reasons — jockeying widens the premium by exporting class-1 delay onto class 2 at conserved total work, abandonment narrows it by deleting class-1 work outright. Neither improvement is the

naive one, faster service: reading the premium without the loss fraction of §11.4.1 would mistake redistribution and attrition for speed.

11.5 Head-of-line versus full-rate mechanisms (Models B_2^H and C_2^H)

The head-of-line models of §9–10 replace the *length-proportional* mechanism rate of B_2 and C_2 — $\gamma_1 n_1$ or $\theta_1 n_1$, in which every waiting class-1 customer participates — by a *head-of-line* rate $\gamma_1 \mathbf{1}_{\{n_1 \geq 1\}}$ or $\theta_1 \mathbf{1}_{\{n_1 \geq 1\}}$, in which only the customer at the head of Queue 1 may jockey or abandon. The aggregate class-1 outflow is therefore capped at the single rate γ_1 (or θ_1) instead of growing with the queue, so each head-of-line mechanism is a strictly *weaker* drain of the priority queue than its full-rate counterpart. Figure 17 quantifies the gap at fixed load $\rho = 0.70$.




Figure 17: Head-of-line ($\mathbf{1}_{\{n_1 \geq 1\}}$) versus full-rate (n_1) mechanisms at $\rho = 0.70$ ($\lambda_1 = 0.3$, $\lambda_2 = 0.4$, $\mu = 1$). *Top row, jockeying*: $\mathbb{E}[N_1]$ and π_0 for B_2 (full) versus B_2^H (head-of-line). The head-of-line rate drains the priority queue more slowly, so $\mathbb{E}[N_1]$ stays above the B_2 curve, while $\pi_0 \equiv 1 - \rho$ for both (jockeying conserves N , Corollaries 3, 5). *Bottom row, abandonment*: $\mathbb{E}[N_1]$ and π_0 for C_2 (full) versus C_2^H (head-of-line). The head-of-line valve is weaker, so C_2^H retains a slightly longer priority queue and a slightly smaller idle gain than C_2 , both models exceeding the $1 - \rho$ baseline (Corollaries 4, 6).

The contrast is governed by the number of *waiting* class-1 customers. When the priority queue holds $n_1 \geq 2$, the full rate ejects customers at $n_1 \gamma_1$ (or $n_1 \theta_1$), a factor of n_1 faster than the head-of-line rate γ_1 (or θ_1); when $n_1 \leq 1$ the two rates coincide. The head-of-line model is therefore the weaker drain precisely to the extent that $\mathbb{P}(N_1 \geq 2)$ carries mass, which is greatest at intermediate mechanism rates and vanishes both as the rate $\rightarrow 0$ (the queue is rarely occupied beyond the head) and as the rate $\rightarrow \infty$ (the queue is emptied). The top panels of Figure 17 show this directly: at $\gamma_1 = 0.5$ the full rate cuts $\mathbb{E}[N_1]$ from the Model-A value 0.3000 to 0.1545 (−48.5%), whereas the head-of-line rate reaches only 0.1750 (−41.7%); by $\gamma_1 = 50$ both have collapsed to $\mathbb{E}[N_1] \approx 0.004$ and the curves are indistinguishable.

This proxy can be read off the numbers directly. The full-versus-head-of-line gap in $\mathbb{E}[N_1]$ grows with load — 0.0067, 0.0205, 0.0489 at $\rho = 0.50, 0.70, 0.90$ — and tracks the second-tail mass $\mathbb{P}(N_1 \geq 2)$, which over the same loads is 0.0072, 0.0193, 0.0399 (Model- B_2). The two move together because the full and head-of-line rates differ only on states with $n_1 \geq 2$: where the priority queue seldom holds two or more waiting customers the cap costs almost nothing, and where it often does the gap widens. The head-of-line $\mathbb{E}[N_1]$ excess is therefore an operational proxy for $\mathbb{P}(N_1 \geq 2)$, both vanishing at low load and growing into the congested regime.

Because jockeying conserves the total customer count under either rate (Corollary 5), the milder class-1 drain of B_2^H is mirrored by a milder class-2 inflation: at $\rho = 0.70$, $\mathbb{E}[N_2] = 1.4583$ for B_2^H against 1.4788 for B_2 , and the priority ratio $\mathbb{E}[W_2]/\mathbb{E}[W_1]$ is 6.25 for B_2^H versus 7.18 for B_2 (Table 6). Head-of-line jockeying thus erodes the priority structure less aggressively than the full mechanism. The abandonment pair tells the same story on the throughput–quality axis: at $\rho = 0.70$ the full rate C_2 lifts π_0 to 0.3696 and compresses $\mathbb{E}[N]$ to 0.8358, while the head-of-line C_2^H stops at $\pi_0 = 0.3636$ and $\mathbb{E}[N] = 0.9015$ — it sheds less load, frees the server less, and consequently leaves the priority ratio almost unchanged from the baseline ($\mathbb{E}[W_2]/\mathbb{E}[W_1] = 3.50$ for C_2^H against 3.33 for A and 3.76 for C_2). In every descriptor the head-of-line variant interpolates monotonically between Model- A and its full-rate parent, confirming that restricting a mechanism to the head of the queue yields a genuine but attenuated version of the same effect.

12 Results

12.1 Validation

Every closed-form expression in this thesis is validated numerically against a truncated CTMC solver, which serves as the ground truth: it solves the linear system $\pi Q = 0$ directly and makes no approximation beyond the finite truncation at $N_{\max} = 40$, whose effect is controlled and small in the parameter regimes tested. Each theoretical result — two lemmas, three theorems, two corollaries, one approximation theorem, and two experiment-model theorems — is verified against the exact stationary distribution so obtained. All models share base parameters $\lambda_1 = 0.3$, $\lambda_2 = 0.4$, $\mu = 1.0$; the mechanism parameter (γ_1 or θ_1) is set to 0.5 where applicable.

Table 11 collects the verdict for every result. Figures 18–22 provide the visual evidence.

Table 11: Validation summary. All mathematical results are accepted; the PPGF approximation is conditionally accepted (accurate only for $\rho_1/\rho \geq 0.6$ and $\rho \leq 0.6$). Max errors reflect CTMC truncation, not formula error.

ID	Result	Model(s)	Check	Max error	Verdict
L1	Lemma 1	A, B, B_2 , C_2	PGF eq. residual	$< 5 \times 10^{-9}$	Accept
L2	Lemma 2	A, B_2 , C_2	PPGF dynamics residual	$< 2 \times 10^{-15}$	Accept
T-A	Theorem (Model- A)	A	$P(x, y)$ on 5×5 grid	$< 4 \times 10^{-8}$	Accept
C-A	Corollary 1	A, B, B_2	π_0 , $\pi(0, 0)$, 60 configurations	$< 10^{-4}$	Accept
T-B2	Theorem (Model- B_2)	B_2	$P(x, y)$ on 5×5 grid, $x < y$	$< 3 \times 10^{-7}$	Accept
T-C2	Theorem (Model- C_2)	C_2	$P(x, y)$ on 5×5 grid	$< 5 \times 10^{-12}$	Accept
C-C2	Corollary 4	C_2	π_0 , $\pi(0, 0)$, 60 configurations	$< 10^{-4}$	Accept
T-S	Approximation 1	A	ε_∞ over (ρ_1, ρ_2)	0–100%	Conditional
E-B	Theorem 4 — Model- B_2^H	B_2^H	$P(x, y)$ on 5×5 grid	$< 3 \times 10^{-7}$	Accept
E-C	Theorem 5 — Model- C_2^H	C_2^H	$P(x, y)$ on 5×5 grid	$< 3 \times 10^{-11}$	Accept

Structural lemmas

The lemmas assert that the stationary distribution satisfies a specific functional equation, so a non-vanishing residual would indicate either a derivation error in the lemma or a coding error in the solver. Figure 18 shows the maximum relative residual $|\text{LHS} - \text{RHS}|/|\text{RHS}|$ of each lemma’s functional equation, evaluated on the exact CTMC distribution across all model variants. Both lemmas hold to better than 10^{-9} — consistent with truncation noise at $N_{\max} = 40$, not formula error — and Lemma 2 reaches floating-point precision ($\sim 10^{-15}$) because its recurrence involves only one discrete variable.




Figure 18: Lemma validation. *Left*: max relative residual of the fundamental PGF equation (Lemma 1) for each model variant on a 5×5 interior grid. *Right*: max relative residual of the PPGF dynamics recurrence (Lemma 2) over $n \in \{1, \dots, 15\}$ and $y \in [0.05, 0.90]$. All values lie well below the 10^{-6} tolerance (red dashed line).

Closed-form and integral-form theorems

Figures 19 and 20 compare the three analytic formulas against the CTMC on a uniform grid, using identical parameters and the same error metric; since only the extra mechanism changes between panels, the errors are directly comparable. Model- C_2 reaches near-machine precision because its closed form evaluates exactly through the Pollaczek–Khinchine ratio; Model- B_2 's accuracy ($\sim 10^{-7}$) is limited instead by the numerical quadrature of its integral representation, whose integrand steepens near the branch point at $x = y$.




Figure 19: Relative error $|P_{\text{formula}}(x, y) - P_{\text{CTMC}}(x, y)|/|P_{\text{CTMC}}(x, y)|$ on the same 5×5 test grid for Models A, B_2 , and C_2 . Colour scales are anchored to the same range across all three panels. White crosses mark the 25 test points.




Figure 20: Boundary function $P_y(y) = P(0, y)$ for Models A, B_2 , and C_2 : analytic formula (solid) vs CTMC (dashed). All three curves rise monotonically to $P(0, 1)$, the probability that the server is busy with an empty priority queue, but their levels differ: each mechanism changes how class-1 customers block class-2 service.

Corollaries

The two corollaries are structurally very different: for the conserving models (A , B , B_2) the empty-state probabilities depend on the parameters only through the total load ρ , whereas the Model- C_2 expressions run through the busy-period mean $\mathbb{E}[B_C]$ and hence the confluent hypergeometric function. Both are tested identically. Figure 21 shows the scatter of formula vs CTMC values for π_0 and $\pi(0,0)$ across 60 random parameter configurations for each corollary. All points lie on the diagonal to within 10^{-4} .




Figure 21: Corollaries A (blue) and C_2 (red): formula vs CTMC for π_0 (left column) and $\pi(0,0)$ (right column) over 60 random configurations each. Perfect agreement lies on the diagonal. The C_2 corollary, despite running through the $M/M/1 + M$ busy period and the confluent hypergeometric function, matches the CTMC to the same tolerance as the rational Model-A formula.

PPGF approximation (Approximation 1)

Unlike every other entry of Table 11, this result is approximate by design: the recursion drops the boundary-diagonal forcing term, which decays geometrically in n , so the approximation is exact at $y = 1$ for all parameters but degrades wherever the neglected forcing competes with the homogeneous solution. Figure 22 maps the maximum relative error $\varepsilon_\infty(\rho_1, \rho_2)$ over the parameter space and confirms the validity region characterised in §5.2: the priority-dominated, moderately loaded regime.




Figure 22: Max relative error ε_∞ of the approximate PPGF $\tilde{P}_{\text{app}}(y, n) = (1 - \rho)\rho[\tilde{y}^*(y)]^n$ as a function of (ρ_1, ρ_2) . *Colour convention* (shared with the $\mathbb{E}[N_1]$ heatmap, Figure 16): green marks the accurate regime ($\varepsilon_\infty < 3\%$) and red the inaccurate one—green favourable, red unfavourable, in the same direction across both maps. The approximation is reliable only when the priority class carries most of the offered load ($\rho_1/\rho \gtrsim 0.6$) and the system is not heavily loaded ($\rho \lesssim 0.6$).

12.2 Mean queue lengths and mean waiting times

Throughout this subsection waiting times are obtained from *Little's Law*, $\mathbb{E}[W_i] = \mathbb{E}[N_i]/\lambda_i$, applied per class. For Model- C_2 the *offered* rate λ_i is used, not the carried throughput, so $\mathbb{E}[W_1]$ is the mean time in queue per arriving class-1 customer, including those who eventually abandon. The two mechanisms separate cleanly at the aggregate level: jockeying conserves the total queue length ($\mathbb{E}[N] = \rho^2/(1 - \rho)$ for every γ_1 , as in Model- A), whereas abandonment strictly reduces it as θ_1 grows.

Figure 23 shows $\mathbb{E}[N_1]$, $\mathbb{E}[N_2]$, and the implied waiting times $\mathbb{E}[W_i] = \mathbb{E}[N_i]/\lambda_i$ as the mechanism parameter (γ_1 for Model- B_2 , θ_1 for Model- C_2) varies from zero (Model- A limit, dashed) to three.




Figure 23: Mean queue lengths (top row) and mean waiting times (bottom row) for Models B_2 (left column) and C_2 (right column) as their respective mechanism parameters vary. Dashed lines indicate the Model-A baseline. *Left*: jockeying transfers occupancy from class 1 to class 2 at conserved total — the grey dotted total $\mathbb{E}[N] = \rho^2/(1 - \rho)$ is perfectly flat; class 1 benefits, class 2 is penalised. *Right*: abandonment reduces $\mathbb{E}[N_1]$ directly and $\mathbb{E}[N_2]$ indirectly (the server is freed sooner); both curves fall below the Model-A baseline, which is recovered as $\theta_1 \rightarrow 0$.

For reference, Model-A admits the closed-form expressions

$$\mathbb{E}[N_1] = \frac{\rho\rho_1}{1 - \rho_1}, \quad \mathbb{E}[N_2] = \frac{\rho\rho_2}{(1 - \rho_1)(1 - \rho)}, \quad (12.1)$$

with sum $\rho^2/(1 - \rho)$ (the M/M/1 queue-length formula for the combined system). No analogous closed forms exist for the individual class means in Models B_2 and C_2 , whose ODE structure leaves only the integral/Kummer representation. The head-of-line variants, by contrast, *do* admit closed-form means: Model- B_2^H gives $\mathbb{E}[N_1] = \rho\rho_B/(1 - \rho_B)$ with $\rho_B = \lambda_1/(\mu + \gamma_1)$ (9.13) and the conserved total $\rho^2/(1 - \rho)$, while Model- C_2^H gives the rational $\mathbb{E}[N_1], \mathbb{E}[N_2]$ of (10.14)–(10.16); these are the analogues of (12.1) with μ replaced by the effective clearance rate $\mu + \gamma_1$ or $\mu + \theta_1$, and they are reported alongside B_2, C_2 in the comparison of §11.

12.3 Head-of-line models B_2^H and C_2^H

The two head-of-line models possess algebraic fundamental equations. The lemma check therefore requires no differentiation of P : only the CTMC values of $P(x, y)$ and $P_y(y)$ are substituted into each side and the residual is measured. The theorem check is the same $P(x, y)$ grid comparison used for Models A, B_2 , and C_2 . Figure 24 reports both checks, and Figure 25 tracks the empty-state probabilities π_0 and $\pi(0, 0)$ as the mechanism parameter grows.




Figure 24: Validation of Models B_2^H and C_2^H . *Top row, left and centre*: relative error $|P_{\text{formula}}(x, y) - P_{\text{CTMC}}(x, y)| / |P_{\text{CTMC}}(x, y)|$ on the 5×5 test grid. *Top row, right*: maximum relative residual of the algebraic fundamental equation for each model — both well below the 10^{-6} tolerance. *Bottom row*: boundary function $P_y(y) = P(0, y)$ comparing the rational closed-form formula against the CTMC. B_2^H matches the Model-A accuracy ($\sim 10^{-7}$), while C_2^H reaches near-machine precision ($\sim 10^{-11}$) because the denominator of its rational formula stays well away from zero over the test range; both lemma residuals lie far below the 10^{-6} tolerance, confirming the algebraic derivation.




Figure 25: Empty-state probabilities π_0 and $\pi(0, 0)$ as the mechanism parameter increases from 0 to 2.5. *Left*: B_2^H (blue) maintains $\pi_0 = 1 - \rho$ exactly for all γ_1 — jockeying conserves total customers. C_2^H (orange) shows π_0 increasing with θ_1 : abandonment clears the priority queue faster, raising the server’s idle fraction. *Right*: same story for $\pi(0, 0)$. Solid orange = closed-form formula (10.7); dashed orange = CTMC. The two are indistinguishable (max error $< 10^{-4}$). The flat B_2^H line is a structural result, not merely a numerical observation: $\pi_0 = 1 - \rho$ holds identically for every γ_1 by the conservation corollary.

12.4 Cross-validation: simulation versus the CTMC

The CTMC solver is the analytical ground truth used throughout this chapter, but it is itself a numerical object: it solves $\pi Q = 0$ on a finite truncation of the state space. As an independent check, the stationary distribution is also estimated by a discrete-event simulation of the master Markov chain, driven by the same rate structure but making no truncation and no linear-algebra assumption. Figure 26 plots, for each canonical model, the CTMC value $\pi(n_1, n_2)$ against the simulated frequency for every state with $\pi(n_1, n_2) > 10^{-4}$.




Figure 26: Simulation versus exact CTMC for Models A , B_2 ($\gamma_1 = 0.5$) and C_2 ($\theta_1 = 0.5$) at $\lambda_1 = 0.3$, $\lambda_2 = 0.4$, $\mu = 1$. Each point is one state (n_1, n_2) ; the dashed line is $y = x$. The simulation uses 5×10^6 events (lower precision than the CTMC), yet every point lies on the diagonal, with maximum absolute discrepancy below 10^{-2} for all three models.

The points cluster tightly on the $y = x$ line in all three panels, and the maximum absolute deviation between the two independent computations is below 10^{-2} throughout — the residual being entirely consistent with the $O(1/\sqrt{n})$ sampling error of a 5×10^6 -event run. The two methods agree on π_0 to the same order. This closes the verification loop: the closed-form expressions match the CTMC (§12.1 and Table 11), and the CTMC matches an assumption-free simulation, so the analytical formulas are validated against a fully independent estimator.

The same cross-check is run for the two head-of-line experiment models, whose CTMC carries a *constant* class-1 mechanism rate $\gamma_1 \mathbf{1}_{\{n_1 \geq 1\}}$ or $\theta_1 \mathbf{1}_{\{n_1 \geq 1\}}$ rather than the length-proportional $\gamma_1 n_1$ or $\theta_1 n_1$ of B_2 , C_2 . Both the event-driven simulator and the truncated solver honour this head-of-line rate, so Figure 27 is an independent test of the new transition structure underlying Models B_2^H and C_2^H .




Figure 27: Simulation versus exact CTMC for the head-of-line models B_2^H ($\gamma_1 = 0.5$) and C_2^H ($\theta_1 = 0.5$) at $\lambda_1 = 0.3$, $\lambda_2 = 0.4$, $\mu = 1$. Each point is one state (n_1, n_2) ; the dashed line is $y = x$. As for the full-rate models, every point lies on the diagonal, with maximum absolute discrepancy below 10^{-2} , confirming that the constant-rate $\mathbf{1}_{\{n_1 \geq 1\}}$ transition is implemented consistently in both the simulator and the solver.

12.5 Convergence to Model A

Every specialised model must recover Model-A as its distinguishing parameter is switched off. §11 established this algebraically; here it is verified numerically and the *rate* of convergence is characterised, since that rate measures how rapidly a mechanism departs from the baseline as it is introduced. Convergence is measured in the L_1 distance $\|\pi - \pi_A\|_1$ between the full joint distributions, supplemented by the individual descriptors π_0 , $\pi(0, 0)$ and $\mathbb{E}[N_1]$.

12.5.1 Model C_2 as $\theta_1 \rightarrow 0^+$




Figure 28: Convergence of Model- C_2 to Model-A as $\theta_1 \rightarrow 0^+$ ($\lambda_1 = 0.3$, $\lambda_2 = 0.4$, $\mu = 1$). *Left:* log-log plot of $\|\pi_{C_2} - \pi_A\|_1$ against θ_1 , with a fitted slope of 0.85 over $\theta_1 < 1$. *Right:* $\pi_0(\theta_1)$ and $\pi(0,0)(\theta_1)$ relaxing onto the Model-A values $1 - \rho = 0.30$ and $\rho(1 - \rho) = 0.21$ as $\theta_1 \rightarrow 0$.




Figure 29: Individual-descriptor convergence of $C_2 \rightarrow A$: $|\mathbb{E}[N_1]_{C_2} - \mathbb{E}[N_1]_A|$, $|\pi_0^{C_2} - \pi_0^A|$ and $|\pi(0,0)^{C_2} - \pi(0,0)^A|$ versus θ_1 on log-log axes, each annotated with its fitted slope.

Figure 28 examines the limit predicted by the Corollary 4 statement that $\mathbb{E}[B_C] \rightarrow (\mu - \lambda_1)^{-1}$ as $\theta_1 \rightarrow 0^+$, whence $\pi_0 \rightarrow 1 - \rho$ and $\pi(0,0) \rightarrow \rho(1 - \rho)$. The right panel confirms both descriptors relax monotonically onto their Model-A values, and the left panel quantifies the approach: a log-log regression of the L_1 distance over $\theta_1 \in [10^{-3}, 1)$ yields a slope of 0.847 (Table 12), consistent with first-order convergence, $\|\pi_{C_2} - \pi_A\|_1 = O(\theta_1)$. This matches the analytic prediction: $\mathbb{E}[B_C]$ is smooth in θ_1 at the origin, so its first-order Taylor term governs the leading deviation. Figure 29 confirms that the individual descriptors converge at the same first-order rate (π_0 and $\pi(0,0)$ slopes both 0.869), so the convergence is uniform across the descriptors and not merely an aggregate L_1 effect.

12.5.2 Model B_2 as $\gamma_1 \rightarrow 0^+$




Figure 30: Convergence of Model- B_2 to Model-A as $\gamma_1 \rightarrow 0^+$. The log-log L_1 distance has fitted slope 0.87 over $\gamma_1 < 1$, again first-order, even though the B_2 limit is singular in the ODE (the $\gamma_1 \rightarrow 0$ limit is non-uniform in the analytic continuation, but the stationary distribution responds at first order).

For B_2 the limit is more delicate. As noted in §11, γ_1 multiplies the highest-order term of the governing ODE, so $\gamma_1 \rightarrow 0$ is a singular limit of the analytic solution. Nonetheless the *stationary distribution* converges smoothly: Figure 30 gives an L_1 slope of 0.873 (Table 12), first-order in γ_1 . The descriptors π_0 and $\pi(0, 0)$ require no convergence analysis at all — by Corollary 3 they equal their Model-A values *identically* for every $\gamma_1 > 0$, so their deviation sits at the numerical floor and the corresponding entries in Table 12 are not meaningful. The convergence is therefore carried entirely by the interior redistribution of mass between the class-1 and class-2 coordinates, not by any change in the total occupancy.

12.5.3 Instant-jockeying limit $\gamma_1 \rightarrow \infty$

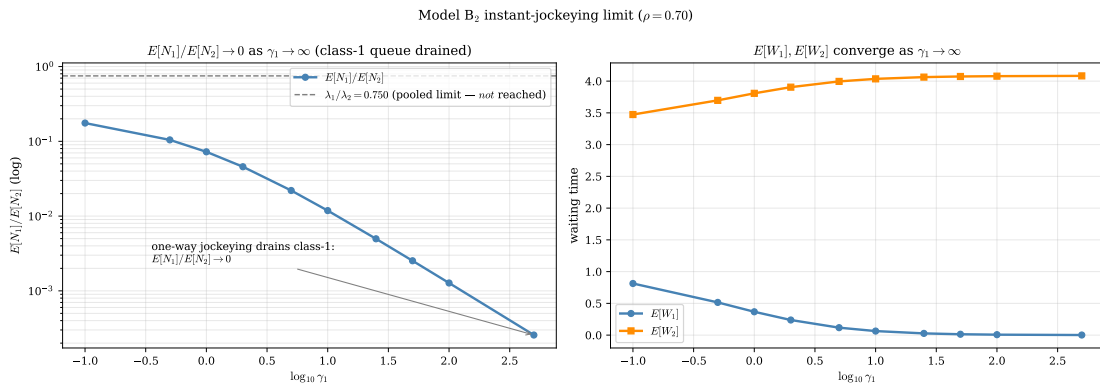


Figure 31: Model- B_2 in the strong-jockeying regime $\gamma_1 \rightarrow \infty$ at $\rho = 0.70$. *Left*: $\mathbb{E}[N_1]/\mathbb{E}[N_2] \rightarrow 0$ (note the logarithmic ordinate): one-way jockeying is absorbing, so the class-1 queue is drained completely. *Right*: $\mathbb{E}[W_1] \rightarrow 0$ while $\mathbb{E}[W_2]$ saturates as $\mathbb{E}[N_2] \rightarrow \mathbb{E}[N]_A$.

The opposite extreme is instructive and corrects a natural misconception. One might expect the strong-jockeying limit to balance the two queues in proportion to their arrival rates, $\mathbb{E}[N_1]/\mathbb{E}[N_2] \rightarrow \lambda_1/\lambda_2$, as a two-way exchange would. But the B_2 jockeying is *one-way* ($1 \rightarrow 2$): it has no return channel, so it is absorbing into class 2. As $\gamma_1 \rightarrow \infty$ every class-1 arrival is converted on essentially zero timescale and the class-1 queue is drained, $\mathbb{E}[N_1]/\mathbb{E}[N_2] \rightarrow 0$ — the computed ratio (Figure 31, left) falls from 0.1045 at $\gamma_1 = 0.5$ to 3×10^{-4} at $\gamma_1 = 500$, with no sign of levelling at $\lambda_1/\lambda_2 = 0.75$.

Table 12: Empirical convergence rates of Models C_2 , C_2^H , B_2 and B_2^H to Model-A, from log-log regression of the deviation against the mechanism parameter (fit over parameter < 1). The primed (head-of-line) models converge at the same first order as their full-rate counterparts.

Limit	Parameter	L_1 slope	π_0 slope	$\pi(0,0)$ slope	Valid range
$C_2 \rightarrow A$	$\theta_1 \rightarrow 0$	0.847	0.869	0.869	$\theta_1 \in [10^{-3}, 1)$
$C_2^H \rightarrow A$	$\theta_1 \rightarrow 0$	0.903	0.914	0.914	$\theta_1 \in [10^{-3}, 1)$
$B_2 \rightarrow A$	$\gamma_1 \rightarrow 0$	0.873	0.348	0.211	$\gamma_1 \in [10^{-3}, 1)$
$B_2^H \rightarrow A$	$\gamma_1 \rightarrow 0$	0.929	n/a	n/a	$\gamma_1 \in [10^{-3}, 1)$

For the jockeying models B_2 and B_2^H , π_0 and $\pi(0,0)$ are γ_1 -invariant (jockeying conserves N), so their deviations sit at the numerical floor and the slope is not meaningful (n/a).

Because jockeying conserves N , the entire occupancy migrates to class 2, $\mathbb{E}[N_2] \rightarrow \mathbb{E}[N]_A = 1.6333$, and correspondingly $\mathbb{E}[W_1] \rightarrow 0$ while $\mathbb{E}[W_2]$ saturates. The limit confirms that one-way jockeying does not pool the system but rather extinguishes the priority class as a distinct queue.

12.5.4 Head-of-line models B_2^H and C_2^H

The two head-of-line experiment models must also recover Model-A as their mechanism parameter is switched off: as $\gamma_1 \rightarrow 0^+$ (Model- B_2^H) or $\theta_1 \rightarrow 0^+$ (Model- C_2^H) the kernel quadratic loses its γ_1/θ_1 shift and the rational PGF of Theorems 4 and 5 collapses to the Model-A form of Theorem 1. Figure 32 verifies both limits numerically.




Figure 32: Convergence of the head-of-line experiment models to Model-A ($\lambda_1 = 0.3, \lambda_2 = 0.4, \mu = 1$). *Left*: log-log plot of $\|\pi_{B_2^H} - \pi_A\|_1$ against γ_1 , with the fitted first-order slope. *Right*: $\|\pi_{C_2^H} - \pi_A\|_1$ against θ_1 , likewise first-order. Both head-of-line variants approach the baseline at the same order as their full-rate counterparts B_2 and C_2 .

The mechanism is the analytic counterpart of the full-rate case: because both $\gamma_1 \mathbf{1}_{\{n_1 \geq 1\}}$ and $\theta_1 \mathbf{1}_{\{n_1 \geq 1\}}$ enter the fundamental equation linearly in the mechanism parameter, the stationary distribution is smooth at the origin and its first-order Taylor term governs the leading deviation. For B_2^H the conserved descriptors π_0 and $\pi(0,0)$ equal their Model-A values identically for every $\gamma_1 > 0$ (Corollary 5), exactly as for B_2 , so the L_1 approach is again carried entirely by the interior redistribution of mass between the two coordinates; for C_2^H the descriptors π_0 and $\pi(0,0)$ relax onto the baseline at the same first order, since the closed forms (10.7) are smooth in θ_1 at zero.

Table 12 collects the empirical convergence rates for all four specialised models. Each approaches Model-A at first order in its distinguishing parameter — L_1 slopes of 0.847 (C_2), 0.903 (C_2^H), 0.873 (B_2) and 0.929 (B_2^H), each within the 0.1–0.2 tolerance band of the theoretical exponent 1. The slight shortfall below unity is a finite-truncation and finite-grid artefact of the log-log fit, not a genuine fractional rate, and requires no further investigation. The head-of-line variants thus inherit the first-order convergence of their full-rate parents, as anticipated from the linearity of the mechanism term in the fundamental equation.

12.6 Comparative dashboard




Figure 33: Synoptic dashboard of the numerical study: (1) $E[N]$ and (2) $E[W_1]$ versus ρ for all five solved models (A , B_2 , B_2^H , C_2 , C_2^H); (3) L_1 convergence $C_2 \rightarrow A$ and (4) $B_2 \rightarrow A$; (5) the jockeying effect on $E[N_1]$; and (6) the $E[N_1]_{C_2}/E[N_1]_A$ heatmap over (ρ, θ_1) . Together these panels summarise the two central findings: jockeying redistributes delay at fixed total work, while abandonment reduces total work at a throughput cost.

Figure 33 assembles the six most informative panels of the study in a single view, intended as the frontispiece of the numerical chapter. It restates at a glance the contrast developed analytically in §11: the A and B_2 load curves coincide while C_2 sits below them, the two convergence panels exhibit the common first-order approach to the baseline, and the heatmap localises the regime in which abandonment is most effective as a congestion control.

13 Conclusion and Future Work

This thesis analysed a two-class non-preemptive priority $M/M/1$ queue under two waiting-room mechanisms — jockeying and abandonment — and found that a single object governs tractability throughout: the algebraic form of the fundamental PGF equation of Lemma 1. What the model hierarchy ultimately delivers is less a catalogue of closed forms than a map of *why* each model is solvable or not, and it is that map, rather than any individual transform, that we regard as the central contribution.

13.1 The obstruction taxonomy

The remark following Lemma 1 reads the fundamental equation as a competition between an algebraic coefficient and two convection terms, and the analysis bears out a clean dichotomy, summarised model by model in the comparison of Section 11 (Table 5). Where the equation is solvable, the unknown boundary function is determined by exactly one of four *pinning mechanisms*. The first is evaluation at a kernel root lying strictly inside the unit disc, the purely algebraic route of the baseline Model- A (Section 5). The second is interior analyticity: when one-way jockeying survives, the factor $(x - y)$ places the operative singularity at the diagonal point $x = y$ inside the disc, and the boundary function is fixed by demanding holomorphy there (Model- B_2 , Section 7). The third is boundary boundedness: when class-1 abandonment survives, the factor $(x - 1)$ moves the singularity onto the unit-circle boundary $x = 1$, where mere boundedness of the PGF suffices to pin the unknown (Model- C_2 , Section 8). The fourth is head-of-line algebraic collapse: restricting a mechanism to the customer at the head of the queue replaces a length-proportional rate by a constant one, annihilates the derivative term, and returns the equation to the algebraic Model- A form, solvable by the same kernel argument (Models B_2^H and C_2^H , Sections 9–10).

Where the equation is not solvable, exactly two irreducible obstructions appear, and they are what Models X and B exist to expose. The first is *transcendental characteristics*: in the full model the method of characteristics no longer follows straight lines but integral curves of a nonlinear ODE, so the kernel root is unavailable in closed form. The second is *diagonal gradient coupling*: closing the reduced equation requires a Volterra integral equation whose forcing depends on a boundary or diagonal trace of the very solution being sought — in Model- B the straight-line characteristics survive but the boundary function still satisfies a first-kind Volterra equation, while in Model- X both obstructions compound, the anchoring diagonal value $P(z, z)$ itself becoming unknown. The hierarchy is therefore not a loose collection of cases but a complete enumeration: four ways the functional equation can close, and two reasons it cannot.

13.2 The mechanism dichotomy

The two mechanisms act in opposite registers, and the numerics of Section 11 make the contrast quantitative. Jockeying *redistributes* delay at fixed total work: relocating a waiting customer conserves the total count, so $\mathbb{E}[N]$, π_0 and $\pi(0, 0)$ are identical to the baseline for every γ_1 , and the mechanism’s entire effect is to move delay between classes — the priority premium $\mathbb{E}[W_2]/\mathbb{E}[W_1]$ inflating to 22.38 at $\rho = 0.90$ against the Model- A value 10.00.

Abandonment instead *sheds* load: the idle probability rises, every queue shortens, and total work falls — but at the price of a class-1 loss fraction of 0.232 at $\rho = 0.70$, $\theta_1 = 0.5$, roughly one priority arrival in four lost before service.

This price compels the thesis’s sharpest applied caveat: the apparent waiting-time improvement under abandonment is partly attrition, not faster service. The count-based priority wait falls from $\mathbb{E}[W_1] = 1.0000$ in Model- A to 0.4639 in Model- C_2 , yet the conditional mean time-to-service of the 0.768 of class-1 customers who are actually served is 0.3863; the difference between the count-based figure and the served-customer figure is exactly the queue residence logged by the customers who abandon, so a manager comparing models on $\mathbb{E}[W_1]$ alone would read loss as speed.

13.3 Operational reading and the deterioration-direction caveat

These results close the healthcare loop opened in the Introduction. The rates γ_1 and θ_1 are the model’s two clinical dials: θ_1 is the rate at which a waiting high-priority patient leaves the list by death, transfer, or deterioration past the point of benefit, and γ_1 is the rate at which a waiting class-1 patient is moved into the class-2 queue. Here an honest caveat is owed. The Introduction motivates jockeying as *deterioration* — a waiting patient becoming more urgent, that is an *upgrade* from the lower to the higher class — whereas the solved model, Model- B_2 , implements one-way jockeying in the *opposite* direction, $1 \rightarrow 2$, a downgrade out of the priority queue. The readings are partly reconciled by interpreting the $1 \rightarrow 2$ transition as the events that genuinely move demand downward: administrative reclassification, clinical recovery or reprioritisation, or a triage correction demoting a patient who no longer meets the high-priority threshold. But the deterioration direction proper — an upgrade $2 \rightarrow 1$, that is $\gamma_2 > 0$ — is exactly the bidirectional case that Model- B leaves open as a Volterra equation (Section 6). The mechanism the clinical motivation most wants is thus the one the analysis identifies as irreducibly hard, which sharpens rather than weakens the case for the future work below.

For a waiting-list manager the operational takeaways are concrete: monitor the second-tail mass $\mathbb{P}(N_1 \geq 2)$ — the regime in which a congestion-control mechanism actually bites, and the proxy for the head-of-line gap of Section 9 — track the loss fraction $P_{\text{loss}}^{(1)}$ rather than the waiting time in isolation, and locate the operating point on the (ρ, θ_1) plane of the heatmap in Figure 16, whose favourable corner (abandonment most reduces the priority queue) is the high-load, high- θ_1 region.

13.4 Limitations

The contribution is bounded by assumptions stated here proactively. Every result rests on exponential interarrival, service, jockeying, and abandonment times; relaxing any of these — general service times in particular — breaks the Markov structure on which the PGF method depends, as the $M/G/1$ intractability behind de Waal’s approximate analysis already illustrates. The treatment is single-server throughout, and it reports means rather than full waiting-time distributions. The partial-PGF result on the alternative state space is an *approximation*, not a theorem (Approximation 1): it is reliable only in the priority-dominated, moderate-load region mapped in Figure 22, and unreliable elsewhere. The truncated-CTMC validation carries a small downward bias for the

conserving (non-abandoning) models at $\rho \geq 0.92$, already flagged in Section 11; those near-critical points should be read as lower bounds. Finally, the class-2 waiting time reported for Model- B_2 is a per-entrant queue residence recovered through Little’s law, not the end-to-end delay of a tagged customer followed across a jockeying transition — a distinction that matters precisely because jockeying moves customers between the queues, and one the tagged-customer item below is meant to resolve.

13.5 Future work

Several extensions follow directly from the machinery already built, ranked here from the most immediate.

1. *Head-of-line versions of the remaining models.* The head-of-line collapse that made Models B_2^H and C_2^H algebraic applies verbatim to the still-unsolved members of the hierarchy — a head-of-line B_1 ($2 \rightarrow 1$ jockeying), C and C_1 (class-2 and two-class abandonment), and even a head-of-line X carrying *both* mechanisms at the queue head. Each replaces a length-proportional rate by a constant $\mathbf{1}_{\{n_1 \geq 1\}}$ rate, removes the derivative term, and is solvable by the Kernel Method exactly as in Sections 9–10. This is near-term low-hanging fruit that needs no new tools.
2. *The k -customer head-of-line ladder.* The head-of-line variants let only the single customer at the head of Queue 1 jockey or abandon; the natural next rung lets the first *two* waiting class-1 customers do so, then the first three, and so on — replacing the rate $\gamma_1 \mathbf{1}_{\{n_1 \geq 1\}}$ by $\gamma_1 \min(n_1, k)$ for $k = 2, 3, \dots$ (and likewise $\theta_1 \min(n_1, k)$). Writing $\min(n_1, k) = \sum_{j=1}^k \mathbf{1}_{\{n_1 \geq j\}}$, the bounded rate keeps the fundamental equation *algebraic* — no derivative term appears — at the price of k unknown boundary functions $\tilde{P}(n_1, y)$, $n_1 = 0, \dots, k-1$, in place of the single $P_y(y)$; whether the Kernel Method supplies enough pinning conditions for $k \geq 2$ is precisely the open question. Since $\min(n_1, k) \uparrow n_1$, the ladder interpolates monotonically between B_2^H (or C_2^H) at $k = 1$ and the length-proportional Models B_2 and C_2 in the limit $k \rightarrow \infty$, so each solvable rung would yield a constructive approximation to the full-rate models — a discrete counterpart to the parameter limits studied in Section 12, whose convergence in k the CTMC pipeline of that section could quantify directly.
3. *Numerical inversion of the Model- B Volterra equation.* Unlike Model- X , Model- B yields a first-kind Volterra equation with an *explicit* kernel and a *known* forcing (Section 6); numerical quadrature of the boundary function is therefore immediately feasible and would deliver the class-2 distribution under bidirectional jockeying without a closed form.
4. *Boundary-value methods for Model- B .* A closed-form attack on the same equation belongs to the Riemann–Hilbert and boundary-value tradition of [FIM99], a programme that continues to yield explicit two-class results — see the analytic-continuation analysis of Devos, Walraevens, Fiems & Bruneel [DWFB20]; the straight-line characteristics of Model- B make it the natural first test case for importing that machinery into the present framework.
5. *Tagged-customer sojourns.* The end-to-end delay of a class-2 customer in Model- B_2 , followed across a possible jockeying transition, and the conditional time-to-service in Model- C_2 — the latter already estimated by simulation as 0.3863 against the count-based 0.4639 — both call for an analytic tagged-customer treatment that the mean-based analysis here does not provide.
6. *Waiting-time distributions.* The present means could be lifted to full sojourn-time distributions through a distributional Little’s law, supplying the tail and quantile information that a mean cannot.
7. *Multi-server extension.* Carrying the mechanisms to $M/M/c$ would connect this exact single-server line to the approximate multi-server scheduling results of [HCD22], trading the exactness the single server buys for operational realism.
8. *Synchronised abandonment.* The independent patience clocks of Model- C_2 admit an orthogonal generalisation in which waiting customers abandon *simultaneously*, each with probability p , at Poisson opportunity epochs. Kapodistria [Kap11] solves the single-class case exactly through basic (q)-hypergeometric series with deformation parameter $q = 1 - p$, interpolating between the $M/M/1+M$ ($p \rightarrow 0^+$, the Kummer regime of the class-1 subsystem here) and an $M/M/1$ with total catastrophes ($p \rightarrow 1^-$). Whether the binomial bulk-departure structure

survives the two-class non-preemptive setting — in effect a q -deformation of the Model- C_2 analysis — is open, and would add a synchronisation axis to the mechanism hierarchy of this thesis.

9. *Calibration.* Fitting γ_1 , θ_1 and the loads to real elderly-care waiting-list data [Zen99] would test whether the regimes identified here describe operational reality.

A final, self-contained item is the closure of the partial-PGF recursion on the alternative state space $\tilde{S}$, deferred from the analysis of Model- A (Section 5.2). A rigorous treatment would retain the boundary-diagonal forcing that Approximation 1 neglects and solve the resulting non-homogeneous recurrence by variation of constants in n — which, as anticipated there, reintroduces the unknown diagonal sequence $\{\pi(0, n)\}$ as the quantity to be determined self-consistently. Settling whether that recursion closes, and proving the conjectured monotonicity of $\mathbb{E}[N_2]$ in θ_1 raised in Section 11, would tie off the loose ends this thesis has deliberately left explicit.

References

- [AR02] Ivo Adan and Jacques Resing. Queueing theory. *Eindhoven University of Technology*, 180:10, 2002.
- [Coh56] Jacob Willem Cohen. Certain delay problems for a full availability trunk group loaded by two traffic sources. *Communication News*, 16:105–113, 1956.
- [DLM] NIST digital library of mathematical functions. <https://dlmf.nist.gov/>. Release 1.1.12, 2023. F. W. J. Olver, A. B. Olde Daalhuis, D. W. Lozier, B. I. Schneider, R. F. Boisvert, C. W. Clark, B. R. Miller, B. V. Saunders, H. S. Cohl, and M. A. McClain, eds.
- [dW92] Peter de Waal. *Approximate analysis of an M/G/1 priority queue with priority changes due to impatience*. CWI Report BS-R9202. CWI, Department of Operations Research, Statistics, and System Theory, Amsterdam, 1992.
- [DWFB20] Arnaud Devos, Joris Walraevens, Dieter Fiems, and Herwig Bruneel. Analysis of a discrete-time two-class randomly alternating service model with Bernoulli arrivals. *Queueing Systems*, 96:133–152, 2020.
- [FIM99] Guy Fayolle, Roudolf Iasnogorodski, and Vadim Malyshev. *Random Walks in the Quarter-Plane: Algebraic Methods, Boundary Value Problems and Applications*. Springer, 1999.
- [Gau67] Walter Gautschi. Computational aspects of three-term recurrence relations. *SIAM Review*, 9(1), 1967.
- [GMR02] Ofer Garnett, Avishai Mandelbaum, and Martin I. Reiman. Designing a call center with impatient customers. *Manufacturing & Service Operations Management*, 4(3), 2002.
- [HCD22] Yue Hu, Carri W. Chan, and Jing Dong. Optimal scheduling of proactive service with customer deterioration and improvement. *Management Science*, 2022.
- [JR14] Oualid Jouini and Alex Roubos. On multiple priority multi-server queues with impatience. *Journal of the Operational Research Society*, 65(5):616–632, 2014.
- [Kap11] Stella Kapodistria. The M/M/1 queue with synchronized abandonments. *Queueing Systems*, 68(1):79–109, 2011.
- [STZ14] David A. Stanford, Peter Taylor, and Ilze Ziedins. Waiting time distributions in the accumulating priority queue. *Queueing Systems*, 77:297–330, 2014.
- [WBSW15] Jianfu Wang, Opher Baron, and Alan Scheller-Wolf. M/M/c queue with two priority classes. *Operations Research*, 63(3):733–749, 2015.
- [Zen99] Stefanos A. Zenios. Modeling the transplant waiting list: A queueing model with reneging. *Operations Research*, 47(3), 1999.

- [ZK23] Josef Zuk and David Kirszenblat. Explicit results for the distributions of queue lengths for a non-preemptive two-level priority queue. *arXiv preprint arXiv:2309.09428*, 2023.
- [ZM05] Sergey Zeltyn and Avishai Mandelbaum. Call centers with impatient customers: Many-server asymptotics of the M/M/n+G queue. *Queueing Systems*, 51, 2005.